
AWS SDK for Go

Developer Guide

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

Table of Contents

AWS SDK for Go	1
Using the AWS SDK for Go with AWS Cloud9	1
More Info	1
Getting Started	2
Install the AWS SDK for Go	2
Get your AWS access keys	2
To get your access key ID and secret access key	2
Import Packages	3
Configuring the SDK	4
Specifying the AWS Region	4
Specifying Credentials	4
IAM Roles for Amazon EC2 Instances	5
Shared Credentials File	5
Environment Variables	7
Hard-Coded Credentials in an Application (Not Recommended)	7
Other Credentials Providers	7
Configuring a Proxy	7
Logging Service Calls	8
Creating a Custom Endpoint	8
Using Cloud9 with the SDK	9
Step 1: Set up Your AWS Account to Use AWS Cloud9	9
Step 2: Set up Your AWS Cloud9 Development Environment	9
Step 3: Set up the AWS SDK for Go	9
Step 4: Download Example Code	10
Step 5: Run Example Code	11
Using Sessions	12
Concurrency	12
Sessions with a Shared Configuration File	12
Creating Sessions	12
Create Sessions with Option Overrides	13
Deprecated New	13
Shared Configuration Fields	14
Environment Variables	7
Adding Request Handlers	14
Copying a Session	14
Using AWS Services	15
Constructing a Service	15
Tagging Service Resources	15
Getting the HTTP Request and Response with Each Service Call	17
Service Operation Calls	18
Calling Operations	18
Calling Operations with the Request Form	19
Handling Operation Response Body	19
Concurrently Using Service Clients	20
Using Pagination Methods	20
Using Waiters	21
Handling Errors	23
Handling Specific Service Error Codes	23
Additional Error Information	24
Specific Error Interfaces	24
Code Examples	25
SDK Request Examples	25
Using context.Context with SDK Requests	25
Using API Field Setters with SDK Requests	26

AWS CloudTrail Examples	26
Listing the CloudTrail Trails	27
Creating a CloudTrail Trail	28
Listing CloudTrail Trail Events	30
Deleting a CloudTrail Trail	31
Amazon CloudWatch Examples	32
Describing CloudWatch Alarms	33
Using Alarms and Alarm Actions in CloudWatch	34
Getting Metrics from CloudWatch	37
Sending Events to Amazon CloudWatch Events	39
AWS CodeBuild Examples	43
Getting Information about All AWS CodeBuild Projects	43
Building an AWS CodeBuild Project	44
Listing Your AWS CodeBuild Project Builds	44
Amazon DynamoDB Examples	45
Listing all Amazon DynamoDB Tables Using the AWS SDK for Go	46
Creating an Amazon DynamoDB Table Using the AWS SDK for Go	47
Creating an Amazon DynamoDB Table Item Using the AWS SDK for Go	48
Creating Amazon DynamoDB Table Items from a JSON File Using the AWS SDK for Go	49
Reading an Amazon DynamoDB Table Item	51
Getting Amazon DynamoDB Table Items Using Expression Builder	52
Updating an Amazon DynamoDB Table Item	54
Deleting an Amazon DynamoDB Table Item	55
Amazon EC2 Examples	56
Creating Amazon EC2 Instances with Tags or without Block Devices	56
Managing Amazon EC2 Instances	58
Working with Amazon EC2 Key Pairs	63
Using Regions and Availability Zones with Amazon EC2	66
Working with Security Groups in Amazon EC2	68
Using Elastic IP Addresses in Amazon EC2	73
Amazon Glacier Examples	77
The Scenario	77
Prerequisites	77
Create a Vault	78
Upload an Archive	78
AWS IAM Examples	78
Managing IAM Users	79
Managing IAM Access Keys	85
Managing IAM Account Aliases	89
Working with IAM Policies	92
Working with IAM Server Certificates	98
AWS KMS Examples	101
Creating a CMK in AWS Key Management Service	101
Encrypting Data with AWS Key Management Service	102
Decrypting a Data Blob in AWS Key Management Service	103
Re-encrypting a Data Blob in AWS Key Management Service	103
AWS Lambda Examples	104
Displaying Information about All Lambda Functions	105
Creating a Lambda Function	105
Running a Lambda Function	106
Configuring a Lambda Function to Receive Notifications	108
Amazon Polly Examples	109
Getting a List of Voices	109
Getting a List of Lexicons	110
Synthesizing Speech	111
Amazon S3 Examples	112
Performing Basic Amazon S3 Bucket Operations	112

Creating Pre-Signed URLs for Amazon S3 Buckets	125
Using an Amazon S3 Bucket as a Static Web Host	128
Working with Amazon S3 CORS Permissions	131
Working with Amazon S3 Bucket Policies	133
Working with Amazon S3 Bucket ACLs	136
Encrypting Amazon S3 Bucket Items	144
Amazon SES Examples	150
Listing Valid Amazon SES Email Addresses	150
Verifying an Email Address in Amazon SES	151
Sending a Message to an Email Address in Amazon SES	152
Deleting an Email Address in Amazon SES	154
Getting Amazon SES Statistics	155
Amazon SQS Examples	156
Using Amazon SQS Queues	156
Sending and Receiving Messages in Amazon SQS	160
Managing Visibility Timeout in Amazon SQS Queues	164
Enabling Long Polling in Amazon SQS Queues	166
Using Dead Letter Queues in Amazon SQS	170
Setting Attributes on an Amazon SQS Queue	172
Amazon WorkDocs Examples	174
Listing Users	174
Listing User Docs	175
SDK Utilities	178
Amazon CloudFront URL Signer	178
Amazon DynamoDB Attributes Converter	178
Amazon Elastic Compute Cloud Metadata	179
Retrieving an Instance's Region	179
Amazon S3 Transfer Managers	180
Upload Manager	180
Download Manager	183
Document History	186

AWS SDK for Go Developer Guide

The AWS SDK for Go provides APIs and utilities that developers can use to build Go applications that use AWS services, such as Amazon Elastic Compute Cloud (Amazon EC2) and Amazon Simple Storage Service (Amazon S3).

The SDK removes the complexity of coding directly against a web service interface. It hides a lot of the lower-level plumbing, such as authentication, request retries, and error handling.

The SDK also includes helpful utilities. For example, the Amazon S3 download and upload manager can automatically break up large objects into multiple parts and transfer them in parallel.

Use the *AWS SDK for Go Developer Guide* to help you install, configure, and use the SDK. The guide provides configuration information, sample code, and an introduction to the SDK utilities.

Using the AWS SDK for Go with AWS Cloud9

AWS Cloud9 is a web-based integrated development environment (IDE) that contains a collection of tools that you use to code, build, run, test, debug, and release software in the cloud.

See [Using AWS Cloud9 with the AWS SDK for Go \(p. 9\)](#) for information on using AWS Cloud9 with the AWS SDK for Go.

More Info

- To learn about everything you need before you can start using the AWS SDK for Go, see [Getting Started with the AWS SDK for Go \(p. 2\)](#).
- For code examples, see [AWS SDK for Go Code Examples \(p. 25\)](#).
- To learn about the SDK utilities, see [Using the AWS SDK for Go Utilities \(p. 178\)](#).
- For learn about the types and functionality that the library provides, see the [AWS SDK for Go API Reference](#).
- To view a video introduction of the SDK and a sample application demonstration, see [AWS SDK For Go: Gophers Get Going with AWS](#) from AWS re:Invent 2015.

Getting Started with the AWS SDK for Go

The AWS SDK for Go requires Go 1.5 or later. You can view your current version of Go by running the `go version` command. For information about installing or upgrading your version of Go, see <https://golang.org/doc/install>.

Install the AWS SDK for Go

To install the SDK and its dependencies, run the following Go command.

```
go get -u github.com/aws/aws-sdk-go/...
```

If you set the `Go vendor experiment` environment variable to 1, you can use the following command to get the SDK. The SDK's runtime dependencies are vendored in the `vendor/` folder.

```
go get -u github.com/aws/aws-sdk-go
```

Get your AWS access keys

Access keys consist of an *access key ID* and *secret access key*, which are used to sign programmatic requests that you make to AWS. If you don't have access keys, you can create them by using the [AWS Management Console](#). We recommend that you use IAM access keys instead of AWS root account access keys. IAM lets you securely control access to AWS services and resources in your AWS account.

Note

To create access keys, you must have permissions to perform the required IAM actions. For more information, see [Granting IAM User Permission to Manage Password Policy and Credentials](#) in the *IAM User Guide*.

To get your access key ID and secret access key

1. Open the [IAM console](#).
2. From the navigation menu, click **Users**.
3. Select your IAM user name.
4. Click **User Actions**, and then click **Manage Access Keys**.
5. Click **Create Access Key**.

Your keys will look something like this:

- Access key ID example: AKIAIOSFODNN7EXAMPLE
- Secret access key example: wJalrXUtnFEMI/K7MDENG/bPxrFc1YEXAMPLEKEY

6. Click **Download Credentials**, and store the keys in a secure location.

Important

Your secret key will no longer be available through the AWS Management Console; you will have the only copy. Keep it confidential in order to protect your account, and never email it. Do not share it outside your organization, even if an inquiry appears to come from AWS or Amazon.com.

No one who legitimately represents Amazon will ever ask you for your secret key.

Related topics

- [What Is IAM?](#) in *IAM User Guide*.
- [AWS Security Credentials](#) in *Amazon Web Services General Reference*.

Import Packages

After you have installed the SDK, you import AWS packages into your Go applications to use the SDK, as shown in the following example:

```
import "github.com/aws/aws-sdk-go/service/s3"
```


Configuring the AWS SDK for Go

In the AWS SDK for Go, you can configure settings for service clients, such as the log level and maximum number of retries. Most settings are optional. However, for each service client, you must specify a region and your credentials. The SDK uses these values to send requests to the correct AWS Region and sign requests with the correct credentials. You can specify these values as part of a session or as environment variables.

Specifying the AWS Region

When you specify the region, you specify where to send requests, such as `us-west-2` or `us-east-2`. For a list of regions for each service, see [Regions and Endpoints](#) in the *Amazon Web Services General Reference*.

The SDK does not have a default region. To specify a region:

- Set the `AWS_REGION` environment variable to the default region
- Set the `AWS_SDK_LOAD_CONFIG` environment variable to **true** to get the region value from the `config` file in the `.aws/` folder in your home directory
- Set the `NewSessionWithOptions` method argument `SharedConfigState` to `SharedConfigEnable` when you create a session to get the region value from the `config` file in the `.aws/` folder in your home directory
- Set the region explicitly when you create a session

If you set a region using all of these techniques, the SDK uses the region you explicitly specified in the session.

The following examples show you how to configure the environment variable.

Linux, OS X, or Unix

```
$ export AWS_REGION=us-west-2
```

Windows

```
> set AWS_REGION=us-west-2
```

The following snippet specifies the region in a session:

```
sess, err := session.NewSession(&aws.Config{Region: aws.String("us-west-2")})
```

Specifying Credentials

The AWS SDK for Go requires credentials (an access key and secret access key) to sign requests to AWS. You can specify your credentials in several different locations, depending on your particular use case. For information about obtaining credentials, see [Setting Up](#) (p. 2).

When you initialize a new service client without providing any credential arguments, the SDK uses the [default credential provider chain](#) to find AWS credentials. The SDK uses the first provider in the chain that returns credentials without an error. The default provider chain looks for credentials in the following order:

1. Environment variables.
2. Shared credentials file.
3. If your application is running on an Amazon EC2 instance, IAM role for Amazon EC2.

The SDK detects and uses the built-in providers automatically, without requiring manual configurations. For example, if you use IAM roles for Amazon EC2 instances, your applications automatically use the instance's credentials. You don't need to manually configure credentials in your application.

As a best practice, AWS recommends that you specify credentials in the following order:

1. Use IAM roles for Amazon EC2 (if your application is running on an Amazon EC2 instance).

IAM roles provide applications on the instance temporary security credentials to make AWS calls. IAM roles provide an easy way to distribute and manage credentials on multiple Amazon EC2 instances.

2. Use a shared credentials file.

This credentials file is the same one used by other SDKs and the AWS CLI. If you're already using a shared credentials file, you can also use it for this purpose.

3. Use environment variables.

Setting environment variables is useful if you're doing development work on a machine other than an Amazon EC2 instance.

4. Hard-code credentials (not recommended).

Hard-coding credentials in your application can make it difficult to manage and rotate those credentials. Use this method only for small personal scripts or testing purposes. Do not submit code with credentials to source control.

IAM Roles for Amazon EC2 Instances

If you are running your application on an Amazon EC2 instance, you can use the instance's [IAM role](#) to get temporary security credentials to make calls to AWS.

If you have configured your instance to use IAM roles, the SDK uses these credentials for your application automatically. You don't need to manually specify these credentials.

Shared Credentials File

A credential file is a plaintext file that contains your access keys. The file must be on the same machine on which you're running your application. The file must be named `credentials` and located in the `.aws/` folder in your home directory. The home directory can vary by operating system. In Windows, you can refer to your home directory by using the environment variable `%UserProfile%`. In Unix-like systems, you can use the environment variable `$HOME` or `~` (tilde).

If you already use this file for other SDKs and tools (like the AWS CLI), you don't need to change anything to use the files in this SDK. If you use different credentials for different tools or applications, you can use *profiles* to configure multiple access keys in the same configuration file.

Creating the Credentials File

If you don't have a shared credentials file (`.aws/credentials`), you can use any text editor to create one in your home directory. Add the following content to your credentials file, replacing `<YOUR_ACCESS_KEY_ID>` and `<YOUR_SECRET_ACCESS_KEY>` with your credentials.

```
[default]
```

```
aws_access_key_id = <YOUR_ACCESS_KEY_ID>
aws_secret_access_key = <YOUR_SECRET_ACCESS_KEY>
```

The `[default]` heading defines credentials for the default profile, which the SDK will use unless you configure it to use another profile.

You can also use temporary security credentials by adding the session tokens to your profile, as shown in the following example:

```
[temp]
aws_access_key_id = <YOUR_TEMP_ACCESS_KEY_ID>
aws_secret_access_key = <YOUR_TEMP_SECRET_ACCESS_KEY>
aws_session_token = <YOUR_SESSION_TOKEN>
```

Specifying Profiles

You can include multiple access keys in the same configuration file by associating each set of access keys with a profile. For example, in your credentials file, you can declare multiple profiles, as follows.

```
[default]
aws_access_key_id = <YOUR_DEFAULT_ACCESS_KEY_ID>
aws_secret_access_key = <YOUR_DEFAULT_SECRET_ACCESS_KEY>

[test-account]
aws_access_key_id = <YOUR_TEST_ACCESS_KEY_ID>
aws_secret_access_key = <YOUR_TEST_SECRET_ACCESS_KEY>

[prod-account]
; work profile
aws_access_key_id = <YOUR_PROD_ACCESS_KEY_ID>
aws_secret_access_key = <YOUR_PROD_SECRET_ACCESS_KEY>
```

By default, the SDK checks the `AWS_PROFILE` environment variable to determine which profile to use. If no `AWS_PROFILE` variable is set, the SDK uses the default profile.

If you have an application named `myapp` that uses the SDK, you can run it with the test credentials by setting the variable to `test-account` `myapp`, as shown in the following command.

```
$ AWS_PROFILE=test-account myapp
```

You can also use the SDK to select a profile by specifying `os.Setenv("AWS_PROFILE", test-account)` before constructing any service clients or by manually setting the credential provider, as shown in the following example.

```
sess, err := session.NewSession(&aws.Config{
    Region:      aws.String("us-west-2"),
    Credentials: credentials.NewSharedCredentials("", "test-account"),
})
```

In addition, checking if your credentials have been found is fairly easy.

```
_, err := sess.Config.Credentials.Get()
```

If `ChainProvider` is being used, set `CredentialsChainVerboseErrors` to `true` in the session config.

Note

If you specify credentials in environment variables, the SDK will always use those credentials, no matter which profile you specify.

Environment Variables

By default, the SDK detects AWS credentials set in your environment and uses them to sign requests to AWS. That way you don't need to manage credentials in your applications.

The SDK looks for credentials in the following environment variables:

- `AWS_ACCESS_KEY_ID`
- `AWS_SECRET_ACCESS_KEY`
- `AWS_SESSION_TOKEN` (optional)

The following examples show how you configure the environment variables.

Linux, OS X, or Unix

```
$ export AWS_ACCESS_KEY_ID=YOUR_AKID
$ export AWS_SECRET_ACCESS_KEY=YOUR_SECRET_KEY
$ export AWS_SESSION_TOKEN=TOKEN
```

Windows

```
> set AWS_ACCESS_KEY_ID=YOUR_AKID
> set AWS_SECRET_ACCESS_KEY=YOUR_SECRET_KEY
> set AWS_SESSION_TOKEN=TOKEN
```

Hard-Coded Credentials in an Application (Not Recommended)

Warning

Do not embed credentials inside an application. Use this method only for testing purposes.

You can hard-code credentials in your application by passing the access keys to a configuration instance, as shown in the following snippet.

```
sess, err := session.NewSession(&aws.Config{
    Region:      aws.String("us-west-2"),
    Credentials: credentials.NewStaticCredentials("AKID", "SECRET_KEY", "TOKEN"),
})
```

Other Credentials Providers

The SDK provides other methods for retrieving credentials in the `aws/credentials` package. For example, you can retrieve temporary security credentials from AWS Security Token Service or credentials from encrypted storage. For more information, see [Credentials](#).

Configuring a Proxy

If you cannot directly connect to the internet, you can use Go-supported environment variables (`HTTP_PROXY`) or create a custom HTTP client to configure your proxy. Use the [Config.HTTPClient](#) struct

to specify a custom HTTP client. For more information about how to create an HTTP client to use a proxy, see the [Transport](#) struct in the Go `http` package.

Logging Service Calls

You can enable logging in a client by setting the `LogLevel` in a configuration instance, as shown in the following snippet, which sets the log level to `LogDebugWithHTTPBody` for a new DynamoDB client.

```
svc := dynamodb.New(sess, aws.NewConfig().WithLogLevel(aws.LogDebugWithHTTPBody))
```

See [LogLevelType](#) for the different log level values.

Creating a Custom Endpoint

In most cases you use the endpoint that is pre-configured for a service. However, you can specify a custom endpoint, such as for pre-release versions of the service, as shown in the following snippet, which sets the `Endpoint` to `https://test.us-west-2.amazonaws.com` for a new DynamoDB client.

```
svc := dynamodb.New(sess, &aws.Config{Endpoint: aws.String("https://test.us-west-2.amazonaws.com")})
```

See [aws.Config](#) for details.

Using AWS Cloud9 with the AWS SDK for Go

You can use AWS Cloud9 with the AWS SDK for Go to write and run your Go code using just a browser. AWS Cloud9 includes tools such as a code editor and terminal. Because the AWS Cloud9 IDE is cloud based, you can work on your projects from your office, home, or anywhere using an internet-connected machine. For general information about AWS Cloud9, see the *AWS Cloud9 User Guide*.

Follow these instructions to set up AWS Cloud9 with the AWS SDK for Go:

- [Step 1: Set up Your AWS Account to Use AWS Cloud9 \(p. 9\)](#)
- [Step 2: Set up Your AWS Cloud9 Development Environment \(p. 9\)](#)
- [Step 3: Set up the AWS SDK for Go \(p. 9\)](#)
- [Step 4: Download Example Code \(p. 10\)](#)
- [Step 5: Run Example Code \(p. 11\)](#)

Step 1: Set up Your AWS Account to Use AWS Cloud9

Start to use AWS Cloud9 by signing in to the AWS Cloud9 console as an AWS Identity and Access Management (IAM) entity (for example, an IAM user) in your AWS account who has access permissions for AWS Cloud9.

To set up an IAM entity in your AWS account to access AWS Cloud9, and to sign in to the AWS Cloud9 console, see [Team Setup for AWS Cloud9](#) in the *AWS Cloud9 User Guide*.

Step 2: Set up Your AWS Cloud9 Development Environment

After you sign in to the AWS Cloud9 console, use the console to create an AWS Cloud9 development environment. After you create the environment, AWS Cloud9 opens the IDE for that environment.

See [Creating an Environment in AWS Cloud9](#) in the *AWS Cloud9 User Guide* for details.

Note

As you create your environment in the console for the first time, we recommend that you choose the option to **Create a new instance for environment (EC2)**. This option tells AWS Cloud9 to create an environment, launch an Amazon EC2 instance, and then connect the new instance to the new environment. This is the fastest way to begin using AWS Cloud9.

Step 3: Set up the AWS SDK for Go

After AWS Cloud9 opens the IDE for your development environment, use the IDE to set up the AWS SDK for Go in your environment, as follows.

1. If the terminal isn't already open in the IDE, open it. On the menu bar in the IDE, choose **Window, New Terminal**.
2. Set your GOPATH environment variable. To do this, add the following code to the end of your shell profile file (for example, `~/ .bashrc` in Amazon Linux, assuming you chose the option to **Create a new instance for environment (EC2)**, earlier in this topic), and then save the file.

```
GOPATH=~/.environment/go  
  
export GOPATH
```

After you save the file, source the `~/ .bashrc` file to finish setting your GOPATH environment variable. To do this, run the following command. (This command assumes you chose the option to **Create a new instance for environment (EC2)**, earlier in this topic.)

```
. ~/.bashrc
```

3. Run the following command to install the AWS SDK for Go.

```
go get -u github.com/aws/aws-sdk-go/...
```

If the IDE can't find Go, run the following commands, one at a time in this order, to install it. (These commands assume you chose the option to **Create a new instance for environment (EC2)**, earlier in this topic. Also, these commands assume the latest stable version of Go at the time this topic was written; for more information, see [Downloads](#) on The Go Programming Language website.)

```
wget https://storage.googleapis.com/golang/go1.9.3.linux-amd64.tar.gz # Download the Go  
installer.  
sudo tar -C /usr/local -xzf ./go1.9.3.linux-amd64.tar.gz # Install Go.  
rm ./go1.9.3.linux-amd64.tar.gz # Delete the Go  
installer, as you no longer need it.
```

After you install Go, add the path to the Go binary to your PATH environment variable. To do this, add the following code to the end of your shell profile file (for example, `~/ .bashrc` in Amazon Linux, assuming you chose the option to **Create a new instance for environment (EC2)**, earlier in this topic), and then save the file.

```
PATH=$PATH:/usr/local/go/bin
```

After you save the file, source the `~/ .bashrc` file so that the terminal can now find the Go binary you just referenced. To do this, run the following command. (This command assumes you chose the option to **Create a new instance for environment (EC2)**, earlier in this topic.)

```
. ~/.bashrc
```

Step 4: Download Example Code

Use the terminal you opened in the previous step to download example code for the AWS SDK for Go into the AWS Cloud9 development environment.

To do this, run the following command. This command downloads a copy of all of the code examples used in the official AWS SDK documentation into your environment's root directory.

```
git clone https://github.com/awsdocs/aws-doc-sdk-examples.git
```

To find code examples for the AWS SDK for Go, use the **Environment** window to open the `ENVIRONMENT_NAME/aws-doc-sdk-examples/go/example_code` directory, where `ENVIRONMENT_NAME` is the name of your development environment.

To learn how to work with these and other code examples, see [AWS SDK for Go Code Examples \(p. 25\)](#).

Step 5: Run Example Code

To run code in your AWS Cloud9 development environment, see [Run Your Code](#) in the *AWS Cloud9 User Guide*.

Using Sessions to Configure Service Clients in the AWS SDK for Go

In the AWS SDK for Go, a session is an object that contains configuration information for [service clients](#) (p. 15), which you use to interact with AWS services. For example, sessions can include information about the region where requests will be sent, which credentials to use, or additional request handlers. Whenever you create a service client, you must specify a session. For more information about sessions, see the [session](#) package in the *AWS SDK for Go API Reference*.

Sessions can be shared across all service clients that share the same base configuration. The session is built from the SDK's default configuration and request handlers.

You should cache sessions when possible. This is because creating a new session loads all configuration values from the environment and configuration files each time the session is created. Sharing the session value across all of your service clients ensures the configuration is loaded the fewest number of times.

Concurrency

Sessions are safe to use concurrently as long as the session isn't being modified. The SDK doesn't modify the session once the session is created. Creating service clients concurrently from a shared session is safe.

Sessions with a Shared Configuration File

Using the previous method, you can create sessions that load the additional configuration file only if the `AWS_SDK_LOAD_CONFIG` environment variable is set. Alternatively you can explicitly create a session with a shared configuration enabled. To do this, you can use `NewSessionWithOptions` to configure how the session is created. Using the `NewSessionWithOptions` with `SharedConfigState` set to `SharedConfigEnabled` will create the session as if the `AWS_SDK_LOAD_CONFIG` environment variable was set.

Creating Sessions

When you create a session, you can pass in optional `aws.Config` values that override the default or that override the current configuration values. This allows you to provide additional or case-based configuration as needed.

By default `NewSession` only loads credentials from the shared credentials file (`~/.aws/credentials`). If the `AWS_SDK_LOAD_CONFIG` environment variable is set to a truthy value, the session is created from the configuration values from the shared credentials (`~/.aws/credentials`) and shared credentials (`~/.aws/credentials`) files. See [Sessions with a Shared Configuration File](#) (p. 12) for more information.

Create a session with the default configuration and request handlers. The following example creates a session with credentials, region, and profile values from either the environment variables or the shared credentials file. It requires that the `AWS_PROFILE` is set, or `default` is used.

```
sess, err := session.NewSession()
```

The SDK provides a [default configuration](#) that all sessions use, unless you override a field. For example, you can specify an AWS Region when you create a session by using the `aws.Config` struct. For more information about the fields you can specify, see the [aws.Config](#) in the *AWS SDK for Go API Reference*.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-east-2")},
)
```

Create an Amazon S3 client instance from a session:

```
sess, err := session.NewSession()
if err != nil {
    // Handle Session creation error
}
svc := s3.New(sess)
```

Create Sessions with Option Overrides

In addition to `NewSession`, you can create sessions using `NewSessionWithOptions`. This function allows you to control and override how the session will be created through code, instead of being driven by environment variables only.

Use [NewSessionWithOptions](#) when you want to provide the config profile, or override the shared credentials state (`AWS_SDK_LOAD_CONFIG`).

```
// Equivalent to session.New
sess, err := session.NewSessionWithOptions(session.Options{})

// Specify profile to load for the session's config
sess, err := session.NewSessionWithOptions(session.Options{
    Profile: "profile_name",
})

// Specify profile for config and region for requests
sess, err := session.NewSessionWithOptions(session.Options{
    Config: aws.Config{Region: aws.String("us-east-2")},
    Profile: "profile_name",
})

// Force enable Shared Config support
sess, err := session.NewSessionWithOptions(session.Options{
    SharedConfigState: SharedConfigEnable,
})

// Assume an IAM role with MFA prompting for token code on stdin
sess := session.Must(session.NewSessionWithOptions(session.Options{
    AssumeRoleTokenProvider: stscreds.StdinTokenProvider,
    SharedConfigState: SharedConfigEnable,
}))
```

Deprecated New

The `New` function has been deprecated because it doesn't provide a good way to return errors that occur when loading the configuration files and values. Because of this, `NewSession` was created so errors can be retrieved when creating a session fails.

Shared Configuration Fields

By default, the SDK loads credentials from the shared credentials file `~/.aws/credentials`. Any other configuration values are provided by the environment variables, SDK defaults, and user-provided `aws.config` values.

If the `AWS_SDK_LOAD_CONFIG` environment variable is set, or the **SharedConfigLoadEnable** option is used to create the session (as shown in the following example), additional configuration information is also loaded from the shared configuration file (`~/.aws/config`), if it exists. If any configuration setting value differs between the two files, the value from the shared credentials file (`~/.aws/credentials`) takes precedence.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))
```

See the [session package's documentation](#) for more information on shared credentials setup.

Environment Variables

When a session is created, you can set several environment variables to adjust how the SDK functions, and what configuration data it loads when creating sessions. Environment values are optional. For credentials, you must set both an access key and a secret access key. Otherwise, Go ignores the one you've set. All environment variable values are strings unless otherwise noted.

See the [session package's documentation](#) for more information on environment variable setup.

Adding Request Handlers

You can add handlers to a session for processing HTTP requests. All service clients that use the session inherit the handlers. For example, the following handler logs every request and its payload made by a service client.

```
// Create a session, and add additional handlers for all service
// clients created with the Session to inherit. Adds logging handler.
sess, err := session.NewSession()
sess.Handlers.Send.PushFront(func(r *request.Request) {
    // Log every request made and its payload
    logger.Println("Request: %s/%s, Payload: %s",
        r.ClientInfo.ServiceName, r.Operation, r.Params)
})
```

Copying a Session

You can use the [Copy](#) method to create copies of sessions. Copying sessions is useful when you want to create multiple sessions that have similar settings. Each time you copy a session, you can specify different values for any field. For example, the following snippet copies the `sess` session while overriding the `Region` field to `us-east-2`:

```
usEast2Sess := sess.Copy(&aws.Config{Region: aws.String("us-east-2")})
```

Using the AWS SDK for Go with AWS Services

To make calls to an AWS service, you must first construct a service client instance with a session. A service client provides low-level access to every API action for that service. For example, you create an Amazon S3 service client to make calls to Amazon S3.

When you call service operations, you pass in input parameters as a struct. A successful call usually results in an output struct that you can use. For example, after you successfully call an Amazon S3 create bucket action, the action returns an output struct with the bucket's location.

For the list of service clients, including their methods and parameters, see the [AWS SDK for Go API Reference](#).

Constructing a Service

To construct a service client instance, use the `NewSession()` function. The following example creates an Amazon S3 service client.

```
sess, err := session.NewSession()
if err != nil {
    fmt.Println("Error creating session ", err)
    return
}
svc := s3.New(sess)
```

After you have a service client instance, you can use it to call service operations. For more information about configurations, see [Configuring the AWS SDK for Go \(p. 4\)](#).

When you create a service client, you can pass in custom configurations so that you don't need to create a session for each configuration. The SDK merges the two configurations, overriding session values with your custom configuration. For example, in the following snippet, the Amazon S3 client uses the `mySession` session but overrides the `Region` field with a custom value (`us-west-2`):

```
svc := s3.New(mySession, aws.NewConfig().WithRegion("us-west-2"))
```

Tagging Service Resources

You can tag service resources, such as Amazon S3 buckets, so that you can determine the costs of your service resources at whatever level of granularity you require.

The following examples shows how to tag the Amazon S3 bucket `MyBucket` with `Cost Center` tag with the value `123456` and `Stack` tag with the value `MyTestStack`.

```
package main

import (
```

```
"github.com/aws/aws-sdk-go/aws"
"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/s3"

"fmt"
)

// Tag S3 bucket MyBucket with cost center tag "123456" and stack tag "MyTestStack".
//
// See:
// http://docs.aws.amazon.com/awsaccountbilling/latest/aboutv2/cost-alloc-tags.html
func main() {
    // Pre-defined values
    bucket := "MyBucket"
    tagName1 := "Cost Center"
    tagValue1 := "123456"
    tagName2 := "Stack"
    tagValue2 := "MyTestStack"

    // Initialize a session in us-west-2 that the SDK will use to load credentials
    // from the shared credentials file. (~/.aws/credentials).
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )
    if err != nil {
        fmt.Println(err.Error())
        return
    }

    // Create S3 service client
    svc := s3.New(sess)

    // Create input for PutBucket method
    input := &s3.PutBucketTaggingInput{
        Bucket: aws.String(bucket),
        Tagging: &s3.Tagging{
            TagSet: []*s3.Tag{
                {
                    Key:   aws.String(tagName1),
                    Value: aws.String(tagValue1),
                },
                {
                    Key:   aws.String(tagName2),
                    Value: aws.String(tagValue2),
                },
            },
        },
    }

    _, err = svc.PutBucketTagging(input)
    if err != nil {
        fmt.Println(err.Error())
        return
    }

    // Now show the tags
    // Create input for GetBucket method
    input := &s3.GetBucketTaggingInput{
        Bucket: aws.String(bucket),
    }

    result, err := svc.GetBucketTagging(input)
    if err != nil {
        fmt.Println(err.Error())
        return
    }
}
```

```
numTags := len(result.TagSet)

if numTags > 0 {
    fmt.Println("Found", numTags, "Tag(s):")
    fmt.Println("")

    for _, t := range result.TagSet {
        fmt.Println("  Key: ", *t.Key)
        fmt.Println("  Value:", *t.Value)
        fmt.Println("")
    }
} else {
    fmt.Println("Did not find any tags")
}
}
```

Note that if a tag of the same name already exists, its value is overwritten by the new value.

Getting the HTTP Request and Response with Each Service Call

You can direct the AWS SDK for Go to display the HTTP request and response it sends and receives for each call by including a configuration option when constructing the service client.

The following examples uses the `DynamoDBListTables` operation to illustrate how to add a custom header to a service call.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/request"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"

    "fmt"
    "os"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load credentials
    // from the shared config file. (~/.aws/credentials).
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )
    if err != nil {
        fmt.Println("Error getting session:")
        fmt.Println(err)
        os.Exit(1)
    }

    // Create DynamoDB client
    // and expose HTTP requests/responses
    svc := dynamodb.New(sess, aws.NewConfig().WithLogLevel(aws.LogDebugWithHTTPBody))

    // Add "CustomHeader" header with value of 10
    svc.Handlers.Send.PushFront(func(r *request.Request) {
        r.HTTPRequest.Header.Set("CustomHeader", fmt.Sprintf("%d", 10))
    })
}
```

```
// Call ListTables just to see HTTP request/response
// The request should have the CustomHeader set to 10
_, _ = svc.ListTables(&dynamodb.ListTablesInput{})
}
```

If you run this program, the output should be similar to the following, where **ACCESS-KEY** is the access key of the user and **TABLE-1**, through **TABLE-N** are the names of the tables.

```
2017/10/25 11:10:57 DEBUG: Request dynamodb/ListTables Details:
---[ REQUEST POST-SIGN ]-----
POST / HTTP/1.1
Host: dynamodb.us-west-2.amazonaws.com
User-Agent: aws-sdk-go/1.10.34 (go1.8; windows; amd64)
Content-Length: 2
Accept-Encoding: identity
Authorization: AWS4-HMAC-SHA256 Credential=ACCESS-KEY/20171025/us-west-2/dynamodb/
aws4_request, SignedHeaders=accept-encoding;content-length;content-type;host;x-amz-date;x-
amz-target, Signature=9c92efe5d6c597cf29e4f7cc74de6dc2e39f8010a0d4957a397c59ef9cde21f2
Content-Type: application/x-amz-json-1.0
CustomHeader: 10
X-Amz-Date: 20171025T181057Z
X-Amz-Target: DynamoDB_20120810.ListTables

{}

2017/10/25 11:10:58 DEBUG: Response dynamodb/ListTables Details:
---[ RESPONSE ]-----
HTTP/1.1 200 OK
Content-Length: 177
Connection: keep-alive
Content-Type: application/x-amz-json-1.0
Date: Wed, 25 Oct 2017 18:10:58 GMT
Server: Server
X-Amz-Crc32: 3023160996
X-Amzn-Requestid: M5B4BM4UU569MVBSDG5O2O9ITJVV4KQNSO5AEMVJF66Q9ASUAAJG

-----
2017/10/25 11:10:58 {"TableNames":["TABLE-1","...", "TABLE-N"]}
```

Service Operation Calls

You can call a service operation directly or with its request form. When you call a service operation, the SDK synchronously validates the input, builds the request, signs it with your credentials, sends it to AWS, and then gets a response or an error. In most cases, you can call service operations directly.

Calling Operations

Calling the operation will sync as the request is built, signed, sent, and the response is received. If an error occurs during the operation, it will be returned. The output or resulting structure won't be valid.

For example, to call the Amazon S3 GET Object API, use the Amazon S3 service client instance and call its [GetObject](#) method:

```
result, err := s3Svc.GetObject(&s3.GetObjectInput{...})
// result is a *s3.GetObjectOutput struct pointer
// err is a error which can be cast to awserr.Error.
```

Passing Parameters to a Service Operation

When calling an operation on a service, you pass in input parameters as option values, similar to passing in a configuration. For example, to retrieve an object, you must specify a bucket and the object's key by passing in the following parameters to the `GetObject` method:

```
svc := s3.New(session.New())
svc.GetObject(&s3.GetObjectInput{
    Bucket: aws.String("bucketName"),
    Key:    aws.String("keyName"),
})
```

Each service operation has an associated input struct and, usually, an output struct. The structs follow the naming pattern *OperationName* Input and *OperationName* Output.

For more information about the parameters of each method, see the service client documentation in the [AWS SDK for Go API Reference](#).

Calling Operations with the Request Form

Calling the request form of a service operation, which follows the naming pattern *OperationName* Request, provides a simple way to control when a request is built, signed, and sent. Calling the request form immediately returns a request object. The request object output is a struct pointer that is not valid until the request is sent and returned successfully.

Calling the request form can be useful when you want to construct a number of pre-signed requests, such as pre-signed Amazon S3 URLs. You can also use the request form to modify how the SDK sends a request.

The following example calls the request form of the `GetObject` method. The `Send` method signs the request before sending it.

```
req, result := s3Svc.GetObjectRequest(&s3.GetObjectInput{...})
// result is a *s3.GetObjectOutput struct pointer, not populated until req.Send() returns
// req is a *aws.Request struct pointer. Used to Send request.
if err := req.Send(); err != nil {
    // process error
    return
}
// Process result
```

Handling Operation Response Body

Some API operations return a response struct that contain a `Body` field that is an `io.ReadCloser`. If you're making requests with these operations, always be sure to call `Close` on the field.

```
resp, err := s3svc.GetObject(&s3.GetObjectInput{...})
if err != nil {
    // handle error
    return
}
// Make sure to always close the response Body when finished
defer resp.Body.Close()

decoder := json.NewDecoder(resp.Body)
if err := decoder.Decode(&myStruct); err != nil {
    // handle error
    return
}
```



```
}
```

Concurrently Using Service Clients

You can create goroutines that concurrently use the same service client to send multiple requests. You can use a service client with as many goroutines as you want. However, you cannot concurrently modify the service client's configuration and request handlers. If you do, the service client operations might encounter race conditions. Define service client settings before you concurrently use it.

In the following example, an Amazon S3 service client is used in multiple goroutines. The example concurrently outputs all objects in `bucket1`, `bucket2`, and `bucket3`, which are all in the same region. To make sure all objects from the same bucket are printed together, the example uses a channel.

```
sess, err := session.NewSession()
if err != nil {
    fmt.Println("Error creating session ", err)
}
var wg sync.WaitGroup
keysCh := make(chan string, 10)

svc := s3.New(sess)
buckets := []string{"bucket1", "bucket2", "bucket3"}
for _, bucket := range buckets {
    params := &s3.ListObjectsInput{
        Bucket: aws.String(bucket),
        MaxKeys: aws.Int64(100),
    }
    wg.Add(1)
    go func(param *s3.ListObjectsInput) {
        defer wg.Done()

        err = svc.ListObjectsPages(params,
            func(page *s3.ListObjectsOutput, last bool) bool {
                // Add the objects to the channel for each page
                for _, object := range page.Contents {
                    keysCh <- fmt.Sprintf("%s:%s", *params.Bucket, *object.Key)
                }
                return true
            },
        )
        if err != nil {
            fmt.Println("Error listing", *params.Bucket, "objects:", err)
        }
    }(params)
}
go func() {
    wg.Wait()
    close(keysCh)
}()
for key := range keysCh {
    // Print out each object key as its discovered
    fmt.Println(key)
}
```

Using Pagination Methods

Typically, when you retrieve a list of items, you might need to check the output for a token or marker to confirm whether AWS returned all results from your request. If present, you use the token or marker

to request the next set of results. Instead of managing these tokens or markers, you can use pagination methods provided by the SDK.

Pagination methods iterate over a list operation until the method retrieves the last page of results or until the callback function returns `false`. The names of these method use the following pattern: *OperationName* Pages. For example, the pagination method for the Amazon S3 list objects operation (`ListObjects`) is `ListObjectPages`.

The following example uses the `ListObjectPages` pagination method to list, at most, three pages of object keys from the `ListObject` operation. Each page consists of at least 10 keys, which is defined by the `MaxKeys` field.

```
svc, err := s3.NewSession(sess)
if err != nil {
    fmt.Println("Error creating session ", err)
}
inputparams := &s3.ListObjectsInput{
    Bucket:  aws.String("mybucket"),
    MaxKeys: aws.Int64(10),
}
pageNum := 0
svc.ListObjectPages(inputparams, func(page *s3.ListObjectsOutput, lastPage bool) bool {
    pageNum++
    for _, value := range page.Contents {
        fmt.Println(*value.Key)
    }
    return pageNum < 3
})
```

Using Waiters

The SDK provides waiters that continuously check for completion of a job. For example, when you send a request to create an Amazon S3 bucket, you can use a waiter to check when the bucket has been successfully created. That way, subsequent operations on the bucket are done only after the bucket has been created.

The following example uses a waiter that waits until specific instances have stopped.

```
sess, err := session.NewSession(aws.NewConfig().WithRegion("us-west-2"))
if err != nil {
    fmt.Println("Error creating session ", err)
}
// Create an EC2 client
ec2client := ec2.New(sess)
// Specify two instances to stop
instanceIDsToStop := aws.StringSlice([]string{"i-12345678", "i-23456789"})
// Send request to stop instances
_, err = ec2client.StopInstances(&ec2.StopInstancesInput{
    InstanceIds: instanceIDsToStop,
})
if err != nil {
    panic(err)
}
// Use a waiter function to wait until the instances are stopped
describeInstancesInput := &ec2.DescribeInstancesInput{
    InstanceIds: instanceIDsToStop,
}
if err := ec2client.WaitUntilInstanceStopped(describeInstancesInput); err != nil {
    panic(err)
}
```

```
fmt.Println("Instances are stopped.")
```

Handling Errors in the AWS SDK for Go

The AWS SDK for Go returns errors that satisfy the Go `error` interface type and the [Error](#) interface in the `aws/awserr` package. You can use the `Error()` method to get a formatted string of the SDK error message without any special handling.

```
if err != nil {
    if awsErr, ok := err.(awserr.Error); ok {
        // process SDK error
    }
}
```

Errors returned by the SDK are backed by a concrete type that will satisfy the `awserr.Error` interface. The interface has the following methods, which provide classification and information about the error.

- `Code` returns the classification code by which related errors are grouped.
- `Message` returns a description of the error.
- `OrigErr` returns the original error of type `error` that is wrapped by the `awserr.Error` interface, such as a standard library error or a service error.

Handling Specific Service Error Codes

The following example demonstrates how to handle error codes that you encounter while using the AWS SDK for Go. The example assumes you have already set up and configured the SDK (that is, all required packages are imported and your credentials and region are set). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

This example highlights how you can use the `awserr.Error` type to perform logic based on specific error codes returned by service API operations.

In this example the `S3 GetObject` API operation is used to request the contents of an object in S3. The example handles the `NoSuchBucket` and `NoSuchKey` error codes, printing custom messages to `stderr`. If any other error is received, a generic message is printed.

```
svc := s3.New(sess)
resp, err := svc.GetObject(&s3.GetObjectInput{
    Bucket: aws.String(os.Args[1]),
    Key:    aws.String(os.Args[2]),
})

if err != nil {
    // Casting to the awserr.Error type will allow you to inspect the error
    // code returned by the service in code. The error code can be used
    // to switch on context specific functionality. In this case a context
    // specific error message is printed to the user based on the bucket
    // and key existing.
    //
    // For information on other S3 API error codes see:
    // http://docs.aws.amazon.com/AmazonS3/latest/API/ErrorResponses.html
    if aerr, ok := err.(awserr.Error); ok {
```

```
switch aerr.Code() {
case s3.ErrCodeNoSuchBucket:
    exitErrorf("bucket %s does not exist", os.Args[1])
case s3.ErrCodeNoSuchKey:
    exitErrorf("object with key %s does not exist in bucket %s", os.Args[2],
os.Args[1])
}
```

You can see the complete example code on [GitHub](#).

Additional Error Information

In addition to the `awserr.Error` interface, you might be able to use other interfaces to get more information about an error.

Specific Error Interfaces

Other packages might provide their own error interfaces. For example, the [service/s3/s3manager](#) package provides a [MultiUploadFailure](#) interface to retrieve the upload ID. This is helpful when you must manually clean up a failed multi-part upload.

```
output, err := s3manager.Upload(svc, input, opts)
if err != nil {
    if multierr, ok := err.(MultiUploadFailure); ok {
        // Process error and its associated uploadID
        fmt.Println("Error:", multierr.Code(), multierr.Message(), multierr.UploadID())
    } else {
        // Process error generically
        fmt.Println("Error:", err.Error())
    }
}
```

For more information, see the [s3Manager.MultiUploadFailure](#) interface in the *AWS SDK for Go API Reference*.

AWS SDK for Go Code Examples

The AWS SDK for Go examples can help you write your own Go applications that use Amazon Web Services. The examples assume you have already set up and configured the SDK (that is, you have imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

Topics

- [AWS SDK for Go Request Examples \(p. 25\)](#)
- [AWS CloudTrail Examples Using the AWS SDK for Go \(p. 26\)](#)
- [Amazon CloudWatch Examples Using the AWS SDK for Go \(p. 32\)](#)
- [AWS CodeBuild Examples Using the AWS SDK for Go \(p. 43\)](#)
- [Amazon DynamoDB Examples Using the AWS SDK for Go \(p. 45\)](#)
- [Amazon EC2 Examples Using the AWS SDK for Go \(p. 56\)](#)
- [Amazon Glacier Examples Using the AWS SDK for Go \(p. 77\)](#)
- [IAM Examples Using the AWS SDK for Go \(p. 78\)](#)
- [AWS Key Management Service Examples Using the AWS SDK for Go \(p. 101\)](#)
- [AWS Lambda Examples Using the AWS SDK for Go \(p. 104\)](#)
- [Amazon Polly Examples Using the AWS SDK for Go \(p. 109\)](#)
- [Amazon S3 Examples Using the AWS SDK for Go \(p. 112\)](#)
- [Amazon SES Examples Using the AWS SDK for Go \(p. 150\)](#)
- [Amazon SQS Examples Using the AWS SDK for Go \(p. 156\)](#)
- [Amazon WorkDocs Examples \(p. 174\)](#)

AWS SDK for Go Request Examples

The AWS SDK for Go examples can help you write your own applications. The examples assume you have already set up and configured the SDK (that is, you have imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

Using context.Context with SDK Requests

In Go 1.7, the `context.Context` type was added to `http.Request`. This type provides an easy way to implement deadlines and cancellations on requests.

To use this pattern with the SDK, you call `WithContext` on the `HTTPRequest` field of the SDK's `request.Request` type, and provide your `Context` value. The following example highlights this process with a timeout on an Amazon `SQS.ReceiveMessage` API call.

```
ctx, cancel := context.WithTimeout(context.Background(), 5*time.Second)
defer cancel()

// SQS ReceiveMessage
params := &sqs.ReceiveMessageInput{ ... }
```

```
req, resp := s.ReceiveMessageRequest(params)
req.HTTPRequest = req.HTTPRequest.WithContext(ctx)
err := req.Send()
```

Using API Field Setters with SDK Requests

In addition to setting API parameters by using struct fields, you can also use chainable setters on the API operation parameter fields. This enables you to use a chain of setters to set the fields of the API struct.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "strings"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, _ := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create S3 service client
    svc := s3.New(sess)

    svc.PutObject((&s3.PutObjectInput{}).
        SetBucket("myBucket").
        SetKey("myKey").
        SetBody(strings.NewReader("object body")).
        SetWebsiteRedirectLocation("https://example.com/something"),
    )
}
```

You can also use this pattern with nested fields in API operation requests.

```
resp, err := svc.UpdateService((&ecs.UpdateServiceInput{}).
    SetService("myService").
    SetDeploymentConfiguration((&ecs.DeploymentConfiguration{}).
        SetMinimumHealthyPercent(80),
    ),
)
```

AWS CloudTrail Examples Using the AWS SDK for Go

CloudTrail is an AWS service that helps you enable governance, compliance, and operational and risk auditing of your AWS account. Actions taken by a user, role, or an AWS service are recorded as events in CloudTrail. Events include actions taken in the AWS Management Console, AWS Command Line Interface, and AWS SDKs and APIs.

The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Listing the CloudTrail Trails \(p. 27\)](#)
- [Creating a CloudTrail Trail \(p. 28\)](#)
- [Listing CloudTrail Trail Events \(p. 30\)](#)
- [Deleting a CloudTrail Trail \(p. 31\)](#)

Listing the CloudTrail Trails

This example uses the [DescribeTrails](#) operation to list the names of the CloudTrail trails and the bucket in which CloudTrail stores information in the `us-west-2` region.

Choose Copy to save the code locally.

Create the file `describe_trails.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudtrail"  
  
    "fmt"  
    "os"  
)
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `.aws/credentials` in your home folder, and create a new service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create CloudTrail client  
svc := cloudtrail.New(sess)
```

Call **DescribeTrails**. If an error occurs, print the error and exit. If no error occurs, loop through the trails, printing the name of each trail and the bucket.

```
resp, err := svc.DescribeTrails(&cloudtrail.DescribeTrailsInput{TrailNameList: nil})  
if err != nil {  
    fmt.Println("Got error calling CreateTrail:")  
    fmt.Println(err.Error())  
    os.Exit(1)  
}  
  
fmt.Println("Found", len(resp.TrailList), "trail(s) in", regionName)  
fmt.Println("")  
  
for _, trail := range resp.TrailList {  
    fmt.Println("Trail name: " + *trail.Name)  
    fmt.Println("Bucket name: " + *trail.S3BucketName)  
    fmt.Println("")  
}
```


See the [complete example](#) on GitHub.

Creating a CloudTrail Trail

This example uses the [CreateTrail](#) operation to create a CloudTrail trail in the `us-west-2` region. It requires two inputs, the name of the trail and the name of the bucket in which CloudTrail stores information. If the bucket does not have the proper policy, include the `-p` flag to attach the correct policy to the bucket.

Choose Copy to save the code locally.

Create the file `create_trail.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "github.com/aws/aws-sdk-go/service/cloudtrail"
    "github.com/aws/aws-sdk-go/service/sts"

    "encoding/json"
    "flag"
    "fmt"
    "os"
)
```

Get the names of the trail and bucket, and whether to attach the policy to the bucket. If either the trail name or bucket name is missing, display an error message and exit.

```
var trailName string
flag.StringVar(&trailName, "n", "", "The name of the trail")
var bucketName string
flag.StringVar(&bucketName, "b", "", "the name of bucket to which the trails are uploaded")

// Option to add CloudTrail policy to bucket
var addPolicy bool
flag.BoolVar(&addPolicy, "p", false, "Whether to add the CloudTrail policy to the bucket")

flag.Parse()

if trailName == "" || bucketName == "" {
    fmt.Println("You must supply a trail name and bucket name.")
    os.Exit(0)
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `.aws/credentials` in your home folder.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)
```

If the `-p` flag was specified, add a policy to the bucket.

```
if addPolicy {
    svc := sts.New(sess)
```

```

input := &sts.GetCallerIdentityInput{}

result, err := svc.GetCallerIdentity(input)
if err != nil {
    fmt.Println("Got error snarfing caller identity:")
    fmt.Println(err.Error())
    os.Exit(1)
}

accountId := aws.StringValue(result.Account)

s3Policy := map[string]interface{}{
    "Version": "2012-10-17",
    "Statement": []map[string]interface{}{
        {
            "Sid": "AWSCloudTrailAclCheck20150319",
            "Effect": "Allow",
            "Principal": map[string]interface{}{
                "Service": "cloudtrail.amazonaws.com",
            },
            "Action": "s3:GetBucketAcl",
            "Resource": "arn:aws:s3:::" + bucketName,
        },
        {
            "Sid": "AWSCloudTrailWrite20150319",
            "Effect": "Allow",
            "Principal": map[string]interface{}{
                "Service": "cloudtrail.amazonaws.com",
            },
            "Action": "s3:PutObject",
            "Resource": "arn:aws:s3:::" + bucketName + "/AWSLogs/" + accountId + "/*",
            "Condition": map[string]interface{}{
                "StringEquals": map[string]interface{}{
                    "s3:x-amz-acl": "bucket-owner-full-control",
                },
            },
        },
    },
}

policy, err := json.Marshal(s3Policy)
if err != nil {
    fmt.Println("Error marshalling request")
    os.Exit(0)
}

// Create S3 service
s3_svc := s3.New(sess)

// Now set the policy on the bucket
_, err = s3_svc.PutBucketPolicy(&s3.PutBucketPolicyInput{
    Bucket: aws.String(bucketName),
    Policy: aws.String(string(policy)),
})
if err != nil {
    fmt.Print("Got error adding bucket policy:")
    fmt.Print(err.Error())
    os.Exit(1)
}

fmt.Printf("Successfully set bucket %q's policy\n", bucketName)
}

```

Create the CloudTrail client, the input for **CreateTrail**, and call **CreateTrail**. If an error occurs, print the error and exit. If no error occurs, print a success message.

```
svc := cloudtrail.New(sess)

input := &cloudtrail.CreateTrailInput{
    Name: aws.String(trailName),
    S3BucketName: aws.String(bucketName),
}

_, err = svc.CreateTrail(input)
if err != nil {
    fmt.Println("Got error calling CreateTrail:")
    fmt.Println(err.Error())
    os.Exit(1)
}

fmt.Println("Created the trail", trailName, "for bucket", bucketName, "in region",
    regionName)
```

See the [complete example](#) on GitHub.

Listing CloudTrail Trail Events

This example uses the [LookupEvents](#) operation to list the CloudTrail trail events in the `us-west-2` region.

Choose Copy to save the code locally.

Create the file `lookup_events.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/cloudtrail"

    "flag"
    "fmt"
    "os"
    "time"
)
```

Get the name of the trail, and whether to display the event. If the trail name is missing, display an error message and exit.

```
// Trail name required
var trailName string
flag.StringVar(&trailName, "n", "", "The name of the trail")

// Option to show event
var showEvent bool
flag.BoolVar(&showEvent, "s", false, "Whether to show the event")

flag.Parse()

if trailName == "" {
    fmt.Println("You must supply a trail name")
    os.Exit(1)
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `.aws/credentials` in your home folder, and create a new service client.

```
// Initialize a session in us-west-2 that the SDK will use to load
// credentials from the shared credentials file ~/.aws/credentials.
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create CloudTrail client
svc := cloudtrail.New(sess)
```

Create the input for and call **LookupEvents**. If an error occurs, print the error and exit. If no error occurs, loop through the events, printing information about each event. If the **-s** flag was specified, print the CloudTrail event.

```
input := &cloudtrail.LookupEventsInput{EndTime: aws.Time(time.Now())}

resp, err := svc.LookupEvents(input)
if err != nil {
    fmt.Println("Got error calling CreateTrail:")
    fmt.Println(err.Error())
    os.Exit(1)
}

fmt.Println("Found", len(resp.Events), "events before now")
fmt.Println("")

for _, event := range resp.Events {
    if showEvents {
        fmt.Println("Event:")
        fmt.Println(aws.StringValue(event.CloudTrailEvent))
        fmt.Println("")
    }

    fmt.Println("Name      ", aws.StringValue(event.EventName))
    fmt.Println("ID:       ", aws.StringValue(event.EventId))
    fmt.Println("Time:     ", aws.TimeValue(event.EventTime))
    fmt.Println("User:     ", aws.StringValue(event.Username))

    fmt.Println("Resources:")

    for _, resource := range event.Resources {
        fmt.Println("  Name:", aws.StringValue(resource.ResourceName))
        fmt.Println("  Type:", aws.StringValue(resource.ResourceType))
    }

    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Deleting a CloudTrail Trail

This example uses the [DeleteTrail](#) operation to delete a CloudTrail trail in the `us-west-2` region. It requires one input, the name of the trail.

Choose Copy to save the code locally.

Create the file `delete_trail.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
```

```
"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/cloudtrail"

"flag"
"fmt"
"os"

)
```

Get the name of the trail. If the trail name is missing, display an error message and exit.

```
var trailName string
flag.StringVar(&trailName, "n", "", "The name of the trail to delete")

flag.Parse()

if trailName == "" {
    fmt.Println("You must supply a trail name")
    os.Exit(1)
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file *.aws/credentials* in your home folder and create the client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create CloudTrail client
svc := cloudtrail.New(sess)
```

Call **DeleteTrail** with the trail name. If an error occurs, print the error and exit. If no error occurs, print a success message.

```
_, err = svc.DeleteTrail(&cloudtrail.DeleteTrailInput{Name: &trailNamePtr})
if err != nil {
    fmt.Println("Got error calling CreateTrail:")
    fmt.Println(err.Error())
    os.Exit(1)
}

fmt.Println("Successfully deleted trail", trailName)
```

See the [complete example](#) on GitHub.

Amazon CloudWatch Examples Using the AWS SDK for Go

Amazon CloudWatch is a web service that monitors your AWS resources and the applications you run on AWS in real time. You can use CloudWatch to collect and track metrics, which are variables you can measure for your resources and applications. CloudWatch alarms send notifications or automatically make changes to the resources you are monitoring based on rules that you define.

The AWS SDK for Go examples show you how to integrate CloudWatch into your Go applications. The examples assume you have already set up and configured the SDK (that is, you have imported all

required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Describing CloudWatch Alarms \(p. 33\)](#)
- [Using Alarms and Alarm Actions in CloudWatch \(p. 34\)](#)
- [Getting Metrics from CloudWatch \(p. 37\)](#)
- [Sending Events to Amazon CloudWatch Events \(p. 39\)](#)

Describing CloudWatch Alarms

This example shows you how to retrieve basic information that describes your CloudWatch alarms.

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

An alarm watches a single metric over a time period you specify. The alarm performs one or more actions based on the value of the metric relative to a given threshold over a number of time periods.

In this example, Go code is used to describe alarms in CloudWatch. The code uses the AWS SDK for Go to describe alarms by using this method of the `AWS.CloudWatch` client class:

- [DescribeAlarms](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with CloudWatch alarms. To learn more, see [Creating Amazon CloudWatch Alarms](#) in the *Amazon CloudWatch User Guide*.

Describe Alarms

Choose **Copy** to save the code locally.

Create the file `describe_alarms.go`. Import the packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudwatch"  
  
    "fmt"  
    "os"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

svc := cloudwatch.New(sess)
```

Call `DescribeAlarms`, and print the results.

```
resp, err := svc.DescribeAlarms(nil)
for _, alarm := range resp.MetricAlarms {
    fmt.Println(*alarm.AlarmName)
}
```

See the [complete example](#) on GitHub.

Using Alarms and Alarm Actions in CloudWatch

These Go examples show you how to change the state of your Amazon EC2 instances automatically based on a CloudWatch alarm, as follows:

- Creating and enabling actions on an alarm
- Disabling actions on an alarm

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

You can use alarm actions to create alarms that automatically stop, terminate, reboot, or recover your Amazon EC2 instances. You can use the stop or terminate actions when you no longer need an instance to be running. You can use the reboot and recover actions to automatically reboot the instance.

In this example, Go code is used to define an alarm action in CloudWatch that triggers the reboot of an Amazon EC2 instance. The code uses the AWS SDK for Go to manage instances by using these methods of `PutMetricAlarm` type:

- [PutMetricAlarm](#)
- [EnableAlarmActions](#)
- [DisableAlarmActions](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with CloudWatch alarm actions. To learn more, see [Create Alarms to Stop, Terminate, Reboot, or Recover an Instance](#) in the *Amazon CloudWatch User Guide*.

Create and Enable Actions on an Alarm

Choose **Copy** to save the code locally.

Create the file `create_enable_alarms.go`.

Import packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudwatch"  
  
    "fmt"  
    "os"  
)
```

Get an instance name, value, and alarm name.

```
if len(os.Args) != 4 {  
    fmt.Println("You must supply an instance name, value, and alarm name")  
    os.Exit(1)  
}  
  
instance := os.Args[1]  
value := os.Args[2]  
name := os.Args[3]
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
// Create new cloudwatch client.  
svc := cloudwatch.New(sess)
```

Create a metric alarm that reboots an instance if its CPU utilization is greater than 70 percent.

```
_, err := svc.PutMetricAlarm(&cloudwatch.PutMetricAlarmInput{  
    AlarmName:      aws.String(name)  
    ComparisonOperator: aws.String(cloudwatch.ComparisonOperatorGreaterThanThreshold),  
    EvaluationPeriods: aws.Int64(1),  
    MetricName:      aws.String("CPUUtilization"),  
    Namespace:      aws.String("AWS/EC2"),  
    Period:          aws.Int64(60),  
    Statistic:       aws.String(cloudwatch.StatisticAverage),  
    Threshold:       aws.Float64(70.0),  
    ActionsEnabled:  aws.Bool(true),  
    AlarmDescription: aws.String("Alarm when server CPU exceeds 70%"),  
    Unit:            aws.String(cloudwatch.StandardUnitSeconds),  
  
    // This is apart of the default workflow actions. This one will reboot the instance, if  
    the  
    // alarm is triggered.  
    AlarmActions: []*string{  
        aws.String(fmt.Sprintf("arn:aws:swf:us-east-1:%s:action/actions/  
AWS_EC2.InstanceId.Reboot/1.0", instance)),  
    },  
    Dimensions: []*cloudwatch.Dimension{  
        &cloudwatch.Dimension{  
            Name:  aws.String("InstanceId"),  
            Value: aws.String(value),  
        },  
    },  
})
```



```
    },  
  })  
}
```

Call `EnableAlarmActions` with the new alarm for the instance, and display a message.

```
result, err := svc.EnableAlarmActions(&cloudwatch.EnableAlarmActionsInput{  
    AlarmNames: []*string{  
        aws.String(name),  
    },  
})  
fmt.Println("Alarm action enabled", result)
```

See the [complete example](#) on GitHub.

Disable Actions on an Alarm

Choose **Copy** to save the code locally.

Create the file `disable_alarm.go`.

Import the packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudwatch"  
  
    "fmt"  
    "os"  
)
```

Get the name of the alarm from the command line.

```
if len(os.Args) != 2 {  
    fmt.Println("You must supply an alarm name")  
    os.Exit(1)  
}  
  
name := os.Args[1]
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
// Create new CloudWatch client.  
svc := cloudwatch.New(sess)
```

Call `DisableAlarmActions` to disable the actions for this alarm and display a message.

```
_, err := svc.DisableAlarmActions(&cloudwatch.DisableAlarmActionsInput{  
    AlarmNames: []*string{  
        aws.String(name),  
    },  
})  
fmt.Println("Success")
```

See the [complete example](#) on GitHub.

Getting Metrics from CloudWatch

These Go examples show you how to retrieve a list of published CloudWatch metrics and publish data points to CloudWatch metrics with the AWS SDK for Go, as follows:

- Listing metrics
- Submitting custom metrics

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

Metrics are data about the performance of your systems. You can enable detailed monitoring of some resources, such as your Amazon EC2 instances, or your own application metrics.

In this example, Go code is used to get metrics from CloudWatch and to send events to CloudWatch Events. The code uses the AWS SDK for Go to get metrics from CloudWatch by using these methods of the CloudWatch type:

- [ListMetrics](#)
- [PutMetricData](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with CloudWatch metrics. To learn more, see [Using Amazon CloudWatch Metrics](#) in the *Amazon CloudWatch User Guide*.

List Metrics

Choose **Copy** to save the code locally.

Create the file `listing_metrics.go`.

Import the packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudwatch"  
  
    "fmt"  
    "os"  
)
```

Get the metric name, namespace, and dimensions from the command line.

```
if len(os.Args) != 4 {  
    fmt.Println("You must supply a metric name, namespace, and dimensions")  
    os.Exit(1)  
}
```

```
metric := os.Args[1]
namespace := os.Args[2]
dimensions := os.Args[3]
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create CloudWatch client
svc := cloudwatch.New(sess)
```

Call `ListMetrics`, supplying the metric name, namespace, and dimensions. Print the metrics returned in the result.

```
result, err := svc.ListMetrics(&cloudwatch.ListMetricsInput{
    MetricName: aws.String(metric),
    Namespace:  aws.String(namespace),
    Dimensions: []*cloudwatch.DimensionFilter{
        &cloudwatch.DimensionFilter{
            Name: aws.String(name),
        },
    },
})
fmt.Println("Metrics", result.Metrics)
```

See the [complete example](#) on GitHub.

Submit Custom Metrics

Choose **Copy** to save the code locally.

Create the file `custom_metrics.go`.

Import the packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/cloudwatch"

    "fmt"
)
```

Get the metric name, namespace, and dimensions from the command line.

```
if len(os.Args) != 4 {
    fmt.Println("You must supply a metric name, dimension, value, and namespace")
    os.Exit(1)
}

metric := os.Args[1]
dimensions := os.Args[3]
namespace := os.Args[2]
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create new cloudwatch client.
svc := cloudwatch.New(sess)
```

Call `PutMetricData`, supplying the metric name, unit, value, and dimensions. Print any errors, or a success message.

```
result, err := svc.PutMetricData(&cloudwatch.PutMetricDataInput{
    MetricData: []*cloudwatch.MetricDatum{
        &cloudwatch.MetricDatum{
            MetricName: aws.String(metric),
            Unit:      aws.String(cloudwatch.StandardUnitNone),
            Value:     aws.Float64(1.0),
            Dimensions: []*cloudwatch.Dimension{
                &cloudwatch.Dimension{
                    Name:  aws.String(dimension),
                    Value: aws.String(value),
                },
            },
        },
    },
    Namespace: aws.String(namespace),
})
fmt.Println("Success", result)
```

See the [complete example](#) on GitHub.

Sending Events to Amazon CloudWatch Events

These Go examples show you how to use the AWS SDK for Go to:

- Create and update a rule used to trigger an event
- Define one or more targets to respond to an event
- Send events that are matched to targets for handling

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

CloudWatch Events delivers a near real-time stream of system events that describe changes in AWS resources to any of various targets. Using simple rules, you can match events and route them to one or more target functions or streams.

In these examples, Go code is used to send events to CloudWatch Events. The code uses the AWS SDK for Go to manage instances by using these methods of the [CloudWatchEvents](#) type:

- [PutRule](#)
- [PutTargets](#)

- [PutEvents](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with CloudWatch Events. To learn more, see [Adding Events with PutEvents](#) in the *Amazon CloudWatch Events User Guide*.

Tasks Before You Start

To set up and run this example, you must first complete these tasks:

1. Create a Lambda function using the hello-world blueprint to serve as the target for events. To learn how, see [Step 1: Create an AWS Lambda function](#) in the CloudWatch Events User Guide.
2. Create an IAM role whose policy grants permission to CloudWatch Events and that includes `events.amazonaws.com` as a trusted entity. For more information about creating an IAM role, see [Creating a Role to Delegate Permissions to an AWS Service](#) in the *IAM User Guide*.

Use the following role policy when creating the IAM role.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "CloudWatchEventsFullAccess",
      "Effect": "Allow",
      "Action": "events:*",
      "Resource": "*"
    },
    {
      "Sid": "IAMPassRoleForCloudWatchEvents",
      "Effect": "Allow",
      "Action": "iam:PassRole",
      "Resource": "arn:aws:iam::*:role/AWS_Events_Invoke_Targets"
    }
  ]
}
```

Use the following trust relationship when creating the IAM role.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "events.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

Create a Scheduled Rule

Choose **Copy** to save the code locally.

Create the file `events_schedule_rule.go`. Import the packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudwatchevents"  
  
    "fmt"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch Events client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
// Create the cloudwatch events client  
svc := cloudwatchevents.New(sess)
```

Call `PutRule`, supplying a name, `DEMO_EVENT`, ARN of the IAM role you created, `IAM_ROLE_ARN`, and an expression defining the schedule. Finally, display the ARN of the rule.

```
result, err := svc.PutRule(&cloudwatchevents.PutRuleInput{  
    Name:            aws.String("DEMO_EVENT"),  
    RoleArn:         aws.String("IAM_ROLE_ARN"),  
    ScheduleExpression: aws.String("rate(5 minutes)"),  
})  
fmt.Println("Rule ARN:", result.RuleArn)
```

See the [complete example](#) on GitHub.

Add a Lambda Function Target

Choose **Copy** to save the code locally.

Create the file `events_put_targets.go`. Import the packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/cloudwatchevents"  
  
    "fmt"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a new CloudWatch Events client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
// Create the cloudwatch events client  
svc := cloudwatchevents.New(sess)
```

Call `PutTargets`, supplying a name for the rule, `DEMO_EVENT`. For the target, specify the ARN of the Lambda function you created, `LAMBDA_FUNCTION_ARN`, and the ID of the rule, `myCloudWatchEventsTarget`. Print any errors, or a success message with any targets that failed.

```
result, err := svc.PutTargets(&cloudwatchevents.PutTargetsInput{
    Rule: aws.String("DEMO_EVENT"),
    Targets: []*cloudwatchevents.Target{
        &cloudwatchevents.Target{
            Arn: aws.String("LAMBDA_FUNCTION_ARN"),
            Id:  aws.String("myCloudWatchEventsTarget"),
        },
    },
})
fmt.Println("Success", result)
```

See the [complete example](#) on GitHub.

Send Events

Choose **Copy** to save the code locally.

Create the file `events_put_events.go`. Import the packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/cloudwatchevents"

    "fmt"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create a CloudWatch Events client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create the cloudwatch events client
svc := cloudwatchevents.New(sess)
```

Call `PutEvents`, supplying key-name value pairs in the `Details` field, and specifying the ARN of the Lambda function you created, `RESOURCE_ARN`. See [PutEventsRequestEntry](#) for a description of the fields. Finally, display the ingested events.

```
result, err := svc.PutEvents(&cloudwatchevents.PutEventsInput{
    Entries: []*cloudwatchevents.PutEventsRequestEntry{
        &cloudwatchevents.PutEventsRequestEntry{
            Detail:      aws.String("{ \"key1\": \"value1\", \"key2\": \"value2\" }"),
            DetailType: aws.String("appRequestSubmitted"),
            Resources: []*string{
                aws.String("RESOURCE_ARN"),
            },
            Source: aws.String("com.company.myapp"),
        },
    },
})
fmt.Println("Ingested events:", result.Entries)
```

See the [complete example](#) on GitHub.

AWS CodeBuild Examples Using the AWS SDK for Go

AWS CodeBuild is a fully managed build service that compiles source code, runs tests, and produces software packages that are ready to deploy. The AWS SDK for Go examples can integrate AWS CodeBuild into your applications. The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Getting Information about All AWS CodeBuild Projects \(p. 43\)](#)
- [Building an AWS CodeBuild Project \(p. 44\)](#)
- [Listing Your AWS CodeBuild Project Builds \(p. 44\)](#)

Getting Information about All AWS CodeBuild Projects

The following example lists the names of up to 100 of your AWS CodeBuild projects.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/codebuild"
    "fmt"
    "os"
)

// Lists a CodeBuild projects in the region configured in the shared config
func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create CodeBuild service client
    svc := codebuild.New(sess)

    // Get the list of projects
    result, err := svc.ListProjects(
        &codebuild.ListProjectsInput{
            SortBy:      aws.String("NAME"),
            SortOrder:  aws.String("ASCENDING", )})

    if err != nil {
        fmt.Println("Got error listing projects: ", err)
        os.Exit(1)
    }
}
```



```
    for _, p := range result.Projects {
        fmt.Println(*p)
    }
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

Building an AWS CodeBuild Project

The following example builds the AWS CodeBuild project specified on the command line. If no command-line argument is supplied, it emits an error and quits.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/codebuild"
    "fmt"
    "os"
)

// Builds a CodeBuild project in the region configured in the shared config
func main() {
    // Requires one argument, the name of the project.
    if len(os.Args) != 2 {
        fmt.Println("Project name required!")
        os.Exit(1)
    }

    project := os.Args[1]

    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create CodeBuild service client
    svc := codebuild.New(sess)

    // Build the project
    _, err := svc.StartBuild(&codebuild.StartBuildInput{ProjectName: aws.String(project)})

    if err != nil {
        fmt.Println("Got error building project: ", err)
        os.Exit(1)
    }

    fmt.Printf("Started build for project %q\n", project)
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

Listing Your AWS CodeBuild Project Builds

The following example displays information about your AWS CodeBuild project builds, including the name of the project, when the build started, and how long each phase of the build took, in seconds.

```
package main
```

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/codebuild"
    "fmt"
    "os"
)

// Lists the CodeBuild builds for all projects in the region configured in the shared
// config
func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create CodeBuild service client
    svc := codebuild.New(sess)

    // Get the list of builds
    names, err := svc.ListBuilds(&codebuild.ListBuildsInput{SortOrder:
aws.String("ASCENDING")})

    if err != nil {
        fmt.Println("Got error listing builds: ", err)
        os.Exit(1)
    }

    // Get information about each build
    builds, err := svc.BatchGetBuilds(&codebuild.BatchGetBuildsInput{Ids: names.Ids})

    if err != nil {
        fmt.Println("Got error getting builds: ", err)
        os.Exit(1)
    }

    for _, build := range builds.Builds {
        fmt.Printf("Project: %s\n", aws.StringValue(build.ProjectName))
        fmt.Printf("Phase:   %s\n", aws.StringValue(build.CurrentPhase))
        fmt.Printf("Status:  %s\n", aws.StringValue(build.BuildStatus))
        fmt.Println("")
    }
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

Amazon DynamoDB Examples Using the AWS SDK for Go

Amazon DynamoDB is a fully managed NoSQL database service that provides fast and predictable performance with seamless scalability. The AWS SDK for Go examples can integrate Amazon DynamoDB into your Go applications. The examples assume you have already set up and configured the SDK (that is, you have imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

The topic also provides a link to a downloadable version of DynamoDB, which includes an interactive web interface so you can experiment with DynamoDB offline.

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Listing all Amazon DynamoDB Tables Using the AWS SDK for Go \(p. 46\)](#)
- [Creating an Amazon DynamoDB Table Using the AWS SDK for Go \(p. 47\)](#)
- [Creating an Amazon DynamoDB Table Item Using the AWS SDK for Go \(p. 48\)](#)
- [Creating Amazon DynamoDB Table Items from a JSON File Using the AWS SDK for Go \(p. 49\)](#)
- [Reading an Amazon DynamoDB Table Item \(p. 51\)](#)
- [Getting Amazon DynamoDB Table Items Using Expression Builder \(p. 52\)](#)
- [Updating an Amazon DynamoDB Table Item \(p. 54\)](#)
- [Deleting an Amazon DynamoDB Table Item \(p. 55\)](#)

Listing all Amazon DynamoDB Tables Using the AWS SDK for Go

The following example uses the `DynamoDBListTables` operation to list all tables for the `us-west-2` region.

Create the file `dynamodb_list_tables.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "fmt"  
    "os"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/dynamodb"  
)
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*</problematic>` `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create DynamoDB client  
svc := dynamodb.New(sess)
```

Call `ListTables`. If an error occurs, print the error and exit. If no error occurs, loop through the tables, printing the name of each table.

```
result, err := svc.ListTables(&dynamodb.ListTablesInput{})  
  
if err != nil {  
    fmt.Println(err)  
    os.Exit(1)  
}  
  
fmt.Println("Tables:")  
fmt.Println("")
```

```
for _, n := range result.TableNames {  
    fmt.Println(*n)  
}
```

See the [complete example](#) on GitHub.

Creating an Amazon DynamoDB Table Using the AWS SDK for Go

The following example uses the DynamoDB [CreateTable](#) operation to create the table **Music** in the `us-west-2` region.

Create the file `dynamodb_create_table.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "fmt"  
    "os"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/dynamodb"  
)
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create DynamoDB client  
svc := dynamodb.New(sess)
```

Call **CreateTable**. If an error occurs, print the error and exit. If no error occurs, print a message that the table was created.

```
input := &dynamodb.CreateTableInput{  
    AttributeDefinitions: []*dynamodb.AttributeDefinition{  
        {  
            AttributeName: aws.String("year"),  
            AttributeType: aws.String("N"),  
        },  
        {  
            AttributeName: aws.String("title"),  
            AttributeType: aws.String("S"),  
        },  
    },  
    KeySchema: []*dynamodb.KeySchemaElement{  
        {  
            AttributeName: aws.String("year"),  
            KeyType:      aws.String("HASH"),  
        },  
        {  
            AttributeName: aws.String("title"),  
            KeyType:      aws.String("RANGE"),  
        },  
    },  
}
```

```
        ProvisionedThroughput: &dynamodb.ProvisionedThroughput{
            ReadCapacityUnits:  aws.Int64(10),
            WriteCapacityUnits: aws.Int64(10),
        },
        TableName: aws.String("Movies"),
    }

    _, err = svc.CreateTable(input)

    if err != nil {
        fmt.Println("Got error calling CreateTable:")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    fmt.Println("Created the table Movies in us-west-2")
```

See the [complete example](#) on GitHub.

Creating an Amazon DynamoDB Table Item Using the AWS SDK for Go

The following example uses the DynamoDB [PutItem](#) operation to create the table item with the year **2015** and title **The Big New Movie** in the *Movies* table in the *us-west-2* region.

Create the file *dynamodb_create_item.go*. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"
    "github.com/aws/aws-sdk-go/service/dynamodb/dynamodbattribute"
```

Create the data structures we use to containing the information about the table item.

```
type ItemInfo struct {
    Plot string`json:"plot"`
    Rating float64`json:"rating"`
}

type Item struct {
    Year int`json:"year"`
    Title string`json:"title"`
    Info ItemInfo`json:"info"`
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file **<problematic>*</problematic>** `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create DynamoDB client
```

```
svc := dynamodb.New(sess)
```

Initialize the structs with the movie data and marshall that data into a map of **AttributeValue** objects.

```
info := ItemInfo{
    Plot: "Nothing happens at all.",
    Rating: 0.0,
}

item := Item{
    Year: 2015,
    Title: "The Big New Movie",
    Info: info,
}

av, err := dynamodbattribute.MarshalMap(item)
```

Create the input for **CreateTable** and call it. If an error occurs, print the error and exit. If no error occurs, print a message that the item was added to the table.

```
input := &dynamodb.PutItemInput{
    Item: av,
    TableName: aws.String("Movies"),
}

_, err = svc.PutItem(input)

if err != nil {
    fmt.Println("Got error calling PutItem:")
    fmt.Println(err.Error())
    os.Exit(1)
}

fmt.Println("Successfully added 'The Big New Movie' (2015) to Movies table")
```

See the [complete example](#) on GitHub.

Creating Amazon DynamoDB Table Items from a JSON File Using the AWS SDK for Go

The following example uses the DynamoDB [PutItem](#) operation in a loop to create the items defined in *movie_data.json* file in the Movies table in the us-west-2 region.

Create the file *dynamodb_load_items.go*. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "encoding/json"
    "fmt"
    "io/ioutil"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"
    "github.com/aws/aws-sdk-go/service/dynamodb/dynamodbattribute"
)
```

Create the data structures we use to containing the information about the table item.

```
type ItemInfo struct {
    Plot string`json:"plot"`
    Rating float64`json:"rating"`
}

type Item struct {
    Year int`json:"year"`
    Title string`json:"title"`
    Info ItemInfo`json:"info"`
}
```

Create a function to get the table items from the JSON file.

```
func getItems() []Item {
    raw, err := ioutil.ReadFile("./movie_data.json")

    if err != nil {
        fmt.Println(err.Error())
        os.Exit(1)
    }

    var items []Item
    json.Unmarshal(raw, &items)
    return items
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

if err != nil {
    fmt.Println("Error creating session:")
    fmt.Println(err.Error())
    os.Exit(1)
}

// Create DynamoDB client
svc := dynamodb.New(sess)
```

Call `getItems` to get the items. Loop through each item, marshall that data into a map of **AttributeValue** objects, add the item to the `Movies` table, and print out the title and year of the movie added to the table.

```
items := getItems()

// Add each item to Movies table:
for _, item := range items {
    av, err := dynamoattribute.MarshalMap(item)

    if err != nil {
        fmt.Println("Got error marshalling map:")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    // Create item in table Movies
    input := &dynamodb.PutItemInput{
```

```
        Item: av,
        TableName: aws.String("Movies"),
    }

    _, err = svc.PutItem(input)

    if err != nil {
        fmt.Println("Got error calling PutItem:")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    fmt.Println("Successfully added '", item.Title, "' (" , item.Year, ") to Movies table")
}
```

See the [complete example](#) and a [sample JSON file](#) on GitHub.

Reading an Amazon DynamoDB Table Item

The following example uses the DynamoDB [GetItem](#) operation to retrieve information about the item with the year **2015** and title **The Big New Movie** in the `movies` table in the `us-west-2` region.

Create the file `dynamodb_read_item.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"
    "github.com/aws/aws-sdk-go/service/dynamodb/dynamodbattribute"
```

Create the data structures we use to contain the information about the table item.

```
type ItemInfo struct {
    Plot string `json:"plot"`
    Rating float64 `json:"rating"`
}

type Item struct {
    Year int `json:"year"`
    Title string `json:"title"`
    Info ItemInfo `json:"info"`
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*</problematic>` `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create DynamoDB client
svc := dynamodb.New(sess)
```

Call **GetItem** to add the item to the table. If we encounter an error, print the error message. Otherwise, display information about the item.


```
result, err := svc.GetItem(&dynamodb.GetItemInput{
    TableName: aws.String("Movies"),
    Key: map[string]*dynamodb.AttributeValue{
        "year": {
            N: aws.String("2015"),
        },
        "title": {
            S: aws.String("The Big New Movie"),
        },
    },
})

if err != nil {
    fmt.Println(err.Error())
    return
}

item := Item{}

err = dynamodbattribute.UnmarshalMap(result.Item, &item)

if err != nil {
    panic(fmt.Sprintf("Failed to unmarshal Record, %v", err))
}

if item.Title == "" {
    fmt.Println("Could not find 'The Big New Movie' (2015)")
    return
}

fmt.Println("Found item:")
fmt.Println("Year: ", item.Year)
fmt.Println("Title: ", item.Title)
fmt.Println("Plot: ", item.Info.Plot)
fmt.Println("Rating:", item.Info.Rating)
```

See the [complete example](#) on GitHub.

Getting Amazon DynamoDB Table Items Using Expression Builder

The following example uses the [DynamoDBScan](#) operation to get items with a rating greater than **8.0** in the year **2011** in the **Movies** table in the **us-west-2** region.

The example uses the [Expression Builder](#) package released in version 1.11.0 of the AWS SDK for Go in September 2017.

Create the file *dynamodb_scan_item.go*. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"
    "github.com/aws/aws-sdk-go/service/dynamodb/dynamodbattribute"
    "github.com/aws/aws-sdk-go/service/dynamodb/expression"
)
```

Create the data structures we use to contain the information about the table item.

```
type ItemInfo struct {
    Plot string`json:"plot"`
    Rating float64`json:"rating"`
}

type Item struct {
    Year int`json:"year"`
    Title string`json:"title"`
    Info ItemInfo`json:"info"`
}
```

Create variables for the minimum rating and year for the table items to retrieve.

```
min_rating := 8.0
year := 2011
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2"),
})

if err != nil {
    fmt.Println("Got error creating session:")
    fmt.Println(err.Error())
    os.Exit(1)
}

// Create DynamoDB client
svc := dynamodb.New(sess)
```

Create the expression defining the year for which we filter the table items to retrieve, and the projection so we get the title, year, and rating for each retrieved item. Then build the expression.

```
filt := expression.Name("year").Equal(expression.Value(year))

// Or we could get by ratings and pull out those with the right year later
// filt := expression.Name("info.rating").GreaterThan(expression.Value(min_rating))

// Get back the title, year, and rating
proj := expression.NamesList(expression.Name("title"), expression.Name("year"),
    expression.Name("info.rating"))

expr, err := expression.NewBuilder().WithFilter(filt).WithProjection(proj).Build()
```

Create the inputs for and call **Scan** to retrieve the items from the table (the movies made in 2011).

```
params := &dynamodb.ScanInput{
    ExpressionAttributeNames: expr.Names(),
    ExpressionAttributeValues: expr.Values(),
    FilterExpression:         expr.Filter(),
    ProjectionExpression:     expr.Projection(),
    TableName:               aws.String("Movies"),
}

// Make the DynamoDB Query API call
```

```
result, err := svc.Scan(params)
```

Loop through the movies from 2011 and display the title and rating for those where the rating is greater than 8.0

```
for _, i := range result.Items {
    item := Item{}

    err = dynamodbattribute.UnmarshalMap(i, &item)

    if err != nil {
        fmt.Println("Got error unmarshalling:")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    // Which ones had a higher rating?
    if item.Info.Rating > min_rating {
        // Or it we had filtered by rating previously:
        // if item.Year == year {
        num_items += 1

        fmt.Println("Title: ", item.Title)
        fmt.Println("Rating:", item.Info.Rating)
        fmt.Println()
    }
}
```

See the [complete example](#) on GitHub.

Updating an Amazon DynamoDB Table Item

The following example uses the DynamoDB [UpdateItem](#) operation to update the rating to **0.5** for the item with the year **2015** and title **The Big New Movie** in the **Movies** table in the **us-west-2** region.

Create the file *dynamodb_update_item.go*. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"
)
```

Initialize the session that the SDK uses to load credentials from the shared credentials file **<problematic>*** and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create DynamoDB client
svc := dynamodb.New(sess)
```

Call **UpdateItem** to add the item to the table. If we encounter an error, print the error message. Otherwise, display a message that the item was updated.

```
input := &dynamodb.UpdateItemInput{
    ExpressionAttributeValues: map[string]*dynamodb.AttributeValue{
        "r": {
            N: aws.String("0.5"),
        },
    },
    TableName: aws.String("Movies"),
    Key: map[string]*dynamodb.AttributeValue{
        "year": {
            N: aws.String("2015"),
        },
        "title": {
            S: aws.String("The Big New Movie"),
        },
    },
    ReturnValues:      aws.String("UPDATED_NEW"),
    UpdateExpression: aws.String("set info.rating = :r"),
}

_, err = svc.UpdateItem(input)

if err != nil {
    fmt.Println(err.Error())
    return
}

fmt.Println("Successfully updated 'The Big New Movie' (2015) rating to 0.5")
```

See the [complete example](#) on GitHub.

Deleting an Amazon DynamoDB Table Item

The following example uses the `DynamoDBUpdateItem` operation to delete the item with the year **2015** and title **The Big New Movie** from the `Movies` table in the `us-west-2` region.

Create the file `dynamodb_update_item.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/dynamodb"
)
```

Initialize the session that the SDK uses to credentials from the shared credentials file `<problematic>*</problematic>` `~/.aws/credentials`, and create a new DynamoDB service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create DynamoDB client
svc := dynamodb.New(sess)
```

Call `DeleteItem` to delete the item from the table. If we encounter an error, print the error message. Otherwise, display a message that the item was deleted.

```
input := &dynamodb.DeleteItemInput{
    Key: map[string]*dynamodb.AttributeValue{
        "year": {
            N: aws.String("2015"),
        },
        "title": {
            S: aws.String("The Big New Movie"),
        },
    },
    TableName: aws.String("Movies"),
}

_, err = svc.DeleteItem(input)

if err != nil {
    fmt.Println("Got error calling DeleteItem")
    fmt.Println(err.Error())
    return
}

fmt.Println("Deleted 'The Big New Movie' (2015)")
```

See the [complete example](#) on GitHub.

Amazon EC2 Examples Using the AWS SDK for Go

Amazon Elastic Compute Cloud (Amazon EC2) is a web service that provides resizable computing capacity—literally servers in Amazon's data centers—that you use to build and host your software systems. The AWS SDK for Go examples can integrate Amazon EC2 into your Go applications. The examples assume you have already set up and configured the SDK (that is, you have imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go](#) (p. 2) and [Configuring the AWS SDK for Go](#) (p. 4).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Creating Amazon EC2 Instances with Tags or without Block Devices](#) (p. 56)
- [Managing Amazon EC2 Instances](#) (p. 58)
- [Working with Amazon EC2 Key Pairs](#) (p. 63)
- [Using Regions and Availability Zones with Amazon EC2](#) (p. 66)
- [Working with Security Groups in Amazon EC2](#) (p. 68)
- [Using Elastic IP Addresses in Amazon EC2](#) (p. 73)

Creating Amazon EC2 Instances with Tags or without Block Devices

This Go example shows you how to:

- Create an Amazon EC2 instance with tags or set up an instance without a block device

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenarios

In these examples, you use a series of Go routines to create Amazon EC2 instances with tags or set up an instance without a block device.

The routines use the AWS SDK for Go to perform these tasks by using these methods of the [EC2](#) type:

- [BlockDeviceMapping](#)
- [RunInstances](#)
- [CreateTags](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.

Create an Instance with Tags

The Amazon EC2 service has an operation for creating instances ([RunInstances](#)) and another for attaching tags to instances ([CreateTags](#)). To create an instance with tags, call both of these operations in succession. The following example creates an instance and then adds a Name tag to it. The Amazon EC2 console displays the value of the Name tag in its list of instances.

```
svc := ec2.New(session.New(&aws.Config{Region: aws.String("us-west-2")}))
// Specify the details of the instance that you want to create.
runResult, err := svc.RunInstances(&ec2.RunInstancesInput{
    // An Amazon Linux AMI ID for t2.micro instances in the us-west-2 region
    ImageId:      aws.String("ami-e7527ed7"),
    InstanceType: aws.String("t2.micro"),
    MinCount:     aws.Int64(1),
    MaxCount:     aws.Int64(1),
})

if err != nil {
    log.Println("Could not create instance", err)
    return
}

log.Println("Created instance", *runResult.Instances[0].InstanceId)

// Add tags to the created instance
_, errtag := svc.CreateTags(&ec2.CreateTagsInput{
    Resources: []*string{runResult.Instances[0].InstanceId},
    Tags: []*ec2.Tag{
        {
            Key:   aws.String("Name"),
            Value: aws.String("MyFirstInstance"),
        },
    },
})

if errtag != nil {
    log.Println("Could not create tags for instance",
runResult.Instances[0].InstanceId, errtag)
    return
}

log.Println("Successfully tagged instance")
```

You can add up to 10 tags to an instance in a single `CreateTags` operation.

Create an Image without a Block Device

Sometimes when you create an Amazon EC2 image, you might want to explicitly exclude certain block devices. To do this, you can use the `NoDevice` parameter in [BlockDeviceMapping](#). When this parameter is set to an empty string `"`, the named device isn't mapped.

The `NoDevice` parameter is compatible only with `DeviceName`, not with any other field in `BlockDeviceMapping`. The request will fail if other parameters are present.

```
func main() {
    svc := ec2.New(session.New())
    opts := &ec2.CreateImageInput{
        Description: aws.String("image description"),
        InstanceId:  aws.String("i-abcdef12"),
        Name:        aws.String("image name"),
        BlockDeviceMappings: []*ec2.BlockDeviceMapping{
            &ec2.BlockDeviceMapping{
                DeviceName: aws.String("/dev/sda1"),
                NoDevice:    aws.String(""),
            },
            &ec2.BlockDeviceMapping{
                DeviceName: aws.String("/dev/sdb"),
                NoDevice:    aws.String(""),
            },
            &ec2.BlockDeviceMapping{
                DeviceName: aws.String("/dev/sdc"),
                NoDevice:    aws.String(""),
            },
        },
    }
    resp, err := svc.CreateImage(opts)
    if err != nil {
        fmt.Println(err)
        return
    }

    fmt.Println("success", resp)
}
```

Managing Amazon EC2 Instances

These Go examples show you how to:

- Describe Amazon EC2 instances
- Manage Amazon EC2 instance monitoring
- Start and stop Amazon EC2 instances
- Reboot Amazon EC2 instances

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In these examples, you use a series of Go routines to perform several basic instance management operations.

The routines use the AWS SDK for Go to perform the operations by using these methods of the Amazon EC2 client class:

- [DescribeInstances](#)
- [MonitorInstances](#)
- [UnmonitorInstances](#)
- [StartInstances](#)
- [StopInstances](#)
- [RebootInstances](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with the lifecycle of Amazon EC2 instances. To learn more, see [Instance Lifecycle](#) in the *Amazon EC2 User Guide for Linux Instances*.

Describe Your Instances

Create a new Go file named `describing_instances.go`.

The Amazon EC2 service has an operation for describing instances, [DescribeInstances](#).

Import the required AWS SDK for Go packages.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

Use the following code to create a session and Amazon EC2 client.

```
func main() {
    // Load session from shared config
    sess := session.Must(session.NewSessionWithOptions(session.Options{
        SharedConfigState: session.SharedConfigEnable,
    }))

    // Create new EC2 client
    ec2Svc := ec2.New(sess)
```

Call `DescribeInstances` to get detailed information for each instance.

```
    result, err := ec2Svc.DescribeInstances(nil)
    if err != nil {
        fmt.Println("Error", err)
    } else {
        fmt.Println("Success", result)
    }
}
```

Manage Instance Monitoring

Create a new Go file named `monitoring_instances.go`.

Import the required AWS SDK for Go packages.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

To access Amazon EC2, create an EC2 client.

```
func main() {
    // Load session from shared config
    sess := session.Must(session.NewSessionWithOptions(session.Options{
        SharedConfigState: session.SharedConfigEnable,
    }))

    // Create new EC2 client
    svc := ec2.New(sess)
```

Based on the value of a command-line argument (ON or OFF), call either the `MonitorInstances` method of the Amazon EC2 service object to begin detailed monitoring of the specified instances, or the `UnmonitorInstances` method. Before you try to change the monitoring of these instances, use the `DryRun` parameter to test whether you have permission to change instance monitoring.

```
if os.Args[1] == "ON" {
    input := &ec2.MonitorInstancesInput{
        InstanceIds: []*string{
            aws.String(os.Args[2]),
        },
        DryRun: aws.Bool(true),
    }
    result, err := svc.MonitorInstances(input)
    awsErr, ok := err.(awserr.Error)

    if ok && awsErr.Code() == "DryRunOperation" {
        input.DryRun = aws.Bool(false)
        result, err = svc.MonitorInstances(input)
        if err != nil {
            fmt.Println("Error", err)
        } else {
            fmt.Println("Success", result.InstanceMonitorings)
        }
    } else {
        fmt.Println("Error", err)
    }
} else if os.Args[1] == "OFF" { // Turn monitoring off
    input := &ec2.UnmonitorInstancesInput{
        InstanceIds: []*string{
            aws.String(os.Args[2]),
        },
        DryRun: aws.Bool(true),
    }
    result, err := svc.UnmonitorInstances(input)
    awsErr, ok := err.(awserr.Error)
    if ok && awsErr.Code() == "DryRunOperation" {
        input.DryRun = aws.Bool(false)
```

```
        result, err = svc.UnmonitorInstances(input)
        if err != nil {
            fmt.Println("Error", err)
        } else {
            fmt.Println("Success", result.InstanceMonitorings)
        }
    } else {
        fmt.Println("Error", err)
    }
}

}
```

Start and Stop Instances

Create a new Go file named `start_stop_instances.go`.

Import the required AWS SDK for Go packages.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awsserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

To access Amazon EC2, create an EC2 client. The user will pass in a state value of START or STOP and the instance ID.

```
func main() {
    // Load session from shared config
    sess := session.Must(session.NewSessionWithOptions(session.Options{
        SharedConfigState: session.SharedConfigEnable,
    }))

    // Create new EC2 client
    svc := ec2.New(sess)
```

Based on the value of a command-line argument (START or STOP), call either the `StartInstances` method of the Amazon EC2 service object to start the specified instances, or the `StopInstances` method to stop them. Before you try to start or stop the selected instances, use the `DryRun` parameter to test whether you have permission to start or stop them.

```
    if os.Args[1] == "START" {
        input := &ec2.StartInstancesInput{
            InstanceIds: []*string{
                aws.String(os.Args[2]),
            },
            DryRun: aws.Bool(true),
        }
        result, err := svc.StartInstances(input)
        awsErr, ok := err.(awsserr.Error)

        if ok && awsErr.Code() == "DryRunOperation" {
            // Let's now set dry run to be false. This will allow us to start the instances
            input.DryRun = aws.Bool(false)
```

```
        result, err = svc.StartInstances(input)
        if err != nil {
            fmt.Println("Error", err)
        } else {
            fmt.Println("Success", result.StartingInstances)
        }
    } else { // This could be due to a lack of permissions
        fmt.Println("Error", err)
    }
} else if os.Args[1] == "STOP" { // Turn instances off
    input := &ec2.StopInstancesInput{
        InstanceIds: []*string{
            aws.String(os.Args[2]),
        },
        DryRun: aws.Bool(true),
    }
    result, err := svc.StopInstances(input)
    awsErr, ok := err.(awserr.Error)
    if ok && awsErr.Code() == "DryRunOperation" {
        input.DryRun = aws.Bool(false)
        result, err = svc.StopInstances(input)
        if err != nil {
            fmt.Println("Error", err)
        } else {
            fmt.Println("Success", result.StoppingInstances)
        }
    } else {
        fmt.Println("Error", err)
    }
}
}
```

Reboot Instances

Create a new Go file named `reboot_instances.go`.

Import the required AWS SDK for Go packages.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

To access Amazon EC2, create an EC2 client. The user will pass in a state value of START or STOP and the instance ID.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create new EC2 client
svc := ec2.New(sess)
```

Based on the value of a command-line argument (START or STOP), call either the `StartInstances` method of the Amazon EC2 service object to start the specified instances, or the `StopInstances`

method to stop them. Before you try to reboot the selected instances, use the `DryRun` parameter to test whether the instance exists and that you have permission to reboot it.

```
input := &ec2.RebootInstancesInput{
    InstanceIds: []*string{
        aws.String(os.Args[1]),
    },
    DryRun: aws.Bool(true),
}
result, err := svc.RebootInstances(input)
awsErr, ok := err.(awserr.Error)
```

If the error code is `DryRunOperation`, it means that you do have the permissions you need to reboot the instance.

```
if ok && awsErr.Code() == "DryRunOperation" {
    input.DryRun = aws.Bool(false)
    result, err = svc.RebootInstances(input)
    if err != nil {
        fmt.Println("Error", err)
    } else {
        fmt.Println("Success", result)
    }
} else { // This could be due to a lack of permissions
    fmt.Println("Error", err)
}
}
```

Working with Amazon EC2 Key Pairs

These Go examples show you how to:

- Describe an Amazon EC2 key pair
- Create an Amazon EC2 key pair
- Delete an Amazon EC2 key pair

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

Amazon EC2 uses public-key cryptography to encrypt and decrypt login information. Public-key cryptography uses a public key to encrypt data, then the recipient uses the private key to decrypt the data. The public and private keys are known as a key pair.

The routines use the AWS SDK for Go to perform these tasks by using these methods of the `EC2` type:

- [CreateKeyPair](#)
- [DeleteKeyPair](#)
- [DescribeKeyPairs](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the SDK.

- You are familiar with Amazon EC2 key pairs. To learn more, see [Amazon EC2 Key Pairs](#) in the *Amazon EC2 User Guide for Linux Instances*.

Describe Your Key Pairs

Create a new Go file named `ec2_describe_keypairs.go`.

Import the required AWS SDK for Go packages.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

Use the following code to create a session and Amazon EC2 client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create an EC2 service client.
    svc := ec2.New(sess)
```

Call `DescribeKeyPairs` to get a list of key pairs and print them out.

```
    result, err := svc.DescribeKeyPairs(nil)
    if err != nil {
        exitErrorf("Unable to get key pairs, %v", err)
    }

    fmt.Println("Key Pairs:")
    for _, pair := range result.KeyPairs {
        fmt.Printf("%s: %s\n", *pair.KeyName, *pair.KeyFingerprint)
    }
}
```

The routine uses the following utility function.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Create a Key Pair

Create a new Go file named `ec2_create_keypair.go`.

Import the required AWS SDK for Go packages.

```
package main

import (
```

```
"fmt"
"os"
"path/filepath"

"github.com/aws/aws-sdk-go/aws"
"github.com/aws/aws-sdk-go/aws/awsserr"
"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/ec2"
)
```

Get the key pair name passed in to the code and, to access Amazon EC2, create an EC2 client.

```
func main() {
    if len(os.Args) != 2 {
        exitErrorf("pair name required\nUsage: %s key_pair_name",
            filepath.Base(os.Args[0]))
    }
    pairName := os.Args[1]
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create an EC2 service client.
    svc := ec2.New(sess)
```

Create a new key pair with the provided name.

```
    result, err := svc.CreateKeyPair(&ec2.CreateKeyPairInput{
        KeyName: aws.String(pairName),
    })
    if err != nil {
        if aerr, ok := err.(awsserr.Error); ok && aerr.Code() ==
            "InvalidKeyPair.Duplicate" {
            exitErrorf("Keypair %q already exists.", pairName)
        }
        exitErrorf("Unable to create key pair: %s, %v.", pairName, err)
    }

    fmt.Printf("Created key pair %q %s\n%s\n",
        *result.KeyName, *result.KeyFingerprint,
        *result.KeyMaterial)
}
```

The routine uses the following utility function.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Delete a Key Pair

Create a new Go file named `ec2_delete_keypair.go`.

Import the required AWS SDK for Go packages.

```
package main

import (
    "fmt"
    "os"
```

```
"path/filepath"

"github.com/aws/aws-sdk-go/aws"
"github.com/aws/aws-sdk-go/aws/awserr"
"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/ec2"

)
```

Get the key pair name passed in to the code and, to access Amazon EC2, create an EC2 client.

```
func main() {
    if len(os.Args) != 2 {
        exitErrorf("pair name required\nUsage: %s key_pair_name",
            filepath.Base(os.Args[0]))
    }
    pairName := os.Args[1]

    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create an EC2 service client.
    svc := ec2.New(sess)
```

Delete the key pair with the provided name.

```
_, err := svc.DeleteKeyPair(&ec2.DeleteKeyPairInput{
    KeyName: aws.String(pairName),
})
if err != nil {
    if aerr, ok := err.(awserr.Error); ok && aerr.Code() ==
        "InvalidKeyPair.Duplicate" {
        exitErrorf("Key pair %q does not exist.", pairName)
    }
    exitErrorf("Unable to delete key pair: %s, %v.", pairName, err)
}

fmt.Printf("Successfully deleted %q key pair\n", pairName)
}
```

The routine uses the following utility function.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Using Regions and Availability Zones with Amazon EC2

These Go examples show you how to retrieve details about AWS Regions and Availability Zones.

An Amazon EC2 security group acts as a virtual firewall that controls the traffic for one or more instances. You add rules to each security group to allow traffic to or from its associated instances. You can modify the rules for a security group at any time; the new rules are automatically applied to all instances that are associated with the security group.

The code in this example uses the AWS SDK for Go to perform these tasks by using these methods of the Amazon EC2 client class:

- [DescribeSecurityGroups](#)
- [AuthorizeSecurityGroupIngress](#)
- [CreateSecurityGroup](#)
- [DescribeVpcs](#)
- [DeleteSecurityGroup](#)

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

Amazon EC2 is hosted in multiple locations worldwide. These locations are composed of AWS Regions and Availability Zones. Each region is a separate geographic area with multiple, isolated locations known as Availability Zones. Amazon EC2 provides the ability to place instances and data in these multiple locations.

In this example, you use Go code to retrieve details about regions and Availability Zones. The code uses the AWS SDK for Go to manage instances by using the following methods of the Amazon EC2 client class:

- [DescribeAvailabilityZones](#)
- [DescribeRegions](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with AWS Regions and Availability Zones. To learn more, see [Regions and Availability Zones](#) in the *Amazon EC2 User Guide for Linux Instances* or [Regions and Availability Zones](#) in the *Amazon EC2 User Guide for Windows Instances*.

List the Groups

This example describes the security groups by IDs that are passed in to the routine. It takes a space separated list of group IDs as input.

To get started, create a new Go file named `regions_and_availability.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

In the `main` function, create a session with credentials from the shared credentials file, `~/.aws/credentials`, and create a new EC2 client.

```
func main() {
    // Load session from shared config
    sess := session.Must(session.NewSessionWithOptions(session.Options{
        SharedConfigState: session.SharedConfigEnable,
```



```
    })))  
  
    // Create new EC2 client  
    svc := ec2.New(sess)
```

Print out the list of regions that work with Amazon EC2 that are returned by calling `DescribeRegions`.

```
resultRegions, err := svc.DescribeRegions(nil)  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}
```

Add a call that retrieves Availability Zones only for the region of the EC2 service object.

```
resultAvalZones, err := svc.DescribeAvailabilityZones(nil)  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}  
  
fmt.Println("Success", resultAvalZones.AvailabilityZones)  
}
```

Working with Security Groups in Amazon EC2

These Go examples show you how to:

- Retrieve information about your security groups
- Create a security group to access an Amazon EC2 instance
- Delete an existing security group

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

An Amazon EC2 security group acts as a virtual firewall that controls the traffic for one or more instances. You add rules to each security group to allow traffic to or from its associated instances. You can modify the rules for a security group at any time; the new rules are automatically applied to all instances that are associated with the security group.

The code in this example uses the AWS SDK for Go to perform these tasks by using these methods of the Amazon EC2 client class:

- [DescribeSecurityGroups](#)
- [AuthorizeSecurityGroupIngress](#)
- [CreateSecurityGroup](#)
- [DescribeVpcs](#)
- [DeleteSecurityGroup](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.

- You are familiar with Amazon EC2 security groups. To learn more, see [Amazon EC2 Amazon Security Groups for Linux Instances](#) in the *Amazon EC2 User Guide for Linux Instances* or [Amazon EC2 Amazon Security Groups for Windows Instances](#) in the *Amazon EC2 User Guide for Windows Instances*.

Describing Your Security Groups

This example describes the security groups by IDs that are passed into the routine. It takes a space separated list of group IDs as input.

To get started, create a new Go file named `ec2_describe_security_groups.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/cloudwatch"

    "fmt"
    "os"
)
```

In the main function, get the security group ID that is passed in.

```
func main() {
    if len(os.Args) < 2 {
        exitErrorf("Security Group ID required\nUsage: %s group_id ...",
            filepath.Base(os.Args[0]))
    }
    groupIds := os.Args[1:]
}
```

Initialize a session and create an EC2 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create an EC2 service client.
svc := ec2.New(sess)
```

Obtain and print out the security group descriptions. You will explicitly check for errors caused by an invalid group ID.

```
result, err := svc.DescribeSecurityGroups(&ec2.DescribeSecurityGroupsInput{
    GroupIds: aws.StringSlice(groupIds),
})
if err != nil {
    if aerr, ok := err.(awserr.Error); ok {
        switch aerr.Code() {
            case "InvalidGroupId.Malformed":
                fallthrough
            case "InvalidGroup.NotFound":
                exitErrorf("%s.", aerr.Message())
        }
    }
    exitErrorf("Unable to get descriptions for security groups, %v", err)
}
```

```
    fmt.Println("Security Group:")
    for _, group := range result.SecurityGroups {
        fmt.Println(group)
    }
}
```

The following utility function is used by this example.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Creating a Security Group

You can create new Amazon EC2 security groups. To do this, you use the [CreateSecurityGroup](#) method.

This example creates a new security group with the given name and description for access to open ports 80 and 22. If a VPC ID is not provided, it associates the security group with the first VPC in the account.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "flag"
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awsserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

Get the parameters (name, description, and optional ID of the VPC) that are passed in to the routine.

```
func main() {
    var name, desc, vpcID string
    flag.StringVar(&name, "n", "", "Group Name")
    flag.StringVar(&desc, "d", "", "Group Description")
    flag.StringVar(&vpcID, "vpc", "", "(Optional) VPC ID to associate security group with")
    flag.Parse()

    if len(name) == 0 || len(desc) == 0 {
        flag.PrintDefaults()
        exitErrorf("Group name and description require")
    }
}
```

Create a session.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create an EC2 service client.
svc := ec2.New(sess)
```

If the VPC ID was not provided, you have to retrieve the first one in the account.

```
if len(vpcID) == 0 {
    // Get a list of VPCs so we can associate the group with the first VPC.
    result, err := svc.DescribeVpcs(nil)
    if err != nil {
        exitErrorrf("Unable to describe VPCs, %v", err)
    }
    if len(result.Vpcs) == 0 {
        exitErrorrf("No VPCs found to associate security group with.")
    }
    vpcID = aws.StringValue(result.Vpcs[0].VpcId)
}
```

Then create the security group with the VPC ID, name, and description.

```
createRes, err := svc.CreateSecurityGroup(&ec2.CreateSecurityGroupInput{
    GroupName:    aws.String(name),
    Description:  aws.String(desc),
    VpcId:        aws.String(vpcID),
})
if err != nil {
    if aerr, ok := err.(awserr.Error); ok {
        switch aerr.Code() {
            case "InvalidVpcID.NotFound":
                exitErrorrf("Unable to find VPC with ID %q.", vpcID)
            case "InvalidGroup.Duplicate":
                exitErrorrf("Security group %q already exists.", name)
        }
    }
    exitErrorrf("Unable to create security group %q, %v", name, err)
}
fmt.Printf("Created security group %s with VPC %s.\n",
    aws.StringValue(createRes.GroupId), vpcID)
```

Add permissions to the security group.

```
_, err = svc.AuthorizeSecurityGroupIngress(&ec2.AuthorizeSecurityGroupIngressInput{
    GroupName:    aws.String(name),
    IpPermissions: []*ec2.IpPermission{
        (&ec2.IpPermission{}).
            SetIpProtocol("tcp").
            SetFromPort(80).
            SetToPort(80).
            SetIpRanges([]*ec2.IpRange{
                {CidrIp: aws.String("0.0.0.0/0")},
            }),
        (&ec2.IpPermission{}).
            SetIpProtocol("tcp").
            SetFromPort(22).
            SetToPort(22).
            SetIpRanges([]*ec2.IpRange{
                (&ec2.IpRange{}).
                    SetCidrIp("0.0.0.0/0"),
            }),
    },
})
if err != nil {
    exitErrorrf("Unable to set security group %q ingress, %v", name, err)
}

fmt.Println("Successfully set security group ingress")
}
```

The following utility function is used by this example.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Deleting a Security Group

You can delete an Amazon EC2 security group in code. To do this, you use the [DeleteSecurityGroup](#) method.

This example deletes a security group with the given group ID.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main  
  
import (  
    "fmt"  
    "os"  
    "path/filepath"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/awsserr"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/ec2"  
)
```

Get the group ID that is passed in to the routine.

```
func main() {  
    if len(os.Args) != 2 {  
        exitErrorf("Security Group ID required\nUsage: %s group_id",  
            filepath.Base(os.Args[0]))  
    }  
    groupID := os.Args[1]
```

Create a session.

```
svc := ec2.New(sess)
```

Then delete the security group with the group ID that is passed in.

```
_, err := svc.DeleteSecurityGroup(&ec2.DeleteSecurityGroupInput{  
    GroupId: aws.String(groupID),  
})  
if err != nil {  
    if aerr, ok := err.(awsserr.Error); ok {  
        switch aerr.Code() {  
            case "InvalidGroupId.Malformed":  
                fallthrough  
            case "InvalidGroup.NotFound":  
                exitErrorf("%s.", aerr.Message())  
            }  
        }  
        exitErrorf("Unable to get descriptions for security groups, %v.", err)  
    }  
  
    fmt.Printf("Successfully delete security group %q.\n", groupID)
```

```
}
```

This example uses the following utility function.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Using Elastic IP Addresses in Amazon EC2

These Go examples show you how to:

- Describe Amazon EC2 instance IP addresses
- Allocate addresses to Amazon EC2 instances
- Release Amazon EC2 instance IP addresses

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

An Elastic IP address is a static IP address designed for dynamic cloud computing that is associated with your AWS account. It is a public IP address, reachable from the Internet. If your instance doesn't have a public IP address, you can associate an Elastic IP address with the instance to enable communication with the Internet.

In this example, you use a series of Go routines to perform several Amazon EC2 operations involving Elastic IP addresses. The routines use the AWS SDK for Go to manage Elastic IP addresses by using these methods of the Amazon EC2 client class:

- [DescribeAddresses](#)
- [AllocateAddress](#)
- [AssociateAddress](#)
- [ReleaseAddress](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with Elastic IP addresses in Amazon EC2. To learn more, see [Elastic IP Addresses](#) in the *Amazon EC2 User Guide for Linux Instances* or [Elastic IP Addresses](#) in the *Amazon EC2 User Guide for Windows Instances*.

Describe Instance IP Addresses

Create a new Go file named `ec2_describe_addresses.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main  
  
import (  
    "fmt"
```

```
"os"

"github.com/aws/aws-sdk-go/aws"
"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/ec2"
)
```

Get the Address Descriptions

This routine prints out the Elastic IP Addresses for the account's VPC. Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`, and create a new EC2 service client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create an EC2 service client.
    svc := ec2.New(sess)
```

Make the API request to EC2 filtering for the addresses in the account's VPC.

```
    result, err := svc.DescribeAddresses(&ec2.DescribeAddressesInput{
        Filters: []*ec2.Filter{
            {
                Name:    aws.String("domain"),
                Values: aws.StringSlice([]string{"vpc"}),
            },
        },
    })
    if err != nil {
        exitErrorf("Unable to elastic IP address, %v", err)
    }

    // Printout the IP addresses if there are any.
    if len(result.Addresses) == 0 {
        fmt.Printf("No elastic IPs for %s region\n", *svc.Config.Region)
    } else {
        fmt.Println("Elastic IPs")
        for _, addr := range result.Addresses {
            fmt.Println(" ", fmtAddress(addr))
        }
    }
}
```

The `fmtAddress` and `exitErrorf` functions are utility functions used in the example.

```
func fmtAddress(addr *ec2.Address) string {
    out := fmt.Sprintf("IP: %s, allocation id: %s",
        aws.StringValue(addr.PublicIp), aws.StringValue(addr.AllocationId))
    if addr.InstanceId != nil {
        out += fmt.Sprintf(", instance-id: %s", *addr.InstanceId)
    }
    return out
}

func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Allocate Addresses to Instances

Create a new Go file named `ec2_allocate_address.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
    "os"
    "path/filepath"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ec2"
)
```

This routine attempts to allocate a VPC Elastic IP address for the current region. The IP address requires and will be associated with the instance ID that is passed in.

```
func main() {
    if len(os.Args) != 2 {
        exitErrorf("instance ID required\nUsage: %s instance_id",
            filepath.Base(os.Args[0]))
    }
    instanceID := os.Args[1]
```

You will need to initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`, and create a new Amazon S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create an EC2 service client.
svc := ec2.New(sess)
```

Call [AllocateAddress](#), passing in "vpc" as the Domain value.

```
allocRes, err := svc.AllocateAddress(&ec2.AllocateAddressInput{
    Domain: aws.String("vpc"),
})
if err != nil {
    exitErrorf("Unable to allocate IP address, %v", err)
}
```

Call [AssociateAddress](#) to associate the new Elastic IP address with an existing Amazon EC2 instance, and print out the results.

```
assocRes, err := svc.AssociateAddress(&ec2.AssociateAddressInput{
    AllocationId: allocRes.AllocationId,
    InstanceId:   aws.String(instanceID),
})
if err != nil {
    exitErrorf("Unable to associate IP address with %s, %v",
        instanceID, err)
}
```



```
    fmt.Printf("Successfully allocated %s with instance %s.\n\tallocation id: %s,\n\tassociation id: %s\n",\n        *allocRes.PublicIp, instanceID, *allocRes.AllocationId, *assocRes.AssociationId)\n}
```

This example also uses the `exitErrorf` utility function.

```
func exitErrorf(msg string, args ...interface{}) {\n    fmt.Fprintf(os.Stderr, msg+"\n", args...)\n    os.Exit(1)\n}
```

Release Instance IP Addresses

This routine releases an Elastic IP address allocation ID. If the address is associated with an Amazon EC2 instance, the association is removed.

Create a new Go file named `ec2_release_address.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main\n\nimport (\n    "fmt"\n    "os"\n    "path/filepath"\n\n    "github.com/aws/aws-sdk-go/aws"\n    "github.com/aws/aws-sdk-go/aws/awsserr"\n    "github.com/aws/aws-sdk-go/aws/session"\n    "github.com/aws/aws-sdk-go/service/ec2"\n)
```

The routine requires that the user pass in the allocation ID of the Elastic IP address.

```
func main() {\n    if len(os.Args) != 2 {\n        exitErrorf("allocation ID required\nUsage: %s allocation_id",\n            filepath.Base(os.Args[0]))\n    }\n    allocationID := os.Args[1]\n}
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`, and create a new EC2 service client.

```
sess, err := session.NewSession(&aws.Config{\n    Region: aws.String("us-west-2")},\n)\n\n// Create an EC2 service client.\nsvc := ec2.New(sess)
```

Attempt to release the Elastic IP address by using the allocation ID.

```
_, err := svc.ReleaseAddress(&ec2.ReleaseAddressInput{\n    AllocationId: aws.String(allocationID),\n})
```

```
    if err != nil {
        if aerr, ok := err.(awserr.Error); ok && aerr.Code() ==
            "InvalidAllocationID.NotFound" {
            exitErrorf("Allocation ID %s does not exist", allocationID)
        }
        exitErrorf("Unable to release IP address for allocation %s, %v",
            allocationID, err)
    }

    fmt.Printf("Successfully released allocation ID %s\n", allocationID)
}
```

This example uses the `fmtAddress` and `exitErrorf` utility functions.

```
func fmtAddress(addr *ec2.Address) string {
    out := fmt.Sprintf("IP: %s, allocation id: %s",
        aws.StringValue(addr.PublicIp), aws.StringValue(addr.AllocationId))
    if addr.InstanceId != nil {
        out += fmt.Sprintf(", instance-id: %s", *addr.InstanceId)
    }
    return out
}

func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Amazon Glacier Examples Using the AWS SDK for Go

Amazon Glacier is a secure, durable, and extremely low-cost cloud storage service for data archiving and long-term backup. The AWS SDK for Go examples can integrate Amazon Glacier into your applications. The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

The Scenario

Amazon Glacier is a secure cloud storage service for data archiving and long-term backup. The service is optimized for infrequently accessed data where a retrieval time of several hours is suitable. These examples show you how to create a vault and upload an archive with Go. The methods used include:

- [CreateVault](#)
- [UploadArchive](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with the Amazon Glacier data model. To learn more, see [Amazon Glacier Data Model](#) in the *Amazon Glacier Developer Guide*.

Create a Vault

The following example uses the Amazon Glacier [CreateVault](#) operation to create a vault named `YOUR_VAULT_NAME`.

```
svc := glacier.New(session.New(&aws.Config{Region: aws.String("us-west-2")})))
_, err := svc.CreateVault(&glacier.CreateVaultInput{
    VaultName: aws.String("YOUR_VAULT_NAME"),
})
if err != nil {
    log.Println(err)
    return
}
log.Println("Created vault!")
```

Upload an Archive

The following example assumes you have a vault named `YOUR_VAULT_NAME`. It uses the Amazon Glacier [UploadArchive](#) operation to upload a single reader object as an entire archive. The AWS SDK for Go automatically computes the tree hash checksum for the data to be uploaded.

```
vaultName := "YOUR_VAULT_NAME"

svc := glacier.New(session.New(&aws.Config{Region: aws.String("us-west-2")})))
result, err := svc.UploadArchive(&glacier.UploadArchiveInput{
    AccountId: aws.String("-"),
    VaultName: &vaultName,
    Body:      bytes.NewReader(make([]byte, 2*1024*1024)), // 2 MB buffer
})
if err != nil {
    log.Println("Error uploading archive.", err)
    return
}

log.Println("Uploaded to archive", *result.ArchiveId)
```

IAM Examples Using the AWS SDK for Go

AWS Identity and Access Management (IAM) is a web service that enables AWS customers to manage users and user permissions in AWS. The service is targeted at organizations with multiple users or systems in the cloud that use AWS products. With IAM, you can centrally manage users, security credentials such as access keys, and permissions that control which AWS resources users can access.

The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go](#) (p. 2) and [Configuring the AWS SDK for Go](#) (p. 4).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Managing IAM Users](#) (p. 79)
- [Managing IAM Access Keys](#) (p. 85)
- [Managing IAM Account Aliases](#) (p. 89)

- [Working with IAM Policies \(p. 92\)](#)
- [Working with IAM Server Certificates \(p. 98\)](#)

Managing IAM Users

This Go example shows you how to create, update, view, and delete IAM users. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In this example, you use a series of Go routines to manage users in IAM. The routines use the AWS SDK for Go IAM client methods that follow:

- [CreateUser](#)
- [ListUsers](#)
- [UpdateUser](#)
- [GetUser](#)
- [DeleteUser](#)
- [GetAccountAuthorizationDetails](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with IAM users. To learn more, see [IAM Users](#) in the *IAM User Guide*.

Create a New IAM User

This code creates a new IAM user.

Create a new Go file named `iam_createuser.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws/awserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

The code takes the new user name as an argument, and then calls `GetUser` with the user name.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)

    _, err := svc.GetUser(&iam.GetUserInput{
```

```
        UserName: &os.Args[1],
    })
```

If you receive a `NoSuchEntity` error, call `CreateUser` because the user doesn't exist.

```
    if awserr, ok := err.(awserr.Error); ok && awserr.Code() ==
iam.ErrCodeNoSuchEntityException {
        result, err := svc.CreateUser(&iam.CreateUserInput{
            UserName: &os.Args[1],
        })

        if err != nil {
            fmt.Println("CreateUser Error", err)
            return
        }

        fmt.Println("Success", result)
    } else {
        fmt.Println("GetUser Error", err)
    }
}
```

List IAM Users in Your Account

You can get a list of the users in your account and print the list to the console.

Create a new Go file named `iam_listusers.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `ListUsers` and print the results.

```
    result, err := svc.ListUsers(&iam.ListUsersInput{
        MaxItems: aws.Int64(10),
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }
```

```
    for i, user := range result.Users {
        if user == nil {
            continue
        }
        fmt.Printf("%d user %s created %v\n", i, *user.UserName, user.CreateDate)
    }
}
```

Update a User's Name

In this example, you change the name of an IAM user to a new value.

Create a new Go file named `iam_updateuser.go`. You must import the relevant Go and `aws-sdk-go` packages by adding the following lines.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `UpdateUser`, passing in the original user name and the new name, and print the results.

```
    result, err := svc.UpdateUser(&iam.UpdateUserInput{
        UserName:    &os.Args[1],
        NewUserName: &os.Args[2],
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Success", result)
}
```

Delete an IAM User

In this example, you delete an IAM user.

Create a new Go file named `iam_deleteuser.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
```

```
"fmt"
"os"

"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/iam"

)
```

Set up a new IAM client.

```
// Initialize a session in us-west-2 that the SDK will use to load
    Region: aws.String("us-west-2")},
)

// Create a IAM service client.
svc := iam.New(sess)

result, err := svc.UpdateUser(&iam.UpdateUserInput{
    UserName:    &os.Args[1],
    NewUserName: &os.Args[2],
})
```

Call `UpdateUser`, passing in the user name, and print the results. If the user doesn't exist, log an error.

```
    fmt.Println("Error", err)
    return
}

fmt.Println("Success", result)
}
```

List the IAM Users who have Administrator Privileges

In this example, you list the IAM users who have administrator privileges (a policy or attached policy of the user or a group to which the user belongs has the name **AdministratorAccess**).

Create a new Go file named `IamGetAdmins.go`. Import the following packages.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"

    "fmt"
    "os"
)
```

Create a method to determine whether a user has a policy that has administrator privileges.

```
func UserPolicyHasAdmin(user *iam.UserDetail, admin string) bool {
    for _, policy := range user.UserPolicyList {
        if *policy.PolicyName == admin {
            return true
        }
    }

    return false
}
```

Create a method to determine whether a user has an attached policy that has administrator privileges.

```
func AttachedUserPolicyHasAdmin(user *iam.UserDetail, admin string) bool {
    for _, policy := range user.AttachedManagedPolicies {
        if *policy.PolicyName == admin {
            return true
        }
    }

    return false
}
```

Create a method that determines whether a group has a policy that has administrator privileges.

```
func GroupPolicyHasAdmin(svc *iam.IAM, group *iam.Group, admin string) bool {
    input := &iam.ListGroupPoliciesInput{
        GroupName: group.GroupName,
    }

    result, err := svc.ListGroupPolicies(input)
    if err != nil {
        fmt.Println("Got error calling ListGroupPolicies for group", group.GroupName)
    }

    // Wade through policies
    for _, policyName := range result.PolicyNames {
        if
            *policyName == admin {
                return true
            }
    }

    return false
}
```

Create a method that determines whether a group has an attached policy that has administrator privileges.

```
func AttachedGroupPolicyHasAdmin(svc *iam.IAM, group *iam.Group, admin string) bool {
    input := &iam.ListAttachedGroupPoliciesInput{GroupName: group.GroupName}
    result, err := svc.ListAttachedGroupPolicies(input)
    if err != nil {
        fmt.Println("Got error getting attached group policies:")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    for _, policy := range result.AttachedPolicies {
        if *policy.PolicyName == admin {
            return true
        }
    }

    return false
}
```

Create a method that determines whether any group to which the user belongs has administrator privileges.

```
func UsersGroupsHaveAdmin(svc *iam.IAM, user *iam.UserDetail, admin string) bool {
    input := &iam.ListGroupsForUserInput{UserName: user.UserName}
    result, err := svc.ListGroupsForUser(input)
    if err != nil {
```



```
        fmt.Println("Got error getting groups for user:")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    for _, group := range result.Groups {
        groupPolicyHasAdmin := GroupPolicyHasAdmin(svc, group, admin)

        if groupPolicyHasAdmin {
            return true
        }

        attachedGroupPolicyHasAdmin := AttachedGroupPolicyHasAdmin(svc, group, admin)

        if attachedGroupPolicyHasAdmin {
            return true
        }
    }

    return false
}
```

Create a method that determines whether a user has administrator privileges.

```
func IsUserAdmin(svc *iam.IAM, user *iam.UserDetail, admin string) bool {
    // Check policy, attached policy, and groups (policy and attached policy)
    policyHasAdmin := UserPolicyHasAdmin(user, admin)
    if policyHasAdmin {
        return true
    }

    attachedPolicyHasAdmin := AttachedUserPolicyHasAdmin(user, admin)
    if attachedPolicyHasAdmin {
        return true
    }

    userGroupsHaveAdmin := UsersGroupsHaveAdmin(svc, user, admin)
    if userGroupsHaveAdmin {
        return true
    }

    return false
}
```

Create a main method with an IAM client in us-west-2. Create variables to keep track of how many users we have and how many of those have administrator privileges.

```
func main() {
    sess, err := session.NewSession()
    if err != nil {
        fmt.Println("Got error creating new session")
        fmt.Println(err.Error())
        os.Exit(1)
    }

    svc := iam.New(sess, &aws.Config{Region: aws.String("us-west-2")})

    numUsers := 0
    numAdmins := 0
}
```

Create the input for and call `GetAccountAuthorizationDetails`. If there is an error, print an error message and quit.

```
user := "User"
input := &iam.GetAccountAuthorizationDetailsInput{Filter: []*string{&user}}
resp, err := svc.GetAccountAuthorizationDetails(input)
if err != nil {
    fmt.Println("Got error getting account details")
    fmt.Println(err.Error())
    os.Exit(1)
}
```

Loop through the users. If a user has administrator privileges, print their name and increment the number of users who have administrator privileges.

```
adminName := "AdministratorAccess"

// Wade through resulting users
for _, user := range resp.UserDetailList {
    numUsers += 1

    isAdmin := IsUserAdmin(svc, user, adminName)

    if isAdmin {
        fmt.Println(*user.UserName)
        numAdmins += 1
    }
}
```

If we did not get all of the users in the first call to `GetAccountAuthorizationDetails`, loop through the next set of users and determine which of those have administrator privileges.

```
// Wade through resulting users
for _, user := range resp.UserDetailList {
    numUsers += 1

    isAdmin := IsUserAdmin(svc, user, adminName)

    if isAdmin {
        fmt.Println(*user.UserName)
        numAdmins += 1
    }
}
```

Finally, display the number of users who have administrator access.

```
fmt.Println("")
fmt.Println("Found", numAdmins, "admin(s) out of", numUsers, "user(s).")
```

See the [complete example](#) on GitHub.

Managing IAM Access Keys

This Go example shows you how to create, modify, view, or rotate IAM access keys. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

Users need their own access keys to make programmatic calls to the AWS SDK for Go. To fill this need, you can create, modify, view, or rotate access keys (access key IDs and secret access keys) for IAM users.

By default, when you create an access key its status is Active, which means the user can use the access key for API calls.

In this example, you use a series of Go routines to manage access keys in IAM. The routines use the AWS SDK for Go IAM client methods that follow:

- [CreateAccessKey](#)
- [ListAccessKeys](#)
- [GetAccessKeyLastUsed](#)
- [UpdateAccessKey](#)
- [DeleteAccessKey](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with IAM access keys. To learn more, see [Managing Access Keys for IAM Users](#) in the *IAM User Guide*.

Create a New IAM Access Key

This code creates a new IAM access key for the IAM user named IAM_USER_NAME.

Create a new Go file named `iam_createaccesskey.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up the session.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `CreateAccessKey` and print the results.

```
    result, err := svc.CreateAccessKey(&iam.CreateAccessKeyInput{
        UserName: aws.String("IAM_USER_NAME"),
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }
```

```
    fmt.Println("Success", *result.AccessKey)
}
```

List a User's Access Keys

In this example, you get a list of the access keys for a user and print the list to the console.

Create a new Go file named `iam_listaccesskeys.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `ListAccessKeys` and print the results.

```
    result, err := svc.ListAccessKeys(&iam.ListAccessKeysInput{
        MaxItems: aws.Int64(5),
        UserName: aws.String("IAM_USER_NAME"),
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Success", result)
}
```

Get the Last Use for an Access Key

In this example, you find out when an access key was last used.

Create a new Go file named `iam_accesskeylastused.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
```

```
"github.com/aws/aws-sdk-go/service/iam"  
)
```

Set up a new IAM client.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a IAM service client.  
    svc := iam.New(sess)
```

Call `GetAccessKeyLastUsed`, passing in the access key ID, and print the results.

```
    result, err := svc.GetAccessKeyLastUsed(&iam.GetAccessKeyLastUsedInput{  
        AccessKeyId: aws.String("ACCESS_KEY_ID"),  
    })  
  
    if err != nil {  
        fmt.Println("Error", err)  
        return  
    }  
  
    fmt.Println("Success", *result.AccessKeyLastUsed)  
}
```

Update Access Key Status

In this example, you delete an IAM user.

Create a new Go file with the name `iam_updateaccesskey.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main  
  
import (  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/iam"  
)
```

Set up a new IAM client.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a IAM service client.  
    svc := iam.New(sess)
```

Call `UpdateAccessKey`, passing in the access key ID, status (making it active in this case), and user name.

```
    _, err := svc.UpdateAccessKey(&iam.UpdateAccessKeyInput{
```

```
        AccessKeyId: aws.String("ACCESS_KEY_ID"),
        Status:      aws.String(iam.StatusTypeActive),
        UserName:    aws.String("USER_NAME"),
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Access Key updated")
}
```

Delete an Access Key

In this example, you delete an access key.

Create a new Go file named `iam_deleteaccesskey.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `DeleteAccessKey`, passing in the access key ID and user name.

```
    result, err := svc.DeleteAccessKey(&iam.DeleteAccessKeyInput{
        AccessKeyId: aws.String("ACCESS_KEY_ID"),
        UserName:    aws.String("USER_NAME"),
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Success", result)
}
```

Managing IAM Account Aliases

This Go example shows you how to create, list, and delete IAM account aliases. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

You can use a series of Go routines to manage aliases in IAM. The routines use the AWS SDK for Go IAM client methods that follow:

- [CreateAccountAlias](#)
- [ListAccountAliases](#)
- [DeleteAccountAlias](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with IAM account aliases. To learn more, see [Your AWS Account ID and Its Alias](#) in the *IAM User Guide*.

Create a New IAM Account Alias

This code creates a new IAM user.

Create a new Go file named `iam_createaccountalias.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a session and an IAM client.

```
// Initialize a session in us-west-2 that the SDK will use to load
Region: aws.String("us-west-2")),
)

// Create a IAM service client.
svc := iam.New(sess)

_, err := svc.CreateAccountAlias(&iam.CreateAccountAliasInput{
    AccountAlias: &os.Args[1],
})
```

The code takes the new alias as an argument, and then calls `CreateAccountAlias` with the alias name.

```
_, err := svc.CreateAccountAlias(&iam.CreateAccountAliasInput{
    AccountAlias: &os.Args[1],
})

if err != nil {
    fmt.Println("Error", err)
    return
}
```

```
    fmt.Printf("Account alias %s has been created\n", os.Args[1])
}
```

List IAM Account Aliases

This code lists the aliases for your IAM account.

Create a new Go file named `iam_listaccountaliases.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a session and an IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

The code calls `ListAccountAliases`, specifying to return a maximum of 10 items.

```
    result, err := svc.ListAccountAliases(&iam.ListAccountAliasesInput{
        MaxItems: aws.Int64(10),
    })

    if err != nil {
        fmt.Println("Error", err)
        return
    }

    for i, alias := range result.AccountAliases {
        if alias == nil {
            continue
        }
        fmt.Printf("Alias %d: %s\n", i, *alias)
    }
}
```

Delete an IAM Account Alias

This code deletes a specified IAM account alias.

Create a new Go file with the name `iam_deleteaccountalias.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main
```



```
import (  
    "fmt"  
    "os"  
  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/iam"  
)
```

Set up a session and an IAM client.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a IAM service client.  
    svc := iam.New(sess)
```

The code calls `ListAccountAliases`, specifying to return a maximum of 10 items.

```
_, err := svc.DeleteAccountAlias(&iam.DeleteAccountAliasInput{  
    AccountAlias: &os.Args[1],  
})  
  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}  
  
fmt.Printf("Alias %s has been deleted\n", os.Args[1])  
}
```

Working with IAM Policies

This Go example shows you how to create, get, attach, and detach IAM policies. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

You grant permissions to a user by creating a policy, which is a document that lists the actions that a user can perform and the resources those actions can affect. Any actions or resources that are not explicitly allowed are denied by default. Policies can be created and attached to users, groups of users, roles assumed by users, and resources.

In this example, you use a series of Go routines to manage policies in IAM. The routines use the AWS SDK for Go IAM client methods that follow:

- [CreatePolicy](#)
- [GetPolicy](#)
- [ListAttachedRolePolicies](#)
- [AttachRolePolicy](#)
- [DetachRolePolicy](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.

- You are familiar with IAM policies. To learn more, see [Overview of Access Management: Permissions and Policies](#) in the *IAM User Guide*.

Create an IAM Policy

This code creates a new IAM Policy. Create a new Go file named `iam_createpolicy.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "encoding/json"
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Define two structs. The first is the definition of the policies to upload to IAM.

```
type PolicyDocument struct {
    Version   string
    Statement []StatementEntry
}
```

The second dictates what this policy will allow or disallow.

```
type StatementEntry struct {
    Effect   string
    Action   []string
    Resource string
}
```

Set up the session and IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Build the policy document using the structures defined earlier.

```
policy := PolicyDocument{
    Version: "2012-10-17",
    Statement: []StatementEntry{
        StatementEntry{
            Effect: "Allow",
            Action: []string{
                "logs:CreateLogGroup", // Allow for creating log groups
            },
            Resource: "RESOURCE ARN FOR logs:*",
        },
        StatementEntry{
```

```
        Effect: "Allow",
        // Allows for DeleteItem, GetItem, PutItem, Scan, and UpdateItem
        Action: []string{
            "dynamodb:DeleteItem",
            "dynamodb:GetItem",
            "dynamodb:PutItem",
            "dynamodb:Scan",
            "dynamodb:UpdateItem",
        },
        Resource: "RESOURCE ARN FOR dynamodb:*",
    },
},
}
```

Marshal the policy to JSON and pass to `CreatePolicy`.

```
b, err := json.Marshal(&policy)
if err != nil {
    fmt.Println("Error marshaling policy", err)
    return
}

result, err := svc.CreatePolicy(&iam.CreatePolicyInput{
    PolicyDocument: aws.String(string(b)),
    PolicyName:     aws.String("myDynamodbPolicy"),
})

if err != nil {
    fmt.Println("Error", err)
    return
}

fmt.Println("New policy", result)
}
```

Get an IAM Policy

In this example, you retrieve an existing policy from IAM. Create a new Go file named `iam_getpolicy.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `GetPolicy`, passing in the ARN for the policy (which is hard coded in this example), and print the results.

```
svc := iam.New(sess)

arn := "arn:aws:iam::aws:policy/AWSLambdaExecute"
result, err := svc.GetPolicy(&iam.GetPolicyInput{
    PolicyArn: &arn,
})

if err != nil {
    fmt.Println("Error", err)
    return
}

fmt.Printf("%s - %s\n", arn, *result.Policy.Description)
}
```

Attach a Managed Role Policy

In this example, you attach an IAM managed role policy. Create a new Go file named `iam_attachuserpolicy.go`. You'll call the `ListAttachedRolePolicies` method of the IAM service object, which returns an array of managed policies.

Then, you'll check the array members to see if the policy you want to attach to the role is already attached. If the policy isn't attached, you'll call the `AttachRolePolicy` method to attach it.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Declare variables to hold the name and ARN of the policy.

```
var pageErr error
policyName := "AmazonDynamoDBFullAccess"
policyArn := "arn:aws:iam::aws:policy/AmazonDynamoDBFullAccess"
```

Paginate through all the role policies. If your role exists on any role policy, you set the `pageErr` and return `false`, stopping the pagination.

```
err := svc.ListAttachedRolePoliciesPages(
```

```
    &iam.ListAttachedRolePoliciesInput{
        RoleName: &os.Args[1],
    },
    func(page *iam.ListAttachedRolePoliciesOutput, lastPage bool) bool {
        if page != nil && len(page.AttachedPolicies) > 0 {
            for _, policy := range page.AttachedPolicies {
                if *policy.PolicyName == policyName {
                    pageErr = fmt.Errorf("%s is already attached to this role",
policyName)
                    return false
                }
            }
            // We should keep paginating because we did not find our role
            return true
        }
        return false
    },
)
```

If your role policy is not attached already, call `AttachRolePolicy`.

```
    if pageErr != nil {
        fmt.Println("Error", pageErr)
        return
    }

    if err != nil {
        fmt.Println("Error", err)
        return
    }

    _, err = svc.AttachRolePolicy(&iam.AttachRolePolicyInput{
        PolicyArn: &policyArn,
        RoleName:  &os.Args[1],
    })

    if err != nil {
        fmt.Println("Unable to attach role policy to role")
        return
    }
    fmt.Println("Role attached successfully")
}
```

Detach a Managed Role Policy

In this example, you detach a role policy. Once again, you call the `ListAttachedRolePolicies` method of the IAM service object, which returns an array of managed policies.

Then, check the array members to see if the policy you want to detach from the role is attached. If the policy is attached, call the `DetachRolePolicy` method to detach it.

Create a new Go file named `iam_detachuserpolicy.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws/session"
```

```
"github.com/aws/aws-sdk-go/service/iam"  
)
```

Set up a new IAM client.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a IAM service client.  
    svc := iam.New(sess)
```

Declare variables to hold the name and ARN of the policy.

```
foundPolicy := false  
policyName := "AmazonDynamoDBFullAccess"  
policyArn := "arn:aws:iam::aws:policy/AmazonDynamoDBFullAccess"
```

Paginate through all the role policies. If the role exists on any role policy, you stop iterating and detach the role.

```
err := svc.ListAttachedRolePoliciesPages(  
    &iam.ListAttachedRolePoliciesInput{  
        RoleName: &os.Args[1],  
    },  
    func(page *iam.ListAttachedRolePoliciesOutput, lastPage bool) bool {  
        if page != nil && len(page.AttachedPolicies) > 0 {  
            for _, policy := range page.AttachedPolicies {  
                if *policy.PolicyName == policyName {  
                    foundPolicy = true  
                    return false  
                }  
            }  
            return true  
        }  
        return false  
    },  
)  
  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}  
  
if !foundPolicy {  
    fmt.Println("Policy was not attached to role")  
    return  
}  
  
_, err = svc.DetachRolePolicy(&iam.DetachRolePolicyInput{  
    PolicyArn: &policyArn,  
    RoleName: &os.Args[1],  
})  
  
if err != nil {  
    fmt.Println("Unable to detach role policy to role")  
    return  
}  
fmt.Println("Role detached successfully")  
}
```

Working with IAM Server Certificates

This Go example shows you how to carry out basic tasks for managing server certificate HTTPS connections with the AWS SDK for Go.

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

To enable HTTPS connections to your website or application on AWS, you need an SSL/TLS server certificate. To use a certificate that you obtained from an external provider with your website or application on AWS, you must upload the certificate to IAM or import it into AWS Certificate Manager.

In this example, you use a series of Go routines to manage policies in IAM. The routines use the AWS SDK for Go IAM client methods that follow:

- [ListServerCertificates](#)
- [GetServerCertificate](#)
- [UpdateServerCertificate](#)
- [DeleteServerCertificate](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with server certificates. To learn more, see [Working with Server Certificates](#) in the *IAM User Guide*.

List Your Server Certificates

This code lists your certificates.

Create a new Go file named `iam_listservercerts.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `ListServerCertificates` and print the details.

```
result, err := svc.ListServerCertificates(nil)
if err != nil {
    fmt.Println("Error", err)
    return
}

for i, metadata := range result.ServerCertificateMetadataList {
    if metadata == nil {
        continue
    }
```

```
        fmt.Printf("Metadata %d: %v\n", i, metadata)
    }
}
```

Get a Server Certificate

In this example, you retrieve an existing server certificate.

Create a new Go file named `iam_getservercert.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/iam"
)
```

Set up a new IAM client.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a IAM service client.
    svc := iam.New(sess)
```

Call `GetServerCertificate`, passing the name of the certificate, and print the results.

```
    result, err := svc.GetServerCertificate(&iam.GetServerCertificateInput{
        ServerCertificateName: aws.String("CERTIFICATE_NAME"),
    })
    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("ServerCertificate:", result)
}
```

Update a Server Certificate

In this example, you update an existing server certificate.

Create a new Go file named `iam_updateservercert.go`. You call the `UpdateServerCertificate` method of the IAM service object to change the name of the certificate.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
```



```
"github.com/aws/aws-sdk-go/aws"  
"github.com/aws/aws-sdk-go/aws/session"  
"github.com/aws/aws-sdk-go/service/iam"  
)
```

Set up a new IAM client.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a IAM service client.  
    svc := iam.New(sess)
```

Update the certificate name.

```
_, err := svc.UpdateServerCertificate(&iam.UpdateServerCertificateInput{  
    ServerCertificateName:    aws.String("CERTIFICATE_NAME"),  
    NewServerCertificateName: aws.String("NEW_CERTIFICATE_NAME"),  
})  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}  
  
fmt.Println("Server certificate updated")  
}
```

Delete a Server Certificate

In this example, you delete an existing server certificate.

Create a new Go file named `iam_deleteservercert.go`. You call the `DeleteServerCertificate` method of the IAM service object to change the name of the certificate.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main  
  
import (  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/iam"  
)
```

Set up a new IAM client.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a IAM service client.  
    svc := iam.New(sess)
```

Call the method to delete the certificate, specifying the name of certificate.

```
_, err := svc.DeleteServerCertificate(&iam.DeleteServerCertificateInput{
    ServerCertificateName: aws.String("CERTIFICATE_NAME"),
})
if err != nil {
    fmt.Println("Error", err)
    return
}

fmt.Println("Server certificate deleted")
}
```

AWS Key Management Service Examples Using the AWS SDK for Go

You can use the following examples to access AWS Key Management Service (AWS KMS) using the AWS SDK for Go. For more information about AWS KMS, see the [Amazon KMS documentation](#). For reference information about the AWS KMS client, see the [New](#) function.

Examples

Topics

- [Creating a CMK in AWS Key Management Service \(p. 101\)](#)
- [Encrypting Data with AWS Key Management Service \(p. 102\)](#)
- [Decrypting a Data Blob in AWS Key Management Service \(p. 103\)](#)
- [Re-encrypting a Data Blob in AWS Key Management Service \(p. 103\)](#)

Creating a CMK in AWS Key Management Service

The following example uses the AWS SDK for Go [CreateKey](#) method, which implements the [CreateKey](#) operation, to create a customer master key (CMK). Since the example only encrypts a small amount of data, a CMK is fine for our purposes. For larger amounts of data, use the CMK to encrypt a data encryption key (DEK).

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/kms"

    "fmt"
    "os"
)

// Create a customer master key (CMK)
// Since we are only encrypting small amounts of data (4 KiB or less) directly,
// a CMK is fine for our purposes.
// For larger amounts of data,
// use the CMK to encrypt a data encryption key (DEK).

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )
```

```
// Create KMS service client
svc := kms.New(sess)

// Create the key
result, err := svc.CreateKey(&kms.CreateKeyInput{
    Tags: []*kms.Tag{
        {
            TagKey:   aws.String("CreatedBy"),
            TagValue: aws.String("ExampleUser"),
        },
    },
})

if err != nil {
    fmt.Println("Got error creating key: ", err)
    os.Exit(1)
}

fmt.Println(*result.KeyMetadata.KeyId)
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

Encrypting Data with AWS Key Management Service

The following example uses the AWS SDK for Go [Encrypt](#) method, which implements the [Encrypt](#) operation, to encrypt the string "1234567890". The example displays a readable version of the resulting encrypted blob.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/kms"

    "fmt"
    "os"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create KMS service client
    svc := kms.New(sess)

    // Encrypt data key
    //
    // Replace the fictitious key ARN with a valid key ID

    keyId := "arn:aws:kms:us-west-2:111122223333:key/1234abcd-12ab-34cd-56ef-1234567890ab"

    text := "1234567890"

    // Encrypt the data key
    result, err := svc.Encrypt(&kms.EncryptInput{
        KeyId:   aws.String(keyId),
        Plaintext: []byte(text),
    })

    if err != nil {
```

```
        fmt.Println("Got error encrypting data: ", err)
        os.Exit(1)
    }

    fmt.Println("Blob (base-64 byte array):")
    fmt.Println(result.CiphertextBlob)
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

Decrypting a Data Blob in AWS Key Management Service

The following example uses the AWS SDK for Go [Decrypt](#) method, which implements the [Decrypt](#) operation, to decrypt the provided string and emits the result.

```
import (
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/kms"

    "fmt"
    "os"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create KMS service client
    svc := kms.New(sess)

    // Encrypted data
    blob := []byte{...}

    // Decrypt the data
    result, err := svc.Decrypt(&kms.DecryptInput{CiphertextBlob: blob})

    if err != nil {
        fmt.Println("Got error decrypting data: ", err)
        os.Exit(1)
    }

    blob_string := string(result.Plaintext)

    fmt.Println(blob_string)
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

Re-encrypting a Data Blob in AWS Key Management Service

The following example uses the AWS SDK for Go [ReEncrypt](#) method, which implements the [ReEncrypt](#) operation, to decrypt encrypted data and then immediately re-encrypt data under a new customer master key (CMK). The operations are performed entirely on the server side within AWS KMS, so they

never expose your plaintext outside of AWS KMS. The example displays a readable version of the resulting re-encrypted blob.

```
import (
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/kms"

    "fmt"
    "os"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create KMS service client
    svc := kms.New(sess)

    // Encrypt data key
    //
    // Replace the fictitious key ARN with a valid key ID

    keyId := "arn:aws:kms:us-west-2:111122223333:key/1234abcd-12ab-34cd-56ef-1234567890ab"

    // Encrypted data
    blob := []byte{...}

    // Re-encrypt the data key
    result, err := svc.ReEncrypt(&kms.ReEncryptInput{CiphertextBlob: blob,
        DestinationKeyId: &keyId})

    if err != nil {
        fmt.Println("Got error re-encrypting data: ", err)
        os.Exit(1)
    }

    fmt.Println("Blob (base-64 byte array):")
    fmt.Println(result.CiphertextBlob)
}
```

Choose Copy to save the code locally. See the [complete example](#) on GitHub.

AWS Lambda Examples Using the AWS SDK for Go

AWS Lambda (Lambda) is a zero-administration compute platform for backend web developers that runs your code for you in the AWS Cloud, and provides you with a fine-grained pricing structure. You can use the following examples to access AWS Lambda (Lambda) using the AWS SDK for Go. For more information about Lambda, see the [Lambda documentation](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Displaying Information about All Lambda Functions \(p. 105\)](#)
- [Creating a Lambda Function \(p. 105\)](#)
- [Running a Lambda Function \(p. 106\)](#)

- [Configuring a Lambda Function to Receive Notifications \(p. 108\)](#)

Displaying Information about All Lambda Functions

The following example displays the names and descriptions of the Lambda functions in the `us-west-2` region.

```
func main() {
    // Initialize a session
    sess := session.Must(session.NewSessionWithOptions(session.Options{
        SharedConfigState: session.SharedConfigEnable,
    }))

    // Create Lambda service client
    svc := lambda.New(sess, &aws.Config{Region: aws.String("us-west-2")})

    result, err := svc.ListFunctions(nil)

    if err != nil {
        fmt.Println("Cannot list functions")
        os.Exit(0)
    }

    fmt.Println("Functions:")

    for _, f := range result.Functions {
        fmt.Println("Name:      " + aws.StringValue(f.FunctionName))
        fmt.Println("Description: " + aws.StringValue(f.Description))
        fmt.Println("")
    }
}
```

See the [complete example](#) on GitHub.

Creating a Lambda Function

The following example creates the Lambda function `functionName` in the `us-west-2` region using the following values:

- Role ARN: `resourceArn`. In most cases, you need to attach only the `AWSLambdaExecute` managed policy to the policy for this role.
- Function entry point: `handler`
- Runtime: `runtime`
- Zip file: `zipFileName + .zip`
- Bucket: `bucketName`
- Key: `zipFileName`

The first step is to create the session and Lambda client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create Lambda service client
svc := lambda.New(sess, &aws.Config{Region: aws.String("us-west-2")})
```

```
contents, err := ioutil.ReadFile(zipFileName + ".zip")

if err != nil {
    fmt.Println("Could not read " + zipFileName + ".zip")
    os.Exit(0)
}
```

Next, create the structures for the input argument to the `CreateFunction` function.

```
createCode := &lambda.FunctionCode{
    S3Bucket:      aws.String(bucketName),
    S3Key:          aws.String(zipFileName),
    S3ObjectVersion: aws.String(""),
    ZipFile:        contents,
}

createArgs := &lambda.CreateFunctionInput{
    Code:          createCode,
    FunctionName:  aws.String(functionName),
    Handler:       aws.String(handler),
    Role:          aws.String(resourceArn),
    Runtime:       aws.String(runtime),
}
```

Finally, call `CreateFunction` and display a message with the result of the call.

```
result, err := svc.CreateFunction(createArgs)

if err != nil {
    fmt.Println("Cannot create function: " + err.Error())
} else {
    fmt.Println(result)
}
```

See the [complete example](#) on GitHub.

Running a Lambda Function

The following example runs the Lambda function `MyGetItemsFunction` in the `us-west-2` region. This Node.js function returns a list of items from a database. The input JSON looks like the following.

```
{
  "SortBy": "name|time",
  "SortOrder": "ascending|descending",
  "Number": 50
}
```

Where:

- `SortBy` is the criteria for sorting the results. Our example uses `time`, which means the returned items are sorted in the order in which they were added to the database.
- `SortOrder` is the order of sorting. Our example uses `descending`, which means the most-recent item is last in the list.
- `Number` is the maximum number of items to retrieve (the default is 50). Our example uses 10, which means get the 10 most-recent items.

The output JSON looks like the following when the function succeeds and two items are returned.

```
{
  "statusCode": 200,
  "body": {
    "result": "success",
    "error": ""
    "data": [
      {
        "item": "item1"
      },
      {
        "item": "item2"
      }
    ]
  }
}
```

Where:

- `statusCode`– An HTTP status code; 200 means the call was successful.
- `body`– The body of the returned JSON.
- `result`– The result of the call, either `success` or `failure`.
- `error`– An error message if `result` is `failure`; otherwise, an empty string.
- `data`– The returned results if `result` is `success`; otherwise, `nil`.
- `item`– An item from the list of results.

The first step is to import the packages we use.

```
import (
    "encoding/json"
    "fmt"
    "os"
    "strconv"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/lambda"
)
```

Then we create the Lambda client we use to invoke the Lambda function.

```
client := lambda.New(sess, &aws.Config{Region: aws.String("us-west-2")})
```

Next, we create the request argument and call `MyGetItemsFunction`. If there is an error, we display a message and quit.

```
// Get the 10 most recent items
request := getItemsRequest{"time", "descending", 10}

payload, err := json.Marshal(request)

if err != nil {
    fmt.Println("Error marshalling MyGetItemsFunction request")
    os.Exit(0)
}

result, err := client.Invoke(&lambda.InvokeInput{FunctionName:
    aws.String("MyGetItemsFunction"), Payload: payload})
```



```
if err != nil {
    fmt.Println("Error calling MyGetItemsFunction")
    os.Exit(0)
}
```

Finally, we parse the response, and if successful, we print out the items.

```
var resp getItemsResponse

err = json.Unmarshal(result.Payload, &resp)

if err != nil {
    fmt.Println("Error unmarshalling MyGetItemsFunction response")
    os.Exit(0)
}

// If the status code is NOT 200, the call failed
if resp.StatusCode != 200 {
    fmt.Println("Error getting items, StatusCode: " + strconv.Itoa(resp.StatusCode))
    os.Exit(0)
}

// If the result is failure, we got an error
if resp.Body.Result == "failure" {
    fmt.Println("Failed to get items")
    os.Exit(0)
}

// Print out items
if len(resp.Body.Data) > 0 {
    for i := range resp.Body.Data {
        fmt.Println(resp.Body.Data[i].Item)
    }
} else {
    fmt.Println("There were no items")
}
```

See the [complete example](#) on GitHub.

Note

The [complete example](#) includes the structures for marshaling the JSON request and unmarshaling the JSON response.

Configuring a Lambda Function to Receive Notifications

The following example configures the Lambda function `functionName` to accept notifications from the resource with the ARN `sourceArn`.

The first step is to create the session and Lambda client.

```
// Initialize a session
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create Lambda service client
svc := lambda.New(sess, &aws.Config{Region: aws.String("us-west-2")})
```

Next, we create the structure for the input argument to the `AddPermission` function.

```
permArgs := &lambda.AddPermissionInput{
    Action:      aws.String("lambda:InvokeFunction"),
    FunctionName: functionName,
    Principal:    aws.String("s3.amazonaws.com"),
    SourceArn:    sourceArn,
    StatementId: aws.String("lambda_s3_notification"),
}
```

Finally, we call `AddPermission` and display a message with the result of the call.

```
result, err := svc.AddPermission(permArgs)

if err != nil {
    fmt.Println("Cannot configure function for notifications")
    os.Exit(0)
} else {
    fmt.Println(result)
}
```

See the [complete example](#) on GitHub.

Amazon Polly Examples Using the AWS SDK for Go

Amazon Polly is a cloud service that converts text into lifelike speech. The AWS SDK for Go examples can integrate Amazon Polly into your applications. The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Getting a List of Voices \(p. 109\)](#)
- [Getting a List of Lexicons \(p. 110\)](#)
- [Synthesizing Speech \(p. 111\)](#)

Getting a List of Voices

This example uses the [DescribeVoices](#) operation to get the list of voices in the `us-west-2` region.

Choose Copy to save the code locally.

Create the file `pollyDescribeVoices.go`. Import the packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/polly"

    "fmt"
    "os"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create an Amazon Polly client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

svc := polly.New(sess)
```

Create the input for and call `DescribeVoices`.

```
input := &polly.DescribeVoicesInput{LanguageCode: aws.String("en-US")}
resp, err := svc.DescribeVoices(input)
```

Display the name and gender of the voices.

```
for _, v := range resp.Voices {
    fmt.Println("Name: " + *v.Name)
    fmt.Println("Gender: " + *v.Gender)
    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Getting a List of Lexicons

This example uses the [ListLexicons](#) operation to get the list of lexicons in the `us-west-2` region.

Choose Copy to save the code locally.

Create the file `pollyListLexicons.go`. Import the packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/polly"

    "fmt"
    "os"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create an Amazon Polly client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

svc := polly.New(sess)
```

Call `ListLexicons` and display the name, alphabet, and language code of each lexicon.

```
resp, err := svc.ListLexicons(nil)

for _, l := range resp.Lexicons {
    fmt.Println(*l.Name)
```

```
    fmt.Println("  Alphabet: " + *l.Attributes.Alphabet)
    fmt.Println("  Language: " + *l.Attributes.LanguageCode)
    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Synthesizing Speech

This example uses the [SynthesizeSpeech](#) operation to get the text from a file and produce an MP3 file containing the synthesized speech.

Choose Copy to save the code locally.

Create the file *pollySynthesizeSpeech.go*. Import the packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/polly"

    "fmt"
    "os"
    "strings"
    "io"
    "io/ioutil"
)
```

Get the name of the text file from the command line.

```
if len(os.Args) != 2 {
    fmt.Println("You must supply an alarm name")
    os.Exit(1)
}

fileName := os.Args[1]
```

Open the text file and read the contents as a string.

```
contents, err := ioutil.ReadFile(fileName)
s := string(contents[:])
```

Initialize a session that the SDK will use to load credentials from the shared credentials file `~/.aws/credentials`, load your configuration from the shared configuration file `~/.aws/config`, and create an Amazon Polly client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

svc := polly.New(sess)
```

Create the input for and call `SynthesizeSpeech`.

```
input := &polly.SynthesizeSpeechInput{OutputFormat: aws.String("mp3"), Text: aws.String(s),
    VoiceId: aws.String("Joanna")}

output, err := svc.SynthesizeSpeech(input)
```

Save the resulting synthesized speech as an MP3 file.

```
names := strings.Split(fileName, ".")
name := names[0]
mp3File := name + ".mp3"

outFile, err := os.Create(mp3File)
defer outFile.Close()
```

Note

The resulting MP3 file is in the MPEG-2 format.

See the [complete example](#) on GitHub.

Amazon S3 Examples Using the AWS SDK for Go

Amazon Simple Storage Service (Amazon S3) is storage for the internet. The AWS SDK for Go examples can integrate Amazon S3 into your applications. The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Performing Basic Amazon S3 Bucket Operations \(p. 112\)](#)
- [Creating Pre-Signed URLs for Amazon S3 Buckets \(p. 125\)](#)
- [Using an Amazon S3 Bucket as a Static Web Host \(p. 128\)](#)
- [Working with Amazon S3 CORS Permissions \(p. 131\)](#)
- [Working with Amazon S3 Bucket Policies \(p. 133\)](#)
- [Working with Amazon S3 Bucket ACLs \(p. 136\)](#)
- [Encrypting Amazon S3 Bucket Items \(p. 144\)](#)

Performing Basic Amazon S3 Bucket Operations

These AWS SDK for Go examples show you how to perform the following operations on Amazon S3 buckets and bucket items:

- List the buckets in your account
- Create a bucket
- List the items in a bucket
- Upload a file to a bucket
- Download a bucket item
- Copy a bucket item to another bucket
- Delete a bucket item
- Delete all the items in a bucket
- Restore a bucket item
- Delete a bucket
- List the users with administrator privileges

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In these examples, a series of Go routines are used to perform operations on your Amazon S3 buckets. The routines use the AWS SDK for Go to perform Amazon S3 bucket operations using the following methods of the Amazon S3 client class, unless otherwise noted:

- [ListBuckets](#)
- [CreateBucket](#)
- [ListObjects](#)
- [Upload](#) (from the `s3manager.NewUploader` class)
- [Download](#) (from the `s3manager.NewDownloader` class)
- [CopyObject](#)
- [DeleteObject](#)
- [DeleteObjects](#)
- [RestoreObject](#)
- [DeleteBucket](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with buckets. To learn more, see [Working with Amazon S3 Buckets](#) in the *Amazon S3 Developer Guide*.

List Buckets

The [ListBuckets](#) function lists the buckets in your account.

The following example lists the buckets in your account. There are no command line arguments.

Create the file `s3_list_buckets.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "fmt"  
    "os"  
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*` and create a new Amazon S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Call `ListBuckets`. Passing `nil` means no filters are applied to the returned list. If an error occurs, call `exitErrorrf`. If no error occurs, loop through the buckets, printing the name and creation date of each bucket.

```
result, err := svc.ListBuckets(nil)
if err != nil {
    exitErrorrf("Unable to list buckets, %v", err)
}

fmt.Println("Buckets:")

for _, b := range result.Buckets {
    fmt.Printf("* %s created on %s\n",
        aws.StringValue(b.Name), aws.TimeValue(b.CreationDate))
}
```

See the [complete example](#) on GitHub.

Create a Bucket

The `CreateBucket` function creates a bucket in your account.

The following example creates a bucket with the name specified as a command line argument. You must specify a globally unique name for the bucket.

Create the file `s3_create_bucket.go`. Import the following Go and AWS SDK for Go packages.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function we use to display errors and exit.

```
func exitErrorrf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

The program requires one argument, the name of the bucket to create.

```
if len(os.Args) != 2 {
    exitErrorrf("Bucket name missing!\nUsage: %s bucket_name", os.Args[0])
}

bucket := os.Args[1]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file

<problematic>*</problematic>

~/aws/credentials, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Call `CreateBucket`, passing in the bucket name defined previously. If an error occurs, call `exitErrorf`. If there are no errors, wait for a notification that the bucket was created.

```
_, err = svc.CreateBucket(&s3.CreateBucketInput{
    Bucket: aws.String(bucket),
})
if err != nil {
    exitErrorf("Unable to create bucket %q, %v", bucket, err)
}

// Wait until bucket is created before finishing
fmt.Printf("Waiting for bucket %q to be created...\n", bucket)

err = svc.WaitUntilBucketExists(&s3.HeadBucketInput{
    Bucket: aws.String(bucket),
})
```

If the `WaitUntilBucketExists` call returns an error, call `exitErrorf`. If there are no errors, notify the user of success.

```
if err != nil {
    exitErrorf("Error occurred while waiting for bucket to be created, %v", bucket)
}

fmt.Printf("Bucket %q successfully created\n", bucket)
```

See the [complete example](#) on GitHub.

List Bucket Items

The `ListObjects` function lists the items in a bucket.

The following example lists the items in the bucket with the name specified as a command line argument.

Create the file `s3_list_objects.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {
```



```
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

The program requires one command line argument, the name of the bucket.

```
if len(os.Args) != 2 {
    exitErrorf("Bucket name required\nUsage: %s bucket_name",
        os.Args[0])
}

bucket := os.Args[1]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*</problematic>` `~/aws/credentials`, and create a new Amazon S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Call [ListObjects](#), passing in the name of the bucket. If an error occurs, call `exitErrorf`. If no error occurs, loop through the items, printing the name, last modified date, size, and storage class of each item.

```
resp, err := svc.ListObjects(&s3.ListObjectsInput{Bucket: aws.String(bucket)})
if err != nil {
    exitErrorf("Unable to list items in bucket %q, %v", bucket, err)
}

for _, item := range resp.Contents {
    fmt.Println("Name:      ", *item.Key)
    fmt.Println("Last modified:", *item.LastModified)
    fmt.Println("Size:      ", *item.Size)
    fmt.Println("Storage class:", *item.StorageClass)
    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Upload a File to a Bucket

The [Upload](#) function uploads an object to a bucket.

The following example uploads a file to a bucket with the names specified as command line arguments.

Create the file `s3_upload_object.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3/s3manager"
    "fmt"
    "os"
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Get the bucket and file name from the command line arguments, open the file, and defer the file closing until we are done with it. If an error occurs, call `exitErrorF`.

```
if len(os.Args) != 3 {  
    exitErrorf("bucket and file name required\nUsage: %s bucket_name filename",  
        os.Args[0])  
}  
  
bucket := os.Args[1]  
filename := os.Args[2]  
  
file, err := os.Open(filename)  
if err != nil {  
    exitErrorf("Unable to open file %q, %v", err)  
}  
  
defer file.Close()
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, and create a `NewUploader` object.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Setup the S3 Upload Manager. Also see the SDK doc for the Upload Manager  
// for more information on configuring part size, and concurrency.  
//  
// http://docs.aws.amazon.com/sdk-for-go/api/service/s3/s3manager/#NewUploader  
uploader := s3manager.NewUploader(sess)
```

Upload the file to the bucket. If an error occurs, call `exitErrorF`. Otherwise, notify the user that the upload succeeded.

```
_, err = uploader.Upload(&s3manager.UploadInput{  
    Bucket: aws.String(bucket),  
    Key: aws.String(filename),  
    Body: file,  
})  
if err != nil {  
    // Print the error and exit.  
    exitErrorf("Unable to upload %q to %q, %v", filename, bucket, err)  
}  
  
fmt.Printf("Successfully uploaded %q to %q\n", filename, bucket)
```

See the [complete example](#) on GitHub.

Download a File from a Bucket

The `Download` function downloads an object from a bucket.

The following example downloads an item from a bucket with the names specified as command line arguments.

Create the file `s3_download_object.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "github.com/aws/aws-sdk-go/service/s3/s3manager"  
  
    "fmt"  
    "os"  
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Get the bucket and file name from the command line arguments. If there aren't two arguments, call `exitErrorf`. Otherwise, create the file and defer file closing until we are done downloading. If an error occurs while creating the file, call `exitErrorf`.

```
if len(os.Args) != 3 {  
    exitErrorf("Bucket and item names required\nUsage: %s bucket_name item_name",  
        os.Args[0])  
}  
  
bucket := os.Args[1]  
item := os.Args[2]
```

Initialize the session in us-west-2 that the SDK uses to load credentials from the shared credentials file `<problematic>*``</problematic>` `~/.aws/credentials`, and create a `NewDownloader` object.

```
sess, _ := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
downloader := s3manager.NewDownloader(sess)
```

Download the item from the bucket. If an error occurs, call `exitErrorf`. Otherwise, notify the user that the download succeeded.

```
numBytes, err := downloader.Download(file,  
    &s3.GetObjectInput{  
        Bucket: aws.String(bucket),  
        Key:    aws.String(item),  
    })  
if err != nil {  
    exitErrorf("Unable to download item %q, %v", item, err)  
}  
  
fmt.Println("Downloaded", file.Name(), numBytes, "bytes")
```

See the [complete example](#) on GitHub.

Copy an Item from one Bucket to Another

The `CopyObject` function copies an object from one bucket to another.

The following example copies an item from one bucket to another with the names specified as command line arguments.

Create the file `s3_copy_object.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "fmt"  
    "os"  
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Get the names of the bucket containing the item, the item to copy, and the name of the bucket to which the item is copied. If there aren't four command line arguments, call `exitErrorf`.

```
if len(os.Args) != 4 {  
    exitErrorf("Bucket, item, and other bucket names required\nUsage: go run s3_copy_object  
    bucket item other-bucket")  
}  
  
bucket := os.Args[1]  
item := os.Args[2]  
other := os.Args[3]  
  
source := bucket + "/" + item
```

Initialize the session that the SDK uses to load credentials from the shared credentials file

`<problematic>*</problematic>`

`~/aws/credentials`, and create a new Amazon S3 service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create S3 service client  
svc := s3.New(sess)
```

Call `CopyObject`, with the names of the bucket containing the item, the item to copy, and the name of the bucket to which the item is copied. If an error occurs, call `exitErrorf`. If no error occurs, wait for the item to be copied.

```
_, err = svc.CopyObject(&s3.CopyObjectInput{Bucket: aws.String(other), CopySource:  
    aws.String(source), Key: aws.String(item)})
```

```
if err != nil {
    exitErrorf("Unable to copy item from bucket %q to bucket %q, %v", bucket, other, err)
}

// Wait to see if the item got copied
err = svc.WaitUntilObjectExists(&s3.HeadObjectInput{Bucket: aws.String(other), Key:
    aws.String(item)})}
```

If the `WaitUntilObjectExists` call returns an error, call `exitErrorf`. Otherwise, notify the user that the copy succeeded.

```
if err != nil {
    exitErrorf("Error occurred while waiting for item %q to be copied to bucket %q, %v",
        bucket, item, other, err)
}

fmt.Printf("Item %q successfully copied from bucket %q to bucket %q\n", item, bucket,
    other)
```

See the [complete example](#) on GitHub.

Delete an Item in a Bucket

The `DeleteObject` function deletes an object from a bucket.

The following example deletes an item from a bucket with the names specified as command line arguments.

Create the file `s3_delete_object.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Get the name of the bucket and object to delete.

```
if len(os.Args) != 3 {
    exitErrorf("Bucket and object name required\nUsage: %s bucket_name object_name",
        os.Args[0])
}

bucket := os.Args[1]
obj := os.Args[2]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, and create a new Amazon S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Call `DeleteObject`, passing in the names of the bucket and object to delete. If an error occurs, call `exitErrorrf`. If no error occurs, wait until the object is deleted.

```
_, err = svc.DeleteObject(&s3.DeleteObjectInput{Bucket: aws.String(bucket), Key:
    aws.String(obj)})
if err != nil {
    exitErrorrf("Unable to delete object %q from bucket %q, %v", obj, bucket, err)
}

err = svc.WaitUntilObjectNotExists(&s3.HeadObjectInput{
    Bucket: aws.String(bucket),
    Key:    aws.String(obj),
})
```

If `WaitUntilObjectNotExists` returns an error, call `exitErrorrf`. Otherwise, inform the user that the object was successfully deleted.

```
fmt.Printf("Object %q successfully deleted\n", obj)
```

See the [complete example](#) on GitHub.

Delete All the Items in a Bucket

The `DeleteObjects` function deletes objects from a bucket.

The following example deletes all the items from a bucket with the bucket name specified as a command line argument.

Create the file `s3_delete_objects.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function we use to display errors and exit.

```
func exitErrorrf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Get the name of the bucket.

```
if len(os.Args) != 2 {
```

```
        exitErrorf("Bucket name required\nUsage: %s BUCKET",
            os.Args[0])
    }

    bucket := os.Args[1]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file
<problematic></problematic>
~/aws/credentials, and create a new Amazon S3 service client.

```
sess, _ := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Since buckets can have more than 1000 items and `ListObjects` returns up to 1000 items, set a flag to indicate when more items are available and keep track of how many items we get.

```
hasMoreObjects := true
// Keep track of how many objects we delete
totalObjects := 0
```

Loop through the bucket items up to 1000 at a time, creating the list of objects to delete from the list of items in the bucket. If an error occurs, call `exitErrorf`.

```
for hasMoreObjects {
    resp, err := svc.ListObjects(&s3.ListObjectsInput{Bucket: aws.String(bucket)})
    if err != nil {
        exitErrorf("Unable to list items in bucket %q, %v", bucket, err)
    }

    numObjs := len(resp.Contents)
    totalObjects += numObjs

    // Create Delete object with slots for the objects to delete
    var items s3.Delete
    var objs = make([]*s3.ObjectIdentifier, numObjs)

    for i, o := range resp.Contents {
        // Add objects from command line to array
        objs[i] = &s3.ObjectIdentifier{Key: aws.String(*o.Key)}
    }

    // Add list of objects to delete to Delete object
    items.SetObjects(objs)
```

Call `DeleteObjects`, passing in the name of the bucket and the list of objects to delete. If an error occurs, call `exitErrorf`. Otherwise, determine whether there are more items in the bucket.

```
_, err = svc.DeleteObjects(&s3.DeleteObjectsInput{Bucket: &bucket, Delete: &items})
if err != nil {
    exitErrorf("Unable to delete objects from bucket %q, %v", bucket, err)
}

hasMoreObjects = *resp.IsTruncated
```

Once all of the items in the bucket have been deleted, inform the user of the number of objects deleted.

```
fmt.Println("Deleted", totalObjects, "object(s) from bucket", bucket)
```

See the [complete example](#) on GitHub.

Restore a Bucket Item

The `RestoreObject` function restores an item in a bucket.

The following example restores the items in a bucket with the names specified as command line arguments.

Create the file `s3_restore_object.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "fmt"  
    "os"  
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

The program requires two arguments, the names of the bucket and object to restore.

```
if len(os.Args) != 3 {  
    exitErrorf("Bucket name and object name required\nUsage: %s bucket_name object_name",  
        os.Args[0])  
}  
  
bucket := os.Args[1]  
obj := os.Args[2]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*</problematic>` `~/.aws/credentials`, and create a new Amazon S3 service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create S3 service client  
svc := s3.New(sess)
```

Call `RestoreObject`, passing in the bucket and object names and the number of days to temporarily restore. If an error occurs, call `exitErrorf`. Otherwise, inform the user that the bucket should be restored in the next four hours or so.

```
_, err = svc.RestoreObject(&s3.RestoreObjectInput{Bucket: aws.String(bucket), Key:  
    aws.String(obj), RestoreRequest: &s3.RestoreRequest{Days: aws.Int64(30)}})  
if err != nil {  
    exitErrorf("Could not restore %s in bucket %s, %v", obj, bucket, err)
```



```
}  
  
fmt.Printf("%q should be restored to %q in about 4 hours\n", obj, bucket)
```

See the [complete example](#) on GitHub.

Delete a Bucket

The `DeleteBucket` function deletes a bucket.

The following example deletes the bucket with the name specified as a command line argument.

Create the file `s3_delete_bucket.go`. Import the following Go and AWS SDK for Go packages.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "fmt"  
    "os"  
)
```

Create a function we use to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

The program requires one argument, the name of the bucket to delete. If the argument is not supplied, call `exitErrorf`.

```
if len(os.Args) != 2 {  
    exitErrorf("bucket name required\nUsage: %s bucket_name", os.Args[0])  
}  
  
bucket := os.Args[1]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*</problematic>` `~/.aws/credentials`, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create S3 service client  
svc := s3.New(sess)
```

Call `DeleteBucket`, passing in the bucket name. If an error occurs, call `exitErrorf`. If there are no errors, wait for a notification that the bucket was deleted.

```
_, err = svc.DeleteBucket(&s3.DeleteBucketInput{  
    Bucket: aws.String(bucket),  
})  
if err != nil {  
    exitErrorf("Unable to delete bucket %q, %v", bucket, err)  
}
```

```
// Wait until bucket is deleted before finishing
fmt.Printf("Waiting for bucket %q to be deleted...\n", bucket)

err = svc.WaitUntilBucketNotExists(&s3.HeadBucketInput{
    Bucket: aws.String(bucket),
})
```

If `WaitUntilBucketNotExists` returns an error, call `exitErrorrf`. Otherwise, inform the user that the bucket was successfully deleted.

```
if err != nil {
    exitErrorrf("Error occurred while waiting for bucket to be deleted, %v", bucket)
}

fmt.Printf("Bucket %q successfully deleted\n", bucket)
```

See the [complete example](#) on GitHub.

Creating Pre-Signed URLs for Amazon S3 Buckets

This Go example shows you how to obtain a pre-signed URL for an Amazon S3 bucket. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In this example, a series of Go routines are used to obtain a pre-signed URL for an Amazon S3 bucket using either `GetObject` or a `PUT` operation. A pre-signed URL allows you to grant temporary access to users who don't have permission to directly run AWS operations in your account. A pre-signed URL is signed with your credentials and can be used by any user.

- [Presign](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the SDK.
- You are familiar with pre-signed URLs. To learn more, see [Uploading Objects Using Pre-Signed URLs](#) in the *Amazon S3 Developer Guide*.

Generate a Pre-Signed URL for a `GetObject` Operation

To generate a pre-signed URL, use the [Presign](#) method on the `request` object. You must set an expiration value because the AWS SDK for Go doesn't set one by default.

The following example generates a pre-signed URL that enables you to temporarily share a file without making it public. Anyone with access to the URL can view the file.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "log"
    "time"
)
```

```
// Downloads an item from an S3 Bucket
//
// Usage:
//   go run s3_download.go BUCKET ITEM
func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create S3 service client
    svc := s3.New(sess)

    req, _ := svc.GetObjectRequest(&s3.GetObjectInput{
        Bucket: aws.String("myBucket"),
        Key:     aws.String("myKey"),
    })
    urlStr, err := req.Presign(15 * time.Minute)

    if err != nil {
        log.Println("Failed to sign request", err)
    }

    log.Println("The URL is", urlStr)
}
```

If the SDK has not retrieved your credentials before calling `Presign`, it will get them to generate the pre-signed URL.

Generate a Pre-Signed URL for an Amazon S3 PUT Operation with a Specific Payload

You can generate a pre-signed URL for a PUT operation that checks whether users upload the correct content. When the SDK pre-signs a request, it computes the checksum of the request body and generates an MD5 checksum that is included in the pre-signed URL. Users must upload the same content that produces the same MD5 checksum generated by the SDK; otherwise, the operation fails. This is not the Content-MD5, but the signature. To enforce Content-MD5, simply add the header to the request.

The following example adds a `Body` field to generate a pre-signed PUT operation that requires a specific payload to be uploaded by users.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "log"
    "strings"
    "time"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create S3 service client
    svc := s3.New(sess)
```

```
req, _ := svc.PutObjectRequest(&s3.PutObjectInput{
    Bucket: aws.String("myBucket"),
    Key:    aws.String("myKey"),
    Body:   strings.NewReader("EXPECTED CONTENTS"),
})
str, err := req.Presign(15 * time.Minute)

log.Println("The URL is:", str, " err:", err)
}
```

If you omit the Body field, users can write any contents to the given object.

The following example shows the enforcing of Content-MD5.

```
package main

import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "encoding/base64"
    "fmt"
    "crypto/md5"
    "strings"
    "time"
    "net/http"
)

// Downloads an item from an S3 Bucket in the region configured in the shared config
// or AWS_REGION environment variable.
//
// Usage:
// go run s3_download.go BUCKET ITEM
func main() {
    h := md5.New()
    content := strings.NewReader("")
    content.WriteTo(h)

    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create S3 service client
    svc := s3.New(sess)

    resp, _ := svc.PutObjectRequest(&s3.PutObjectInput{
        Bucket: aws.String("testBucket"),
        Key:    aws.String("testKey"),
    })

    md5s := base64.StdEncoding.EncodeToString(h.Sum(nil))
    resp.HTTPRequest.Header.Set("Content-MD5", md5s)

    url, err := resp.Presign(15 * time.Minute)
    if err != nil {
        fmt.Println("error presigning request", err)
        return
    }

    req, err := http.NewRequest("PUT", url, strings.NewReader(""))
    req.Header.Set("Content-MD5", md5s)
    if err != nil {
```

```
        fmt.Println("error creating request", url)
        return
    }

    defClient, err := http.DefaultClient.Do(req)
    fmt.Println(defClient, err)
}
```

Using an Amazon S3 Bucket as a Static Web Host

This AWS SDK for Go example shows you how to configure an Amazon S3 bucket to act as a static web host. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In this example, you use a series of Go routines to configure any of your buckets to act as a static web host. The routines use the AWS SDK for Go to configure a selected Amazon S3 bucket using these methods of the Amazon S3 client class:

- `GetBucketWebsite`
- `PutBucketWebsite`
- `DeleteBucketWebsite`

For more information about using an Amazon S3 bucket as a static web host, see [Hosting a Static Website on Amazon S3](#) in the *Amazon S3 Developer Guide*.

Prerequisites

- You have [set up](#) (p. 2), and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with hosting static websites on Amazon S3. To learn more, see [Hosting a Static Website on Amazon S3](#) in the *Amazon S3 Developer Guide*.

Retrieve a Bucket's Website Configuration

Create a new Go file named `s3_get_bucket_website.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awsserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

This routine requires you to pass in an argument containing the name of the bucket for which you want to get website configuration.

```
if len(os.Args) != 2 {
    exitErrorrf("bucket name required\nUsage: %s bucket_name", os.Args[0])
}

bucket := os.Args[1]
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Call [GetBucketWebsite](#) to get the bucket configuration. You should check for the `NoSuchWebsiteConfiguration` error code, which tells you that the bucket doesn't have a website configured.

```
result, err := svc.GetBucketWebsite(&s3.GetBucketWebsiteInput{
    Bucket: aws.String(bucket),
})
if err != nil {
    // Check for the NoSuchWebsiteConfiguration error code telling us
    // that the bucket does not have a website configured.
    if awsErr, ok := err.(awserr.Error); ok && awsErr.Code() ==
        "NoSuchWebsiteConfiguration" {
        exitErrorf("Bucket %s does not have website configuration\n", bucket)
    }
    exitErrorf("Unable to get bucket website config, %v", err)
}
```

Print the bucket's website configuration.

```
fmt.Println("Bucket Website Configuration:")
fmt.Println(result)
```

Set a Bucket's Website Configuration

Create a Go file named `s3_set_bucket_website.go` and add the code below. The Amazon [S3PutBucketWebsite](#) operation sets the website configuration on a bucket, replacing any existing configuration.

Create a new Go file named `s3_create_bucket.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
    "path/filepath"
)
```

This routine requires you to pass in an argument containing the name of the bucket and the index suffix page.

```
if len(os.Args) != 3 {
    exitErrorf("bucket name and index suffix page required\nUsage: %s bucket_name\nindex_page [error_page]",
        filepath.Base(os.Args[0]))
}
```

```
bucket := fromArgs(os.Args, 1)
indexSuffix := fromArgs(os.Args, 2)
errorPage := fromArgs(os.Args, 3)
```

Initialize a session that the SDK will use to load configuration, credentials, and region information from the shared credentials file, `~/.aws/credentials`, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Create the parameters to be passed in to [PutBucketWebsite](#), including the bucket name and index document details. If the user passed in an error page when calling the routine, also add that to the parameters.

```
params := s3.PutBucketWebsiteInput{
    Bucket: aws.String(bucket),
    WebsiteConfiguration: &s3.WebsiteConfiguration{
        IndexDocument: &s3.IndexDocument{
            Suffix: aws.String(indexSuffix),
        },
    },
}

// Add the error page if set on CLI
if len(errorPage) > 0 {
    params.WebsiteConfiguration.ErrorDocument = &s3.ErrorDocument{
        Key: aws.String(errorPage),
    }
}
```

Call `PutBucketWebsite`, passing in the parameters you just defined. If an error occurs, print the error details and exit the routine.

```
_, err = svc.PutBucketWebsite(&params)
if err != nil {
    exitErrorf("Unable to set bucket %q website configuration, %v",
        bucket, err)
}

fmt.Printf("Successfully set bucket %q website configuration\n", bucket)
```

Delete Website Configuration on a Bucket

Create a Go file named `s3_delete_bucket_website.go`. Import the relevant Go and AWS SDK for Go packages.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
    "path/filepath"
)
```

This routine requires you to pass in the name of the bucket for which you want to delete the website configuration.

```
if len(os.Args) != 2 {
    exitErrorf("bucket name required\nUsage: %s bucket_name",
        filepath.Base(os.Args[0]))
}

bucket := os.Args[1]
```

Initialize a session that the SDK will use to load configuration, credentials, and region information from the shared credentials file, `~/.aws/credentials`, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create S3 service client
svc := s3.New(sess)
```

Call `DeleteBucketWebsite` and pass in the name of the bucket to complete the action.

```
_, err = svc.DeleteBucketWebsite(&s3.DeleteBucketWebsiteInput{
    Bucket: aws.String(bucket),
})
if err != nil {
    exitErrorf("Unable to delete bucket %q website configuration, %v",
        bucket, err)
}

fmt.Printf("Successfully delete bucket %q website configuration\n", bucket)
```

Working with Amazon S3 CORS Permissions

This AWS SDK for Go example shows you how to list Amazon S3 buckets and configure CORS and bucket logging. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In this example, a series of Go routines are used to list your Amazon S3 buckets and to configure CORS and bucket logging. The routines use the AWS SDK for Go to configure a selected Amazon S3 bucket using these methods of the Amazon S3 client class:

- [GetBucketCors](#)
- [PutBucketCors](#)

If you are unfamiliar with using CORS configuration with an Amazon S3 bucket, it's worth your time to read [Cross-Origin Resource Sharing \(CORS\)](#) in the *Amazon S3 Developer Guide* before proceeding.

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with using CORS configuration with an Amazon S3 bucket. To learn more, see [Cross-Origin Resource Sharing \(CORS\)](#) in the *Amazon S3 Developer Guide*.

Configure CORS on the Bucket

Create a new Go file named `s3_set_cors.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "flag"  
    "fmt"  
    "os"  
    "strings"  
)
```

This routine configures CORS rules for a bucket by setting the allowed HTTP methods.

It requires the bucket name and can also take a space-separated list of HTTP methods. Using the Go Language's `flag` package, it parses the input and validates the bucket name.

```
var bucket string  
// Setup the CLI options and validation  
flag.StringVar(&bucket, "b", "", "Bucket to set CORS on, (required)")  
flag.Parse()  
  
if len(bucket) == 0 {  
    exitErrorf("-b <bucket> Bucket name required")  
}  
  
methods := filterMethods(flag.Args())
```

Notice the helper method, `filterMethods`, which ensures the methods passed in are uppercase.

```
func filterMethods(methods []string) []string {  
    filtered := make([]string, 0, len(methods))  
    for _, m := range methods {  
        v := strings.ToUpper(m)  
        switch v {  
            case "POST", "GET", "PUT", "PATCH", "DELETE":  
                filtered = append(filtered, v)  
            }  
    }  
  
    return filtered  
}
```

Initialize a session that the SDK will use to load credentials, from the shared credentials file, `~/.aws/credentials`, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create S3 service client  
svc := s3.New(sess)
```

Create a new [CORSRule](#) to set up the CORS configuration.

```
rule := s3.CORSRule{
```

```
AllowedHeaders: aws.StringSlice([]string{"Authorization"}),
AllowedOrigins: aws.StringSlice([]string{"*"}),
MaxAgeSeconds:  aws.Int64(3000),

// Add HTTP methods CORS request that were specified in the CLI.
AllowedMethods: aws.StringSlice(methods),
}
```

Add the `CORSRule` to the `PutBucketCorsInput` structure, call `PutBucketCors` with that structure, and print a success or error message.

```
params := s3.PutBucketCorsInput{
    Bucket: aws.String(bucket),
    CORSConfiguration: &s3.CORSConfiguration{
        CORSRules: []*s3.CORSRule{&rule},
    },
}

_, err = svc.PutBucketCors(&params)
if err != nil {
    // Print the error message
    exitErrorf("Unable to set Bucket %q's CORS, %v", bucket, err)
}

// Print the updated CORS config for the bucket
fmt.Printf("Updated bucket %q CORS for %v\n", bucket, methods)
```

Working with Amazon S3 Bucket Policies

This AWS SDK for Go example shows you how to retrieve, display, and set Amazon S3 bucket policies. You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In this example, you use a series of Go routines to retrieve or set a bucket policy on an Amazon S3 bucket. The routines use the AWS SDK for Go to configure policy for a selected Amazon S3 bucket using these methods of the Amazon S3 client class:

- [GetBucketPolicy](#)
- [PutBucketPolicy](#)
- [DeleteBucketPolicy](#)

For more information about bucket policies for Amazon S3 buckets, see [Using Bucket Policies and User Policies](#) in the *Amazon S3 Developer Guide*.

Prerequisites

- You have [set up \(p. 2\)](#), and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with Amazon S3 bucket policies. To learn more, see [Using Bucket Policies and User Policies](#) in the *Amazon S3 Developer Guide*.

Retrieve and Display a Bucket Policy

Create a new Go file named `s3_get_bucket_policy.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/awser"   
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "bytes"  
    "encoding/json"  
    "fmt"  
    "os"  
    "path/filepath"  
)
```

Creat the `exitError` function to deal with errors.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)   
    os.Exit(1)  
}
```

This routine prints the policy for a bucket. If the bucket doesn't exist, or there was an error, an error message is printed instead. It requires the bucket name as input.

```
if len(os.Args) != 2 {  
    exitErrorf("bucket name required\nUsage: %s bucket_name",  
        filepath.Base(os.Args[0]))  
}  
  
bucket := os.Args[1]
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`, and create a new S3 service client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create S3 service client  
svc := s3.New(sess)
```

Call `GetBucketPolicy` to fetch the policy, then display any errors.

```
result, err := svc.GetBucketPolicy(&s3.GetBucketPolicyInput{  
    Bucket: aws.String(bucket),  
})  
if err != nil {  
    // Special error handling for the when the bucket doesn't  
    // exists so we can give a more direct error message from the CLI.  
    if aerr, ok := err.(awser.Error); ok {  
        switch aerr.Code() {  
            case s3.ErrCodeNoSuchBucket:  
                exitErrorf("Bucket %q does not exist.", bucket)  
            case "NoSuchBucketPolicy":  
                exitErrorf("Bucket %q does not have a policy.", bucket)  
        }  
    }  
    exitErrorf("Unable to get bucket %q policy, %v.", bucket, err)  
}
```

Use Go's JSON package to print the Policy JSON returned by the call.

```
out := bytes.Buffer{}
policyStr := aws.StringValue(result.Policy)
if err := json.Indent(&out, []byte(policyStr), "", " "); err != nil {
    exitErrorf("Failed to pretty print bucket policy, %v.", err)
}

fmt.Printf("%q's Bucket Policy:\n", bucket)
fmt.Println(out.String())
```

Set Bucket Policy

This routine sets the policy for a bucket. If the bucket doesn't exist, or there was an error, an error message will be printed instead. It requires the bucket name as input. It also requires the same Go and AWS SDK for Go packages as the previous example, except for the bytes Go package.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awsserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "encoding/json"
    "fmt"
    "os"
    "path/filepath"
)
```

Add the main function and parse the arguments to get the bucket name.

```
if len(os.Args) != 2 {
    exitErrorf("bucket name required\nUsage: %s bucket_name",
        filepath.Base(os.Args[0]))
}
bucket := os.Args[1]
```

Initialize a session that the SDK will use to load configuration, credentials, and region information from the shared credentials file, ~/.aws/credentials, and create a new S3 service client.

```
    Region: aws.String("us-west-2")),
)

// Create S3 service client
svc := s3.New(sess)
```

Create a policy using the map interface, filling in the bucket as the resource.

```
"Version": "2012-10-17",
"Statement": []map[string]interface{}{
    {
        "Sid":      "AddPerm",
        "Effect":   "Allow",
        "Principal": "*",
        "Action": []string{
            "s3:GetObject",
        },
        "Resource": []string{
```

```
        fmt.Sprintf("arn:aws:s3:::%s/*", bucket),  
    },  
    },  
}
```

Use Go's JSON package to marshal the policy into a JSON value so that it can be sent to S3.

```
if err != nil {  
    exitErrorf("Failed to marshal policy, %v", err)  
}
```

Call the S3 client's [PutBucketPolicy](#) to PUT the policy for the bucket and print the results.

```
    Bucket: aws.String(bucket),  
    Policy: aws.String(string(policy)),  
})  
if err != nil {  
    if aerr, ok := err.(awserr.Error); ok && aerr.Code() == s3.ErrCodeNoSuchBucket {  
        // Special error handling for the when the bucket doesn't  
        // exists so we can give a more direct error message from the CLI.  
        exitErrorf("Bucket %q does not exist", bucket)  
    }  
    exitErrorf("Unable to set bucket %q policy, %v", bucket, err)  
}  
  
fmt.Printf("Successfully set bucket %q's policy\n", bucket)  
}
```

The `exitError` function is used to deal with printing any errors.

```
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Working with Amazon S3 Bucket ACLs

The following examples use AWS SDK for Go functions to:

- Get the access control lists (ACLs) on a bucket
- Get the ACLs on a bucket item
- Add a new user to the ACLs on a bucket
- Add a new user to the ACLs on a bucket item

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

In these examples, a series of Go routines are used to manage ACLs on your Amazon S3 buckets. The routines use the AWS SDK for Go to perform Amazon S3 bucket operations using the following methods of the Amazon S3 client class:

- [GetBucketAcl](#)
- [GetObjectAcl](#)
- [PutBucketAcl](#)
- [PutObjectAcl](#)

Prerequisites

- You have [set up \(p. 2\)](#), and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with Amazon S3 bucket ACLs. To learn more, see [Managing Access with ACLs](#) in the *Amazon S3 Developer Guide*.

Get a Bucket ACL

The [GetBucketAcl](#) function gets the ACLs on a bucket.

The following example gets the ACLs on a bucket with the name specified as a command line argument.

Create the file `s3_get_bucket_acl.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
    "fmt"  
    "os"  
)
```

Create a function to display errors and exit.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

This example requires one input parameter, the name of the bucket. If the name is not supplied, call the error function and exit.

```
if len(os.Args) != 2 {  
    exitErrorf("Bucket name required\nUsage: go run", os.Args[0], "BUCKET")  
}  
  
bucket := os.Args[1]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, the region from the shared configuration file `~/.aws/config`, and create a new Amazon S3 service client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
// Create S3 service client  
svc := s3.New(sess)
```

Call `GetBucketAcl`, passing in the name of the bucket. If an error occurs, call `exitErrorrf`. If no error occurs, loop through the results and print out the name, type, and permission for the grantees.

```
result, err := svc.GetBucketAcl(&s3.GetBucketAclInput{Bucket: &bucket})
if err != nil {
    exitErrorrf(err.Error())
}

fmt.Println("Owner:", *result.Owner.DisplayName)
fmt.Println("")
fmt.Println("Grants")

for _, g := range result.Grants {
    // If we add a canned ACL, the name is nil
    if g.Grantee.DisplayName == nil {
        fmt.Println("  Grantee:    EVERYONE")
    } else {
        fmt.Println("  Grantee:   ", *g.Grantee.DisplayName)
    }

    fmt.Println("  Type:      ", *g.Grantee.Type)
    fmt.Println("  Permission:", *g.Permission)
    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Set a Bucket ACL

The `PutBucketAcl` function sets the ACLs on a bucket.

The following example gives a user access by email address to a bucket with the bucket name and email address specified as command line arguments. The user can also supply a permission argument. However, if it isn't supplied, the user is given READ access to the bucket.

Create the file `s3_put_bucket_acl.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function to display errors and exit.

```
func exitErrorrf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Get the two required input parameters. If the optional permission parameter is supplied, make sure it is one of the allowed values. If not, print an error message and quit.

```
if len(os.Args) < 3 {
    exitErrorrf("Bucket name and email address required; permission optional (READ if omitted)\nUsage: go run", os.Args[0], "BUCKET EMAIL [PERMISSION]")
}
```

```
}

bucket := os.Args[1]
address := os.Args[2]

permission := "READ"

if len(os.Args) == 4 {
    permission = os.Args[3]

    if !(permission == "FULL_CONTROL" || permission == "WRITE" || permission == "WRITE_ACP" || permission == "READ" || permission == "READ_ACP") {
        fmt.Println("Illegal permission value. It must be one of:")
        fmt.Println("FULL_CONTROL, WRITE, WRITE_ACP, READ, or READ_ACP")
        os.Exit(1)
    }
}
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, the region from the shared configuration file `~/.aws/config`, and create a new Amazon S3 service client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create S3 service client
svc := s3.New(sess)
```

Get the existing ACLs and append the new user to the list. If there is an error while retrieving the list, print an error message and quit.

```
result, err := svc.GetBucketAcl(&s3.GetBucketAclInput{Bucket: &bucket})
if err != nil {
    exitErrorf(err.Error())
}

owner := *result.Owner.DisplayName
ownerId := *result.Owner.ID

// Existing grants
grants := result.Grants

// Create new grantee to add to grants
var newGrantee = s3.Grantee{EmailAddress: &address, Type: &userType}
var newGrant = s3.Grant{Grantee: &newGrantee, Permission: &permission}

// Add them to the grants
grants = append(grants, &newGrant)
```

Build the parameter list for the call based on the existing ACLs and the new user information.

```
params := &s3.PutBucketAclInput{
    Bucket: &bucket,
    AccessControlPolicy: &s3.AccessControlPolicy{
        Grants: grants,
        Owner: &s3.Owner{
            DisplayName: &owner,
            ID:          &ownerId,
        },
    },
},
```



```
}
```

Call `PutBucketAcl`, passing in the parameter list. If an error occurs, display a message and quit. Otherwise, display a message indicating success.

```
_, err = svc.PutBucketAcl(params)
if err != nil {
    exitErrorrf(err.Error())
}

fmt.Println("Congratulations. You gave user with email address", address, permission,
"permission to bucket", bucket)
```

See the [complete example](#) on GitHub.

Making a Bucket Public using a Canned ACL

As in the previous example, the `PutBucketAcl` function sets the ACLs on a bucket.

The following example gives everyone read-only access to the items in the bucket with the bucket name specified as a command line argument.

Create the file `s3_make_bucket_public.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"

    "fmt"
    "os"
)
```

Create a function to display errors and exit.

```
func exitErrorrf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Get the required input parameter and create the **acl**.

```
if len(os.Args) < 2 {
    exitErrorrf("Bucket name required.\nUsage: go run", os.Args[0], "BUCKET")
}

bucket := os.Args[1]
acl := "public-read"
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials` the region from the shared configuration file `~/.aws/config`, and create a new Amazon S3 service client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))
svc := s3.New(sess)
```

Create the input for and call `PutBucketAcl`. If an error occurs, display a message and quit. Otherwise, display a message indicating success.

```
params := &s3.PutBucketAclInput{
    Bucket: &bucket,
    ACL: &acl,
}

// Set bucket ACL
_, err := svc.PutBucketAcl(params)
if err != nil {
    exitErrorrf(err.Error())
}

fmt.Println("Bucket " + bucket + " is now public")
```

See the [complete example](#) on GitHub.

Get a Bucket Object ACL

The `PutObjectAcl` function sets the ACLs on a bucket item.

The following example gets the ACLs for a bucket item with the bucket and item name specified as command line arguments.

Create the file `s3_get_bucket_object_acl.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function to display errors and exit.

```
func exitErrorrf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

This example requires two input parameters, the names of the bucket and object. If either name is not supplied, call the error function and exit.

```
if len(os.Args) != 3 {
    exitErrorrf("Bucket and object names required\nUsage: go run", os.Args[0], "BUCKET\nOBJECT")
}

bucket := os.Args[1]
key := os.Args[2]
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `~/.aws/credentials`, the region from the shared configuration file `~/.aws/config`, and create a new Amazon S3 service client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
```

```
    SharedConfigState: session.SharedConfigEnable,
  )))

// Create S3 service client
svc := s3.New(sess)
```

Call `GetObjectAcl`, passing in the names of the bucket and object. If an error occurs, call `exitErrorrf`. If no error occurs, loop through the results and print out the name, type, and permission for the grantees.

```
result, err := svc.GetObjectAcl(&s3.GetObjectAclInput{Bucket: &bucket, Key: &key})
if err != nil {
    exitErrorrf(err.Error())
}

fmt.Println("Owner:", *result.Owner.DisplayName)
fmt.Println("")
fmt.Println("Grants")

for _, g := range result.Grants {
    fmt.Println("  Grantee:   ", *g.Grantee.DisplayName)
    fmt.Println("  Type:      ", *g.Grantee.Type)
    fmt.Println("  Permission:", *g.Permission)
    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Set a Bucket Object ACL

The `PutObjectAcl` function sets the ACLs on a bucket item.

The following example gives a user access by email address to a bucket item, with the bucket and item names and email address specified as command line arguments. The user can also supply a permission argument. However, if it isn't supplied, the user is given READ access to the bucket.

Create the file `s3_put_bucket_object_acl.go`. Add the following statements to import the Go and AWS SDK for Go packages used in the example.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "fmt"
    "os"
)
```

Create a function to display errors and exit.

```
func exitErrorrf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Get the three required input parameters. If the optional permission parameter is supplied, make sure it is one of the allowed values. If not, print an error message and quit.

```
if len(os.Args) < 4 {
    exitErrorrf("Bucket name, object name, and email address required; permission optional\n(READ if omitted)\nUsage: go run", os.Args[0], "BUCKET OBJECT EMAIL [PERMISSION]")
}
```

```
}

bucket := os.Args[1]
key := os.Args[2]
address := os.Args[3]

permission := "READ"

if len(os.Args) == 5 {
    permission = os.Args[4]

    if !(permission == "FULL_CONTROL" || permission == "WRITE" || permission == "WRITE_ACP" || permission == "READ" || permission == "READ_ACP") {
        fmt.Println("Illegal permission value. It must be one of:")
        fmt.Println("FULL_CONTROL, WRITE, WRITE_ACP, READ, or READ_ACP")
        os.Exit(1)
    }
}
}
```

Initialize the session that the SDK uses to load credentials from the shared credentials file `<problematic>*</problematic>` `~/.aws/credentials`, and create a new Amazon S3 service client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

// Create S3 service client
svc := s3.New(sess)
```

Build the parameter list for the call based on the existing ACLs and the new user information.

```
result, err := svc.GetObjectAcl(&s3.GetObjectAclInput{Bucket: &bucket, Key: &key})
if err != nil {
    exitErrorf(err.Error())
}

owner := *result.Owner.DisplayName
ownerId := *result.Owner.ID

// Existing grants
grants := result.Grants

// Create new grantee to add to grants
userType := "AmazonCustomerByEmail"
var newGrantee = s3.Grantee{EmailAddress: &address, Type: &userType}
var newGrant = s3.Grant{Grantee: &newGrantee, Permission: &permission}

// Add them to the grants
grants = append(grants, &newGrant)

params := &s3.PutObjectAclInput{
    Bucket: &bucket,
    Key: &key,
    AccessControlPolicy: &s3.AccessControlPolicy{
        Grants: grants,
        Owner: &s3.Owner{
            DisplayName: &owner,
            ID: &ownerId,
        },
    },
},
}
```

Call [PutObjectAcl](#), passing in the parameter list. If an error occurs, display a message and quit. Otherwise, display a message indicating success.

```
_, err = svc.PutObjectAcl(params)
if err != nil {
    exitErrorrf(err.Error())
}

fmt.Println("Congratulations. You gave user with email address", address, permission,
    "permission to bucket", bucket, "object", key)
```

See the [complete example](#) on GitHub.

Encrypting Amazon S3 Bucket Items

Amazon S3 supports encrypting Amazon S3 bucket objects on both the client and the server. To encrypt objects on the client, you perform the encryption yourself, either using keys that you create or keys that AWS Key Management Service (AWS KMS) manages for you.

To encrypt objects on the server, you have more options.

- You can have Amazon S3 automatically encrypt objects as you upload them to a bucket. Once you configure a bucket with this option, every object that you upload--from that point on--is encrypted.
- You can have Amazon S3 encrypt an object when you upload it to a bucket. The disadvantage with this approach is that you can still upload objects that are not encrypted.
- You can have Amazon S3 reject objects that you attempt to upload to a bucket without requesting Amazon S3 encrypt the items.

The following examples describe some of these options.

Learn about encryption in Amazon S3 at [Protecting Data Using Encryption](#).

Topics

- [Setting Default Server-Side Encryption for an Amazon S3 Bucket \(p. 144\)](#)
- [Requiring Encryption on the Server to Upload Amazon S3 Bucket Objects \(p. 145\)](#)
- [Encrypting an Amazon S3 Bucket Object on the Server Using AWS KMS \(p. 147\)](#)
- [Encrypting an Amazon S3 Bucket Object with a User-Supplied AWS KMS Key \(p. 148\)](#)
- [Decrypting an Amazon S3 Bucket Object with a User-Supplied AWS KMS Key \(p. 149\)](#)

Setting Default Server-Side Encryption for an Amazon S3 Bucket

The following example uses the [PutBucketEncryption](#) method to enable KMS server-side encryption on any objects added to `myBucket`.

The only exception is if the user configures their request to explicitly use server-side encryption. In that case, the specified encryption takes precedence.

Choose Copy to save the code locally.

Create the file `set_default_encryption.go`.

Import the required packages.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
  
    "fmt"  
    "os"  
)
```

Get the KMS key from the command line, where key is a KMS key ID as created in the [Creating a CMK in AWS Key Management Service \(p. 101\)](#) example, and set the bucket name.

```
if len(os.Args) != 2 {  
    fmt.Println("You must supply a key")  
    os.Exit(1)  
}  
  
key := os.Args[1]  
bucket := "myBucket"
```

Create a session and Amazon S3 client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
svc := s3.New(sess)
```

Create the input for and call `PutBucketEncryption`, and display a success message.

```
defEnc := &s3.ServerSideEncryptionByDefault{KMSEncryption: aws.String(key), SSEAlgorithm:  
    aws.String(s3.ServerSideEncryptionAwsKms)}  
rule := &s3.ServerSideEncryptionRule{ApplyServerSideEncryptionByDefault: defEnc}  
rules := []*s3.ServerSideEncryptionRule{rule}  
serverConfig := &s3.ServerSideEncryptionConfiguration{Rules: rules}  
input := &s3.PutBucketEncryptionInput{Bucket: aws.String(bucket),  
    ServerSideEncryptionConfiguration: serverConfig}  
_, err := svc.PutBucketEncryption(input)  
fmt.Println("Bucket " + bucket + " now has KMS encryption by default")
```

See the [complete example](#) on GitHub.

Requiring Encryption on the Server to Upload Amazon S3 Bucket Objects

The following example uses the `PutBucketPolicy` method to require that objects uploaded to an Amazon S3 bucket have Amazon S3 encrypt the object with an AWS KMS key. Attempts to upload an object without specifying that Amazon S3 encrypt the object with an AWS KMS key raise an `Aws::S3::Errors::AccessDenied` exception.

Avoid using this configuration option if you use default server-side encryption as described in [Setting Default Server-Side Encryption for an Amazon S3 Bucket \(p. 144\)](#) as they could conflict and result in unexpected results.

Choose Copy to save the code locally.

Create the file `require_server_encryption.go`.

Import the required packages.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
  
    "flag"  
    "fmt"  
    "os"  
    "encoding/json"  
)
```

Set the name of the bucket, create a session, and create an Amazon S3 client.

```
bucket := "myBucket"  
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
svc := s3.New(sess)
```

Create an Amazon S3 policy that requires server-side KMS encryption on objects uploaded to the bucket.

```
PolicyDoc := map[string]interface{}{  
    "Version": "2012-10-17",  
    "Statement": []map[string]interface{}{  
        {  
            "Sid": "DenyIncorrectEncryptionHeader",  
            "Effect": "Deny",  
            "Principal": "*",  
            "Action": "s3:PutObject",  
            "Resource": "arn:aws:s3:::" + bucket + "/*",  
            "Condition": map[string]interface{}{  
                "StringNotEquals": map[string]interface{}{  
                    "s3:x-amz-server-side-encryption": "aws:kms",  
                },  
            },  
        },  
        {  
            "Sid": "DenyUnEncryptedObjectUploads",  
            "Effect": "Deny",  
            "Principal": "*",  
            "Action": "s3:PutObject",  
            "Resource": "arn:aws:s3:::" + bucket + "/*",  
            "Condition": map[string]interface{}{  
                "Null": map[string]interface{}{  
                    "s3:x-amz-server-side-encryption": "true",  
                },  
            },  
        },  
    },  
}
```

Convert the policy into JSON, create the input for and call `PutBucketPolicy`, apply the policy to the bucket, and print a success message.

```
policy, err := json.Marshal(PolicyDoc)  
input := &s3.PutBucketPolicyInput{  
    Bucket: bucketPtr,  
    Policy: aws.String(string(policy)),  
}
```

```
}  
_, err = svc.PutBucketPolicy(input)  
fmt.Println("Set policy for " + bucket)
```

See the [complete example](#) on GitHub.

Encrypting an Amazon S3 Bucket Object on the Server Using AWS KMS

The following example uses the [PutObject](#) method to add the object `myItem` to the bucket `myBucket` with server-side encryption set to AWS KMS.

Note that this differs from [Setting Default Server-Side Encryption for an Amazon S3 Bucket \(p. 144\)](#), is in that case, the objects are encrypted without you having to explicitly perform the operation.

Choose Copy to save the code locally.

Create the file `encrypt_object_on_server.go`.

Add the required packages.

```
import (  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/s3"  
  
    "fmt"  
    "os"  
    "strings"  
)
```

Get the KMS key from the command line, where `key` is a KMS key ID as created in the [Creating a CMK in AWS Key Management Service \(p. 101\)](#) example, and set the bucket and object names.

```
if len(os.Args) != 2 {  
    fmt.Println("You must supply a key")  
    os.Exit(1)  
}  
  
key := os.Args[1]  
bucket := "myBucket"  
object := "myItem"
```

Create a session and Amazon S3 client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{  
    SharedConfigState: session.SharedConfigEnable,  
}))  
  
svc := s3.New(sess)
```

Create input for and call `put_object`. Notice that the `server_side_encryption` property is set to `aws:kms`, indicating that Amazon S3 encrypts the object using AWS KMS, and display a success message to the user.

```
input := &s3.PutObjectInput{  
    Body: strings.NewReader(object),
```



```
    Bucket:      aws.String(bucket),
    Key:         aws.String(object),
    ServerSideEncryption: aws.String("aws:kms"),
    SSEKMSKeyId:  aws.String(key),
}

_, err := svc.PutObject(input)
fmt.Println("Added object " + obj + " to bucket " + bucket + " with AWS KMS encryption")
```

See the [complete example](#) on GitHub.

Encrypting an Amazon S3 Bucket Object with a User-Supplied AWS KMS Key

The following example uses the [PutObject](#) method to add the object `myObject` to the bucket `myBucket`.

Choose Copy to save the code locally.

Create the file `encrypt_object.go`.

Include the required packages.

```
import (
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/kms"
    "github.com/aws/aws-sdk-go/service/s3"
    "github.com/aws/aws-sdk-go/service/s3/s3crypto"

    "fmt"
    "os"
    "strings"
)
```

Get AWS KMS key from the command line, where `key` is an AWS KMS key ID as created in the [Creating a CMK in AWS Key Management Service \(p. 101\)](#) example, and set the bucket and object names.

```
if len(os.Args) != 2 {
    fmt.Println("You must supply a bucket name, object name, and key")
    os.Exit(1)
}

key := os.Args[0]
bucket := "myBucket"
object := "myObject"
```

Create a session and encryption client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))

handler := s3crypto.NewKMSKeyGenerator(kms.New(sess), key)
svc := s3crypto.NewEncryptionClient(sess, s3crypto.AESGCMContentCipherBuilder(handler))
```

Create the input for and call `PutObject` to upload the object to the bucket, and display a success message.

```
input := &s3.PutObjectInput{
    Body:    strings.NewReader(object),
    Bucket:  aws.String(bucket),
    Key:     aws.String(object),
}

_, err := svc.PutObject(input)
fmt.Println("Added object " + object + " to bucket " + bucket + " with AWS KMS encryption on the client")
```

See the [complete example](#) on GitHub.

Decrypting an Amazon S3 Bucket Object with a User-Supplied AWS KMS Key

The following example uses the [GetObject](#) method to get the object `myObject` from the bucket `myBucket`.

Choose Copy to save the code locally.

Create the file `decrypt_object.go`.

Import the required packages.

```
import (
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "github.com/aws/aws-sdk-go/service/s3/s3crypto"

    "fmt"
    "os"
    "bytes"
)
```

Get the AWS KMS key from the command line, where `key` is an AWS KMS key ID as created in the [Creating a CMK in AWS Key Management Service \(p. 101\)](#) example and must be the same value you used to encrypt the object, and set the name of the bucket and object.

```
if len(os.Args) != 2 {
    fmt.Println("You must supply a key")
    os.Exit(1)
}

key := os.Args[1]
bucket := "myBucket"
object := "myObject"
```

Create a session and an Amazon S3 encryption client.

```
sess := session.Must(session.NewSessionWithOptions(session.Options{
    SharedConfigState: session.SharedConfigEnable,
}))
svc := s3crypto.NewDecryptionClient(sess)
```

Create input for and call `get_object` to get the object.

```
input := &s3.GetObjectInput{
```

```
    Bucket: aws.String(bucket),
    Key:    aws.String(object),
}

resp, err := svc.GetObject(input)
```

Save the object and display a success message.

```
outFile, err := os.Create(obj)
defer outFile.Close()
_, err = io.Copy(outFile, resp.Body)
fmt.Println("Saved " + obj + " from bucket " + bucket + ":")
```

See the [complete example](#) on GitHub.

Amazon SES Examples Using the AWS SDK for Go

Amazon Simple Email Service (Amazon SES) is an email platform that provides an easy, cost-effective way for you to send and receive email using your own email addresses and domains. You can use the following examples to access Amazon SES using the AWS SDK for Go. For more information about Amazon SES, see the [Amazon SES documentation](#).

Topics

- [Listing Valid Amazon SES Email Addresses \(p. 150\)](#)
- [Verifying an Email Address in Amazon SES \(p. 151\)](#)
- [Sending a Message to an Email Address in Amazon SES \(p. 152\)](#)
- [Deleting an Email Address in Amazon SES \(p. 154\)](#)
- [Getting Amazon SES Statistics \(p. 155\)](#)

Listing Valid Amazon SES Email Addresses

The following example demonstrates how to use the AWS SDK for Go to list the valid Amazon SES email addresses.

```
package main

import (
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ses"
)

func main() {
    // Initialize a session in us-west-2 that the SDK will use to load
    // credentials from the shared credentials file ~/.aws/credentials.
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create SES service client
    svc := ses.New(sess)
```

```
    result, err := svc.ListIdentities(&ses.ListIdentitiesInput{IdentityType:
aws.String("EmailAddress")})

    if err != nil {
        fmt.Println(err)
        os.Exit(1)
    }

    for _, email := range result.Identities {
        var e = []*string{email}

        verified, err :=
svc.GetIdentityVerificationAttributes(&ses.GetIdentityVerificationAttributesInput{Identities:
e})

        if err != nil {
            fmt.Println(err)
            os.Exit(1)
        }

        for _, va := range verified.VerificationAttributes {
            if *va.VerificationStatus == "Success" {
                fmt.Println(*email)
            }
        }
    }
}
```

See the [complete example](#) on GitHub.

Verifying an Email Address in Amazon SES

The following example demonstrates how to use the AWS SDK for Go to verify an Amazon SES email address.

```
package main

import (
    "fmt"

    //go get -u github.com/aws/aws-sdk-go
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ses"
    "github.com/aws/aws-sdk-go/aws/awsserr"
)

const (
    // Replace sender@example.com with your "From" address.
    // This address must be verified with Amazon SES.
    Sender = "sender@example.com"

    // Replace recipient@example.com with a "To" address. If your account
    // is still in the sandbox, this address must be verified.
    Recipient = "recipient@example.com"
)

func main() {
    // Create a new session in the us-west-2 region.
    // Replace us-west-2 with the AWS Region you're using for Amazon SES.
    sess, err := session.NewSession(&aws.Config{
        Region:aws.String("us-west-2")},
    )
```

```
// Create an SES session.
svc := ses.New(sess)

// Attempt to send the email.
_, err = svc.VerifyEmailAddress(&ses.VerifyEmailAddressInput{EmailAddress:
aws.String(Recipient)})

// Display error messages if they occur.
if err != nil {
    if aerr, ok := err.(awserr.Error); ok {
        switch aerr.Code() {
            case ses.ErrCodeMessageRejected:
                fmt.Println(ses.ErrCodeMessageRejected, aerr.Error())
            case ses.ErrCodeMailFromDomainNotVerifiedException:
                fmt.Println(ses.ErrCodeMailFromDomainNotVerifiedException, aerr.Error())
            case ses.ErrCodeConfigurationSetDoesNotExistException:
                fmt.Println(ses.ErrCodeConfigurationSetDoesNotExistException, aerr.Error())
            default:
                fmt.Println(aerr.Error())
        }
    } else {
        // Print the error, cast err to awserr.Error to get the Code and
        // Message from an error.
        fmt.Println(err.Error())
    }
}

return
}

fmt.Println("Verification sent to address: " + Recipient)
}
```

See the [complete example](#) on GitHub.

Sending a Message to an Email Address in Amazon SES

The following example demonstrates how to use the AWS SDK for Go to send a message to an Amazon SES email address.

```
package main

import (
    "fmt"

    //go get -u github.com/aws/aws-sdk-go
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ses"
    "github.com/aws/aws-sdk-go/aws/awserr"
)

const (
    // Replace sender@example.com with your "From" address.
    // This address must be verified with Amazon SES.
    Sender = "sender@example.com"

    // Replace recipient@example.com with a "To" address. If your account
    // is still in the sandbox, this address must be verified.
    Recipient = "recipient@example.com"
}
```

```
// Specify a configuration set. To use a configuration
// set, comment the next line and line 92.
//ConfigurationSet = "ConfigSet"

// The subject line for the email.
Subject = "Amazon SES Test (AWS SDK for Go)"

// The HTML body for the email.
HtmlBody = "<h1>Amazon SES Test Email (AWS SDK for Go)</h1><p>This email was sent with
" +
    "<a href='https://aws.amazon.com/ses/'>Amazon SES</a> using the " +
    "<a href='https://aws.amazon.com/sdk-for-go/'>AWS SDK for Go</a>.</p>"

//The email body for recipients with non-HTML email clients.
TextBody = "This email was sent with Amazon SES using the AWS SDK for Go."

// The character encoding for the email.
CharSet = "UTF-8"
)

func main() {
    // Create a new session in the us-west-2 region.
    // Replace us-west-2 with the AWS Region you're using for Amazon SES.
    sess, err := session.NewSession(&aws.Config{
        Region:aws.String("us-west-2")},
    )

    // Create an SES session.
    svc := ses.New(sess)

    // Assemble the email.
    input := &ses.SendEmailInput{
        Destination: &ses.Destination{
            CcAddresses: []*string{
            },
            ToAddresses: []*string{
                aws.String(Recipient),
            },
        },
        Message: &ses.Message{
            Body: &ses.Body{
                Html: &ses.Content{
                    CharSet: aws.String(CharSet),
                    Data:    aws.String(HtmlBody),
                },
                Text: &ses.Content{
                    CharSet: aws.String(CharSet),
                    Data:    aws.String(TextBody),
                },
            },
            Subject: &ses.Content{
                CharSet: aws.String(CharSet),
                Data:    aws.String(Subject),
            },
        },
        Source: aws.String(Sender),
        // Uncomment to use a configuration set
        //ConfigurationSetName: aws.String(ConfigurationSet),
    }

    // Attempt to send the email.
    result, err := svc.SendEmail(input)

    // Display error messages if they occur.
    if err != nil {
        if aerr, ok := err.(awserr.Error); ok {
```

```
        switch aerr.Code() {
        case ses.ErrCodeMessageRejected:
            fmt.Println(ses.ErrCodeMessageRejected, aerr.Error())
        case ses.ErrCodeMailFromDomainNotVerifiedException:
            fmt.Println(ses.ErrCodeMailFromDomainNotVerifiedException, aerr.Error())
        case ses.ErrCodeConfigurationSetDoesNotExistException:
            fmt.Println(ses.ErrCodeConfigurationSetDoesNotExistException, aerr.Error())
        default:
            fmt.Println(aerr.Error())
        }
    } else {
        // Print the error, cast err to awserr.Error to get the Code and
        // Message from an error.
        fmt.Println(err.Error())
    }

    return
}

fmt.Println("Email Sent to address: " + Recipient)
fmt.Println(result)
}
```

See the [complete example](#) on GitHub.

Deleting an Email Address in Amazon SES

The following example demonstrates how to use the AWS SDK for Go to delete an Amazon SES email address.

```
package main

import (
    "fmt"
    "os"

    //go get -u github.com/aws/aws-sdk-go
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ses"
)

const (
    // Replace sender@example.com with your "From" address
    Sender = "sender@example.com"

    // Replace recipient@example.com with a "To" address
    Recipient = "recipient@example.com"
)

func main() {
    // Create a new session in the us-west-2 region
    // Replace us-west-2 with the AWS Region you're using for Amazon SES
    sess, err := session.NewSession(&aws.Config{
        Region:aws.String("us-west-2")},
    )

    if err != nil {
        fmt.Println("Got error creating SES session:")
        fmt.Println(err.Error())
        os.Exit(1)
    }
}
```

```
// Create an SES session
svc := ses.New(sess)

// Remove email address
_, delErr :=
svc.DeleteVerifiedEmailAddress(&ses.DeleteVerifiedEmailAddressInput{EmailAddress:
aws.String(Recipient)})

// Display error message if it occurs
if delErr != nil {
    fmt.Println("Got error attempting to remove email address: " + Recipient)
    fmt.Println(delErr.Error())
    os.Exit(1)
}

// Display success message
fmt.Println("Removed email address: " + Recipient)
}
```

See the [complete example](#) on GitHub.

Getting Amazon SES Statistics

The following example demonstrates how to use the AWS SDK for Go to get statistics about Amazon SES. Use this information to avoid damaging your reputation when emails are bounced or rejected.

```
package main

import (
    "fmt"

    //go get -u github.com/aws/aws-sdk-go
    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/ses"
    "github.com/aws/aws-sdk-go/aws/awserr"
)

func main() {
    // Create a new session in the us-west-2 region.
    // Replace us-west-2 with the AWS Region you're using for Amazon SES.
    sess, err := session.NewSession(&aws.Config{
        Region:aws.String("us-west-2")},
    )

    // Create an SES session.
    svc := ses.New(sess)

    // Attempt to send the email.
    result, err := svc.GetSendStatistics(nil)

    // Display error messages if they occur.
    if err != nil {
        if aerr, ok := err.(awserr.Error); ok {
            switch aerr.Code() {
            case ses.ErrCodeMessageRejected:
                fmt.Println(ses.ErrCodeMessageRejected, aerr.Error())
            case ses.ErrCodeMailFromDomainNotVerifiedException:
                fmt.Println(ses.ErrCodeMailFromDomainNotVerifiedException, aerr.Error())
            case ses.ErrCodeConfigurationSetDoesNotExistException:
                fmt.Println(ses.ErrCodeConfigurationSetDoesNotExistException, aerr.Error())
            default:
                fmt.Println(aerr.Error())
            }
        }
    }
}
```



```
    }
    } else {
        // Print the error, cast err to awserr.Error to get the Code and
        // Message from an error.
        fmt.Println(err.Error())
    }

    return
}

dps := result.SendDataPoints

fmt.Println("Got", len(dps), "datapoints")
fmt.Println("")

for _, dp := range dps {
    fmt.Println("Timestamp: ", dp.Timestamp)
    fmt.Println("Attempts:  ", aws.Int64Value(dp.DeliveryAttempts))
    fmt.Println("Bounces:   ", aws.Int64Value(dp.Bounces))
    fmt.Println("Complaints:", aws.Int64Value(dp.Complaints))
    fmt.Println("Rejects:   ", aws.Int64Value(dp.Rejects))
    fmt.Println("")
}
}
```

See the [complete example](#) on GitHub.

Amazon SQS Examples Using the AWS SDK for Go

Amazon Simple Queue Service (Amazon SQS) is a fully managed message queuing service that makes it easy to decouple and scale microservices, distributed systems, and serverless applications. The AWS SDK for Go examples can integrate Amazon SQS into your applications. The examples assume you have already set up and configured the SDK (that is, you've imported all required packages and set your credentials and region). For more information, see [Getting Started with the AWS SDK for Go \(p. 2\)](#) and [Configuring the AWS SDK for Go \(p. 4\)](#).

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Topics

- [Using Amazon SQS Queues \(p. 156\)](#)
- [Sending and Receiving Messages in Amazon SQS \(p. 160\)](#)
- [Managing Visibility Timeout in Amazon SQS Queues \(p. 164\)](#)
- [Enabling Long Polling in Amazon SQS Queues \(p. 166\)](#)
- [Using Dead Letter Queues in Amazon SQS \(p. 170\)](#)
- [Setting Attributes on an Amazon SQS Queue \(p. 172\)](#)

Using Amazon SQS Queues

These AWS SDK for Go examples show you how to:

- List Amazon SQS queues
- Create Amazon SQS queues
- Get Amazon SQS queue URLs
- Delete Amazon SQS queues

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

These examples demonstrate how to work with Amazon SQS queues.

The code uses these methods of the Amazon SQS client class:

- [CreateQueue](#)
- [ListQueues](#)
- [GetQueueUrl](#)
- [DeleteQueue](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with using Amazon SQS. To learn more, see [How Queues Work](#) in the *Amazon SQS Developer Guide*.

List Queues

Create a new Go file named `sqs_listqueues.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a SQS service client.
    svc := sqs.New(sess)
```

Call `ListQueues` passing in `nil` to return all queues. Print any errors or a success message and loop through the queue URLs to print them.

```
    result, err := svc.ListQueues(nil)
    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Success")
    // As these are pointers, printing them out directly would not be useful.
```

```
    for i, urls := range result.QueueUrls {
        // Avoid dereferencing a nil pointer.
        if urls == nil {
            continue
        }
        fmt.Printf("%d: %s\n", i, *urls)
    }
}
```

Create Queues

Create a new Go file named `sqs_createqueues.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a SQS service client.
    svc := sqs.New(sess)
```

Call `CreateQueue` passing in the new queue name and queue attributes. Print any errors or a success message.

```
    result, err := svc.CreateQueue(&sqs.CreateQueueInput{
        QueueName: aws.String("SQS_QUEUE_NAME"),
        Attributes: map[string]*string{
            "DelaySeconds":      aws.String("60"),
            "MessageRetentionPeriod": aws.String("86400"),
        },
    })
    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Success", *result.QueueUrl)
}
```

Get a Queue URL

Create a new Go file named `sqs_getqueueurl.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main
```

```
import (  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/sqs"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a SQS service client.  
    svc := sqs.New(sess)
```

Call `GetQueueUrl` passing in the queue name. Print any errors or a success message.

```
    result, err := svc.GetQueueUrl(&sqs.GetQueueUrlInput{  
        QueueName: aws.String("SQS_QUEUE_NAME"),  
    })  
  
    if err != nil {  
        fmt.Println("Error", err)  
        return  
    }  
  
    fmt.Println("Success", *result.QueueUrl)  
}
```

Delete a Queue

Create a new Go file named `sqs_deletequeue.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main  
  
import (  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/sqs"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {  
    sess, err := session.NewSession(&aws.Config{  
        Region: aws.String("us-west-2")},  
    )  
  
    // Create a SQS service client.  
    svc := sqs.New(sess)
```

Call `DeleteQueue` passing in the queue name. Print any errors or a success message.

```
result, err := svc.DeleteQueue(&sqs.DeleteQueueInput{
    QueueUrl: aws.String("SQS_QUEUE_URL"),
})

if err != nil {
    fmt.Println("Error", err)
    return
}

fmt.Println("Success", result)
}
```

Sending and Receiving Messages in Amazon SQS

These AWS SDK for Go examples show you how to:

- Send a message to an Amazon SQS queue
- Receive and delete a message from an Amazon SQS queue
- Send and receive messages from an Amazon SQS queue

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

The Scenario

These examples demonstrate sending, receiving, and deleting messages from an Amazon SQS queue.

The code uses these methods of the Amazon SQS client class:

- [SendMessage](#)
- [ReceiveMessage](#)
- [DeleteMessage](#)
- [GetQueueUrl](#)

Prerequisites

- You have [set up](#) (p. 2) and [configured](#) (p. 4) the AWS SDK for Go.
- You are familiar with the details of Amazon SQS messages. To learn more, see [Sending a Message to an Amazon SQS Queue](#) and [Receiving and Deleting a Message from an Amazon SQS Queue](#) in the *Amazon SQS Developer Guide*.

Send a Message to a Queue

Create a new Go file named `sqs_sendmessage.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"
```

```
"github.com/aws/aws-sdk-go/aws"  
"github.com/aws/aws-sdk-go/aws/session"  
"github.com/aws/aws-sdk-go/service/sqs"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {  
    sess := session.Must(session.NewSessionWithOptions(session.Options{  
        SharedConfigState: session.SharedConfigEnable,  
    }))  
  
    svc := sqs.New(sess)  
  
    // URL to our queue  
    qURL := "QueueURL"
```

Now you're ready to send your message. In the example, the message input passed to `SendMessage` represents information about a fiction best seller for a particular week and defines title, author, and weeks on the list values.

```
result, err := svc.SendMessage(&sqs.SendMessageInput{  
    DelaySeconds: aws.Int64(10),  
    MessageAttributes: map[string]*sqs.MessageAttributeValue{  
        "Title": &sqs.MessageAttributeValue{  
            DataType:    aws.String("String"),  
            StringValue: aws.String("The Whistler"),  
        },  
        "Author": &sqs.MessageAttributeValue{  
            DataType:    aws.String("String"),  
            StringValue: aws.String("John Grisham"),  
        },  
        "WeeksOn": &sqs.MessageAttributeValue{  
            DataType:    aws.String("Number"),  
            StringValue: aws.String("6"),  
        },  
    },  
    MessageBody: aws.String("Information about current NY Times fiction bestseller for  
week of 12/11/2016."),  
    QueueUrl:    &qURL,  
})  
  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}  
  
fmt.Println("Success", *result.MessageId)  
}
```

Receive and Delete a Message from a Queue

Create a new Go file named `sqs_deletemessage.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main
```

```
import (  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/sqs"  
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {  
    sess := session.Must(session.NewSessionWithOptions(session.Options{  
        SharedConfigState: session.SharedConfigEnable,  
    }))  
  
    svc := sqs.New(sess)
```

Now you're ready to receive a message from a queue specified by a queue URL. In the example, the `qURL` variable would hold the URL for the queue containing the message.

```
qURL := "QueueURL"  
  
result, err := svc.ReceiveMessage(&sqs.ReceiveMessageInput{  
    AttributeNames: []*string{  
        aws.String(sqs.MessageSystemAttributeNameSentTimestamp),  
    },  
    MessageAttributeNames: []*string{  
        aws.String(sqs.QueueAttributeNameAll),  
    },  
    QueueUrl:           &qURL,  
    MaxNumberOfMessages: aws.Int64(1),  
    VisibilityTimeout:   aws.Int64(36000), // 10 hours  
    WaitTimeSeconds:     aws.Int64(0),  
})  
  
if err != nil {  
    fmt.Println("Error", err)  
    return  
}  
  
if len(result.Messages) == 0 {  
    fmt.Println("Received no messages")  
    return  
}
```

After retrieving the message, delete it from the queue with `DeleteMessage`, passing the `ReceiptHandle` returned from the previous call.

```
resultDelete, err := svc.DeleteMessage(&sqs.DeleteMessageInput{  
    QueueUrl:      &qURL,  
    ReceiptHandle: result.Messages[0].ReceiptHandle,  
})  
  
if err != nil {  
    fmt.Println("Delete Error", err)  
    return  
}  
  
fmt.Println("Message Deleted", resultDelete)  
}
```

Send and Receive Messages

Create a new Go file named `sqs_longpolling_receive_message.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "flag"
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awsserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

Get the queue name and timeout passed from the command.

```
func main() {
    var name string
    var timeout int64
    flag.StringVar(&name, "n", "", "Queue name")
    flag.Int64Var(&timeout, "t", 20, "(Optional) Timeout in seconds for long polling")
    flag.Parse()

    if len(name) == 0 {
        flag.PrintDefaults()
        exitErrorf("Queue name required")
    }
}
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create a SQS service client.
svc := sqs.New(sess)
```

Get the Queue. You need to convert the queue name into a URL. You can use the `GetQueueUrl` API call to retrieve the URL. This is needed for receiving messages from the queue. Print any errors.

```
resultURL, err := svc.GetQueueUrl(&sqs.GetQueueUrlInput{
    QueueName: aws.String(name),
})
if err != nil {
    if aerr, ok := err.(awsserr.Error); ok && aerr.Code() ==
sqs.ErrCodeQueueDoesNotExist {
        exitErrorf("Unable to find queue %q.", name)
    }
    exitErrorf("Unable to queue %q, %v.", name, err)
}
```

Call `ReceiveMessage` to get the latest message from the queue.

```
result, err := svc.ReceiveMessage(&sqs.ReceiveMessageInput{
```



```
    QueueUrl: resultURL.QueueUrl,
    AttributeNames: aws.StringSlice([]string{
        "SentTimestamp",
    }),
    MaxNumberOfMessages: aws.Int64(1),
    MessageAttributeNames: aws.StringSlice([]string{
        "All",
    }),
    WaitTimeSeconds: aws.Int64(timeout),
})
if err != nil {
    exitErrorf("Unable to receive message from queue %q, %v.", name, err)
}

fmt.Printf("Received %d messages.\n", len(result.Messages))
if len(result.Messages) > 0 {
    fmt.Println(result.Messages)
}
}
```

The example uses this utility function.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Managing Visibility Timeout in Amazon SQS Queues

This AWS SDK for Go example shows you how to:

- Change visibility timeout with Amazon SQS queues

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

This example manages visibility timeout with Amazon SQS queues. It uses these methods of the Amazon SQS client class:

- [CreateQueue](#)
- [ListQueues](#)
- [GetQueueUrl](#)
- [DeleteQueue](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with using Amazon SQS visibility timeout. To learn more, see [Visibility Timeout](#) in the *Amazon SQS Developer Guide*.

Change the Visibility Timeout

Create a new Go file named `sqs_changingvisibility.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a SQS service client.
    svc := sqs.New(sess)
```

Get a message from the queue. Call `ReceiveMessage`. Pass in the URL of the queue to return details of the next message in the queue. Print any errors, or a message if no message was received.

```
qURL := "QueueURL"

result, err := svc.ReceiveMessage(&sqs.ReceiveMessageInput{
    AttributeNames: []*string{
        aws.String(sqs.MessageSystemAttributeNameSentTimestamp),
    },
    MaxNumberOfMessages: aws.Int64(1),
    MessageAttributeNames: []*string{
        aws.String(sqs.QueueAttributeNameAll),
    },
    QueueUrl: &qURL,
})

if err != nil {
    fmt.Println("Error", err)
    return
}

// Check if we have any messages
if len(result.Messages) == 0 {
    fmt.Println("Received no messages")
    return
}
```

If a message was returned, use its receipt handle to set the timeout to 10 hours.

```
duration := int64(36000)
resultVisibility, err := svc.ChangeMessageVisibility(&sqs.ChangeMessageVisibilityInput{
    ReceiptHandle: result.Messages[0].ReceiptHandle,
    QueueUrl:      &qURL,
    VisibilityTimeout: &duration,
})

if err != nil {
    fmt.Println("Visibility Error", err)
```

```
    return
}

fmt.Println("Time Changed", resultVisibility)
}
```

Enabling Long Polling in Amazon SQS Queues

These AWS SDK for Go examples show you how to:

- Enable long polling when you create an Amazon SQS queue
- Enable long polling on an existing Amazon SQS queue
- Enable long polling when a message is received

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

Long polling reduces the number of empty responses by allowing Amazon SQS to wait a specified time for a message to become available in the queue before sending a response. Also, long polling eliminates false empty responses by querying all of the servers instead of a sampling of servers. To enable long polling, you must specify a non-zero wait time for received messages. You can do this by setting the `ReceiveMessageWaitTimeSeconds` parameter of a queue or by setting the `WaitTimeSeconds` parameter on a message when it is received.

The code uses these methods of the Amazon SQS client class:

- [SetQueueAttributes](#)
- [ReceiveMessage](#)
- [CreateQueue](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with Amazon SQS polling. To learn more, see [Long Polling](#) in the *Amazon SQS Developer Guide*.

Enable Long Polling When Creating a Queue

This example creates a queue with long polling enabled. If the queue already exists, no error is returned.

Create a new Go file named `sqs_longpolling_create_queue.go`. You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "flag"
    "fmt"
    "os"
    "strconv"

    "github.com/aws/aws-sdk-go/aws"
```

```
"github.com/aws/aws-sdk-go/aws/session"
"github.com/aws/aws-sdk-go/service/sqs"
)
```

Get the queue name passed in by the user.

```
func main() {
    var name string
    var timeout int
    flag.StringVar(&name, "n", "", "Queue name")
    flag.IntVar(&timeout, "t", 20, "(Optional) Timeout in seconds for long polling")
    flag.Parse()

    if len(name) == 0 {
        flag.PrintDefaults()
        exitErrorf("Queue name required")
    }
}
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create a SQS service client.
svc := sqs.New(sess)
```

Create the queue with long polling enabled. Print any errors or a success message.

```
result, err := svc.CreateQueue(&sqs.CreateQueueInput{
    QueueName: aws.String(name),
    Attributes: aws.StringMap(map[string]string{
        "ReceiveMessageWaitTimeSeconds": strconv.Itoa(timeout),
    }),
})
if err != nil {
    exitErrorf("Unable to create queue %q, %v.", name, err)
}

fmt.Printf("Succesffully created queue %q. URL: %s\n", name,
    aws.StringValue(result.QueueUrl))
}
```

The example uses this utility function.

```
func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Enable Long Polling on an Existing Queue

Create a new Go file named `sqs_longpolling_existing_queue.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main
```

```
import (  
    "flag"  
    "fmt"  
    "os"  
    "strconv"  
  
    "github.com/aws/aws-sdk-go/aws"  
    "github.com/aws/aws-sdk-go/aws/awsserr"  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/sqs"  
)
```

This example takes two flags, the -n flag is the queue name, and the -t flag contains the timeout value.

```
func main() {  
    var name string  
    var timeout int  
    flag.StringVar(&name, "n", "", "Queue name")  
    flag.IntVar(&timeout, "t", 20, "(Optional) Timeout in seconds for long polling")  
    flag.Parse()  
  
    if len(name) == 0 {  
        flag.PrintDefaults()  
        exitErrorf("Queue name required")  
    }  
}
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, ~/.aws/credentials.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create a SQS service client.  
svc := sqs.New(sess)
```

You need to convert the queue name into a URL. Make the `GetQueueUrl` API call to retrieve the URL. This is needed for setting attributes on the queue.

```
resultURL, err := svc.GetQueueUrl(&sqs.GetQueueUrlInput{  
    QueueName: aws.String(name),  
})  
if err != nil {  
    if aerr, ok := err.(awsserr.Error); ok && aerr.Code() ==  
sqs.ErrCodeQueueDoesNotExist {  
        exitErrorf("Unable to find queue %q.", name)  
    }  
    exitErrorf("Unable to get queue %q, %v.", name, err)  
}
```

Update the queue to enable long polling with a call to `SetQueueAttributes`, passing in the queue URL. Print any errors or a success message.

```
_, err = svc.SetQueueAttributes(&sqs.SetQueueAttributesInput{  
    QueueUrl: resultURL.QueueUrl,  
    Attributes: aws.StringMap(map[string]string{  
        "ReceiveMessageWaitTimeSeconds": strconv.Itoa(timeout),  
    }),  
})  
if err != nil {
```

```
        exitErrorf("Unable to update queue %q, %v.", name, err)
    }

    fmt.Printf("Succesffuly updated queue %q.\n", name)
}

func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Enable Long Polling on Message Receipt

Create a new Go file named `sqs_longpolling_receive_message.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "flag"
    "fmt"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/awsserr"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

This example takes two flags, the `-n` flag is the queue name, and the `-t` flag contains the timeout value.

```
func main() {
    var name string
    var timeout int64
    flag.StringVar(&name, "n", "", "Queue name")
    flag.Int64Var(&timeout, "t", 20, "(Optional) Timeout in seconds for long polling")
    flag.Parse()

    if len(name) == 0 {
        flag.PrintDefaults()
        exitErrorf("Queue name required")
    }
}
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
sess, err := session.NewSession(&aws.Config{
    Region: aws.String("us-west-2")},
)

// Create a SQS service client.
svc := sqs.New(sess)
```

You need to convert the queue name into a URL. Make the `GetQueueUrl` API call to retrieve the URL. This is needed for setting attributes on the queue.

```
resultURL, err := svc.GetQueueUrl(&sqs.GetQueueUrlInput{
    QueueName: aws.String(name),
})
```

```
if err != nil {
    if aerr, ok := err.(awserr.Error); ok && aerr.Code() ==
sqs.ErrCodeQueueDoesNotExist {
        exitErrorf("Unable to find queue %q.", name)
    }
    exitErrorf("Unable to queue %q, %v.", name, err)
}
```

Receive a message from the queue with long polling enabled with a call to `ReceiveMessage`, passing in the queue URL. Print any errors or a success message.

```
result, err := svc.ReceiveMessage(&sqs.ReceiveMessageInput{
    QueueUrl: resultURL.QueueUrl,
    AttributeNames: aws.StringSlice([]string{
        "SentTimestamp",
    }),
    MaxNumberOfMessages: aws.Int64(1),
    MessageAttributeNames: aws.StringSlice([]string{
        "All",
    }),
    WaitTimeSeconds: aws.Int64(timeout),
})
if err != nil {
    exitErrorf("Unable to receive message from queue %q, %v.", name, err)
}

fmt.Printf("Received %d messages.\n", len(result.Messages))
if len(result.Messages) > 0 {
    fmt.Println(result.Messages)
}
}

func exitErrorf(msg string, args ...interface{}) {
    fmt.Fprintf(os.Stderr, msg+"\n", args...)
    os.Exit(1)
}
```

Using Dead Letter Queues in Amazon SQS

This AWS SDK for Go example shows you how to configure source Amazon SQS queues that send messages to a dead letter queue.

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

A dead letter queue is one that other (source) queues can target for messages that can't be processed successfully. You can set aside and isolate these messages in the dead letter queue to determine why their processing didn't succeed. You must individually configure each source queue that sends messages to a dead letter queue. Multiple queues can target a single dead letter queue.

The code uses this method of the Amazon SQS client class:

- [SetQueueAttributes](#)

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.

- You are familiar with Amazon SQS dead letter queues. To learn more, see [Using Amazon SQS Dead Letter Queues](#) in the *Amazon SQS Developer Guide*.

Configure Source Queues

After you create a queue to act as a dead letter queue, you must configure the other queues that route unprocessed messages to the dead letter queue. To do this, specify a redrive policy that identifies the queue to use as a dead letter queue and the maximum number of receives by individual messages before they are routed to the dead letter queue.

Create a new Go file with the name `sqs_deadletterqueue.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "encoding/json"
    "fmt"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
func main() {
    sess, err := session.NewSession(&aws.Config{
        Region: aws.String("us-west-2")},
    )

    // Create a SQS service client.
    svc := sqs.New(sess)
```

Define the redrive policy for the queue, then marshal the policy to use as input for the `SetQueueAttributes` call.

```
policy := map[string]string{
    "deadLetterTargetArn": "SQS_QUEUE_ARN",
    "maxReceiveCount":     "10",
}

b, err := json.Marshal(policy)
if err != nil {
    fmt.Println("Failed to marshal policy:", err)
    return
}
```

Set the policy on the queue.

```
result, err := svc.SetQueueAttributes(&sqs.SetQueueAttributesInput{
    QueueUrl: aws.String("SQS_QUEUE_URL"),
    Attributes: map[string]*string{
        sqs.QueueAttributeNameRedrivePolicy: aws.String(string(b)),
    },
})
```



```
    if err != nil {
        fmt.Println("Error", err)
        return
    }

    fmt.Println("Success", result)
}
```

Setting Attributes on an Amazon SQS Queue

This Go example shows you how to set attributes on an existing Amazon SQS queue.

You can download complete versions of these example files from the [aws-doc-sdk-examples](#) repository on GitHub.

Scenario

This example updates an existing Amazon SQS queue to use long polling.

Long polling reduces the number of empty responses by allowing Amazon SQS to wait a specified time for a message to become available in the queue before sending a response. Also, long polling eliminates false empty responses by querying all of the servers instead of a sampling of servers. To enable long polling, you must specify a non-zero wait time for received messages. You can do this by setting the *ReceiveMessageWaitTimeSeconds* parameter of a queue or by setting the *WaitTimeSeconds* parameter on a message when it is received.

The code uses these methods of the Amazon SQS client class:

- [GetQueueUrl](#)
- [SetQueueAttributes](#)

If you are unfamiliar with using Amazon SQS long polling, you should read [Long Polling](#) in the *Amazon SQS Developer Guide* before proceeding.

Prerequisites

- You have [set up \(p. 2\)](#) and [configured \(p. 4\)](#) the AWS SDK for Go.
- You are familiar with using Amazon SQS long polling. To learn more, see [Long Polling](#) in the *Amazon SQS Developer Guide*.

Set Attributes on Queue

Create a new Go file named `sqs_longpolling_existing_queue.go`.

You must import the relevant Go and AWS SDK for Go packages by adding the following lines.

```
package main

import (
    "flag"
    "fmt"
    "os"
    "strconv"
)
```

```
"github.com/aws/aws-sdk-go/aws"  
"github.com/aws/aws-sdk-go/aws/awsserr"  
"github.com/aws/aws-sdk-go/aws/session"  
"github.com/aws/aws-sdk-go/service/sqs"  
)
```

Get the queue name and timeout passed in by the user.

```
func main() {  
    var name string  
    var timeout int  
    flag.StringVar(&name, "n", "", "Queue name")  
    flag.IntVar(&timeout, "t", 20, "(Optional) Timeout in seconds for long polling")  
    flag.Parse()  
  
    if len(name) == 0 {  
        flag.PrintDefaults()  
        exitErrorf("Queue name required")  
    }  
}
```

Initialize a session that the SDK will use to load credentials from the shared credentials file, `~/.aws/credentials`.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create a SQS service client.  
svc := sqs.New(sess)
```

Get the queue. You need to convert the queue name into a URL. You can use the `GetQueueUrl` API call to retrieve the URL. This is needed for setting attributes on the queue. Print any errors.

```
resultURL, err := svc.GetQueueUrl(&sqs.GetQueueUrlInput{  
    QueueName: aws.String(name),  
})  
if err != nil {  
    if aerr, ok := err.(awsserr.Error); ok && aerr.Code() ==  
sqs.ErrCodeQueueDoesNotExist {  
        exitErrorf("Unable to find queue %q.", name)  
    }  
    exitErrorf("Unable to get queue %q, %v.", name, err)  
}
```

Update the queue to enable long polling.

```
_ , err = svc.SetQueueAttributes(&sqs.SetQueueAttributesInput{  
    QueueUrl: resultURL.QueueUrl,  
    Attributes: aws.StringMap(map[string]string{  
        "ReceiveMessageWaitTimeSeconds": strconv.Itoa(timeout),  
    }},  
})  
if err != nil {  
    exitErrorf("Unable to update queue %q, %v.", name, err)  
}  
  
fmt.Printf("Successfully updated queue %q.\n", name)  
}
```

The example uses this utility function.

```
func exitErrorf(msg string, args ...interface{}) {  
    fmt.Fprintf(os.Stderr, msg+"\n", args...)  
    os.Exit(1)  
}
```

Amazon WorkDocs Examples

You can use the following examples to access Amazon WorkDocs (Amazon WorkDocs) using the AWS SDK for Go. For more information about Amazon WorkDocs, see the [Amazon WorkDocs documentation](#).

You need your organization ID to use these examples. Get your organization ID from the AWS console using the following steps:

- Select the AWS Directory Service
- Select Directories

The organization ID is the Directory ID corresponding to your Amazon WorkDocs site.

Examples

Topics

- [Listing Users \(p. 174\)](#)
- [Listing User Docs \(p. 175\)](#)

Listing Users

The following example lists the names of all users, or lists additional details about a user if a user name is specified on the command line. Choose Copy to save the code locally, or see the link to the complete example at the end of this topic.

Import the following Go packages.

```
import (  
    "flag"  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/workdocs"  
)
```

- `flag` is for getting user input, in this case the name of the user
- `fmt` is for formatting output
- `session` is for creating a session
- `workdocs` is for using the WorkDocs APIs

Create a session and Amazon WorkDocs client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)
```

```
// Create a Workdocs service client.  
svc := workdocs.New(sess)
```

Create the input arguments for the `DescribeUsers` method and add your organization ID.

```
input := new(workdocs.DescribeUsersInput)  
  
// Replace with your organization ID  
org_id := "d-123456789c"  
input.OrganizationId = &org_id
```

If we have a user name, add that to the input arguments so we only get information about that user.

```
user_ptr := flag.String("u", "", "User for whom info is retrieved")  
flag.Parse()  
  
// Show all users if we don't get a user name  
if *user_ptr == "" {  
    fmt.Println("Getting info about all users")  
} else {  
    fmt.Println("Getting info about user " + *user_ptr)  
    input.Query = user_ptr  
}
```

Run the `DescribeUsers` method and display the information for the user or all users.

```
result, err := svc.DescribeUsers(input)  
  
if err != nil {  
    fmt.Println("Error getting user info", err)  
    return  
}  
  
if *user_ptr == "" {  
    fmt.Println("Found", *result.TotalNumberOfUsers, "users")  
    fmt.Println("")  
}  
  
for _, user := range result.Users {  
    fmt.Println("Username:   " + *user.Username)  
  
    if *user_ptr != "" {  
        fmt.Println("Firstname:  " + *user.GivenName)  
        fmt.Println("Lastname:   " + *user.Surname)  
        fmt.Println("Email:      " + *user.EmailAddress)  
        fmt.Println("Root folder " + *user.RootFolderId)  
    }  
  
    fmt.Println("")  
}
```

See the [complete example](#) on GitHub.

Listing User Docs

The following example lists the documents for the user whose name is specified on the command line. Choose Copy to save the code locally, or see the link to the complete example at the end of this topic.

Import the following Go packages.

```
import (  
    "flag"  
    "fmt"  
  
    "github.com/aws/aws-sdk-go/aws/session"  
    "github.com/aws/aws-sdk-go/service/workdocs"  
)
```

- `flag` is for getting user input, in this case the name of the user
- `fmt` is for formatting output
- `session` is for creating a session
- `workdocs` is for using the WorkDocs APIs

Create a session and Amazon WorkDocs client.

```
sess, err := session.NewSession(&aws.Config{  
    Region: aws.String("us-west-2")},  
)  
  
// Create a Workdocs service client.  
svc := workdocs.New(sess)
```

Check that we have a user name, and get the root folder for that user.

```
user_ptr := flag.String("u", "", "User for whom info is retrieved")  
flag.Parse()  
  
// Show all users if we don't get a user name  
if *user_ptr == "" {  
    fmt.Println("You must supply a user name")  
    return  
}  
  
// Replace with your organization ID  
org_id := "d-123456789c"  
  
input := new(workdocs.DescribeUsersInput)  
input.OrganizationId = &org_id  
input.Query = user_ptr  
  
result, err := svc.DescribeUsers(input)  
  
if err != nil {  
    fmt.Println("Error getting user info", err)  
    return  
}  
  
var folder_id = ""  
  
if *result.TotalNumberOfUsers == 1 {  
    for _, user := range result.Users {  
        folder_id = *user.RootFolderId  
    }  
}
```

Run the `DescribeFolderContents` method and display the name, size, and last modified information for each document.

```
result, err := svc.DescribeFolderContents(&workdocs.DescribeFolderContentsInput{FolderId:  
    &folder_id})
```

```
if err != nil {
    fmt.Println("Error getting docs for user", err)
    return
}

fmt.Println(*user_ptr + " docs:")
fmt.Println("")

for _, doc := range result.Documents {
    fmt.Println(*doc.LatestVersionMetadata.Name)
    fmt.Println("  Size:           ", *doc.LatestVersionMetadata.Size, "(bytes)")
    fmt.Println("  Last modified:", *doc.LatestVersionMetadata.ModifiedTimestamp)
    fmt.Println("")
}
```

See the [complete example](#) on GitHub.

Using the AWS SDK for Go Utilities

The AWS SDK for Go includes the following utilities to help you more easily use AWS services. Find the SDK utilities in their related AWS service package.

Amazon CloudFront URL Signer

The Amazon CloudFront URL signer simplifies the process of creating signed URLs. A signed URL includes information, such as an expiration date and time, that enables you to control access to your content. Signed URLs are useful when you want to distribute content through the internet, but want to restrict access to certain users (for example, to users who have paid a fee).

To sign a URL, create a `URLSigner` instance with your CloudFront key pair ID and the associated private key. Then call the `Sign` or `SignWithPolicy` method and include the URL to sign. For more information about Amazon CloudFront key pairs, see [Creating CloudFront Key Pairs for Your Trusted Signers](#) in the *Amazon CloudFront Developer Guide*.

The following example creates a signed URL that's valid for one hour after it is created.

```
signer := sign.NewURLSigner(keyID, privKey)

signedURL, err := signer.Sign(rawURL, time.Now().Add(1*time.Hour))
if err != nil {
    log.Fatalf("Failed to sign url, err: %s\n", err.Error())
    return
}
```

For more information about the signing utility, see the [sign](#) package in the *AWS SDK for Go API Reference*.

Amazon DynamoDB Attributes Converter

The attributes converter simplifies converting Amazon DynamoDB attribute values to and from concrete Go types. Conversions make it easy to work with attribute values in Go and to write values to Amazon DynamoDB tables. For example, you can create records in Go and then use the converter when you want to write those records as attribute values to a DynamoDB table.

The following example converts a structure to an `AmazonDynamoDBAttributeValues` map and then puts the data to the `exampleTable`.

```
type Record struct {
    MyField string
    Letters []string
    A2Num    map[string]int
}

r := Record{
    MyField: "dynamodbattribute.ConvertToX example",
    Letters: []string{"a", "b", "c", "d"},
}
```

```
A2Num:    map[string]int{"a": 1, "b": 2, "c": 3},
}

//...

svc := dynamodb.New(session.New(&aws.Config{Region: aws.String("us-west-2")}))
item, err := dynamodbattribute.ConvertToMap(r)
if err != nil {
    fmt.Println("Failed to convert", err)
    return
}
result, err := svc.PutItem(&dynamodb.PutItemInput{
    Item:      item,
    TableName: aws.String("exampleTable"),
})
fmt.Println("Item put to dynamodb", result, err)
```

For more information about the converter utility, see the [dynamodbattribute](#) package in the *AWS SDK for Go API Reference*.

Amazon Elastic Compute Cloud Metadata

EC2Metadata is a client that interacts with the Amazon EC2 metadata service. The client can help you easily retrieve information about instances on which your applications run, such as its region or local IP address. Typically, you must create and submit HTTP requests to retrieve instance metadata. Instead, create an EC2Metadata service client.

```
c := ec2metadata.New(session.New())
```

Then use the service client to retrieve information from a metadata category like `local-ipv4` (the private IP address of the instance).

```
localip, err := c.GetMetadata("local-ipv4")
if err != nil {
    log.Printf("Unable to retrieve the private IP address from the EC2 instance: %s\n",
        err)
    return
}
```

For a list of all metadata categories, see [Instance Metadata Categories](#) in the *Amazon EC2 User Guide for Linux Instances*.

Retrieving an Instance's Region

There's no instance metadata category that returns only the region of an instance. Instead, use the included `Region` method to easily return an instance's region.

```
region, err := ec2metadata.New(session.New()).Region()
if err != nil {
    log.Printf("Unable to retrieve the region from the EC2 instance %v\n", err)
}
```

For more information about the EC2 metadata utility, see the [ec2metadata](#) package in the *AWS SDK for Go API Reference*.

Amazon S3 Transfer Managers

The Amazon Simple Storage Service upload and download managers can break up large objects so they can be transferred in multiple parts, in parallel. This makes it easy to resume interrupted transfers.

Upload Manager

The Amazon Simple Storage Service upload manager determines if a file can be split into smaller parts and uploaded in parallel. You can customize the number of parallel uploads and the size of the uploaded parts.

Example: Uploading a File

The following example uses the `AmazonS3Uploader` to upload a file. Using `Uploader` is similar to the `s3.PutObject()` operation.

```
mySession, _ := session.NewSession()
uploader := s3manager.NewUploader(mySession)
result, err := uploader.Upload(&s3manager.UploadInput{
    Bucket: &uploadBucket,
    Key:    &uploadFileKey,
    Body:   uploadFile,
})
```

Configuration Options

When you instantiate an `Uploader` instance, you can specify several configuration options (`UploadOptions`) to customize how objects are uploaded:

- `PartSize`– Specifies the buffer size, in bytes, of each part to upload. The minimum size per part is 5 MB.
- `Concurrency`– Specifies the number of parts to upload in parallel.
- `LeavePartsOnError`– Indicates whether to leave successfully uploaded parts in Amazon S3.

Tweak the `PartSize` and `Concurrency` configuration values to find the optimal configuration. For example, systems with high-bandwidth connections can send bigger parts and more uploads in parallel.

For more information about `Uploader` and its configurations, see the [s3manager](#) package in the *AWS SDK for Go API Reference*.

UploadInput Body Field (io.ReadSeeker vs. io.Reader)

The `Body` field of the `s3manager.UploadInput` struct is an `io.Reader` type. However, the field also satisfies the `io.ReadSeeker` interface.

For `io.ReadSeeker` types, the `Uploader` doesn't buffer the body contents before sending it to Amazon S3. `Uploader` calculates the expected number of parts before uploading the file to Amazon S3. If the current value of `PartSize` requires more than 10,000 parts to upload the file, `Uploader` increases the part size value so that fewer parts are required.

For `io.Reader` types, the bytes of the reader must buffer each part in memory before the part is uploaded. When you increase the `PartSize` or `Concurrency` value, the required memory (RAM) for the `Uploader` increases significantly. The required memory is approximately ```PartSize``` * ```Concurrency```. For example, if you specify 100 MB for `PartSize` and 10 for `Concurrency`, the required memory will be at least 1 GB.

Because an `io.Reader` type cannot determine its size before reading its bytes, `Uploader` cannot calculate how many parts must be uploaded. Consequently, `Uploader` can reach the Amazon S3 upload limit of 10,000 parts for large files if you set the `PartSize` too low. If you try to upload more than 10,000 parts, the upload stops and returns an error.

Handling Partial Uploads

If an upload to Amazon S3 fails, by default, `Uploader` uses the `AmazonS3AbortMultipartUpload` operation to remove the uploaded parts. This functionality ensures that failed uploads do not consume Amazon S3 storage.

You can set `LeavePartsOnError` to `true` so that the `Uploader` doesn't delete successfully uploaded parts. This is useful for resuming partially completed uploads. To operate on uploaded parts, you must get the `UploadID` of the failed upload. The following example demonstrates how to use the `s3manager.MultiUploadFailure` message to get the `UploadID`.

```
u := s3manager.NewUploader(session.New())
output, err := u.upload(input)
if err != nil {
    if multierr, ok := err.(s3manager.MultiUploadFailure); ok {
        // Process error and its associated uploadID
        fmt.Println("Error:", multierr.Code(), multierr.Message(), multierr.UploadID())
    } else {
        // Process error generically
        fmt.Println("Error:", err.Error())
    }
}
```

Example: Upload a Folder to Amazon S3

The following example uses the `path/filepath` package to recursively gather a list of files and upload them to the specified Amazon S3 bucket. The keys of the Amazon S3 objects are prefixed with the file's relative path.

```
package main

import (
    "log"
    "os"
    "path/filepath"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3/s3manager"
)

var (
    localPath string
    bucket    string
    prefix    string
)

func init() {
    if len(os.Args) != 4 {
        log.Fatalf("Usage:", os.Args[0], "<local path> <bucket> <prefix>")
    }
    localPath = os.Args[1]
    bucket = os.Args[2]
    prefix = os.Args[3]
}
```

```
func main() {
    walker := make(fileWalk)
    go func() {
        // Gather the files to upload by walking the path recursively
        if err := filepath.Walk(localPath, walker.Walk); err != nil {
            log.Fatalln("Walk failed:", err)
        }
        close(walker)
    }()

    // For each file found walking, upload it to S3
    uploader := s3manager.NewUploader(session.New())
    for path := range walker {
        rel, err := filepath.Rel(localPath, path)
        if err != nil {
            log.Fatalln("Unable to get relative path:", path, err)
        }
        file, err := os.Open(path)
        if err != nil {
            log.Println("Failed opening file", path, err)
            continue
        }
        defer file.Close()
        result, err := uploader.Upload(&s3manager.UploadInput{
            Bucket: &bucket,
            Key:    aws.String(filepath.Join(prefix, rel)),
            Body:   file,
        })
        if err != nil {
            log.Fatalln("Failed to upload", path, err)
        }
        log.Println("Uploaded", path, result.Location)
    }
}

type fileWalk chan string

func (f fileWalk) Walk(path string, info os.FileInfo, err error) error {
    if err != nil {
        return err
    }
    if !info.IsDir() {
        f <- path
    }
    return nil
}
```

Example: Upload a File to Amazon S3 and Send its Location to Amazon SQS

The following example uploads a file to an Amazon S3 bucket and then sends a notification message of the file's location to an Amazon Simple Queue Service queue.

```
package main

import (
    "log"
    "os"

    "github.com/aws/aws-sdk-go/aws"
    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3/s3manager"
    "github.com/aws/aws-sdk-go/service/sqs"
)
```

```

)

// Uploads a file to a specific bucket in S3 with the file name
// as the object's key. After it's uploaded, a message is sent
// to a queue.
func main() {
    if len(os.Args) != 4 {
        log.Fatalln("Usage:", os.Args[0], "<bucket> <queue> <file>")
    }

    file, err := os.Open(os.Args[3])
    if err != nil {
        log.Fatal("Open failed:", err)
    }
    defer file.Close()

    // Upload the file to S3 using the S3 Manager
    uploader := s3manager.NewUploader(session.New())
    uploadRes, err := uploader.Upload(&s3manager.UploadInput{
        Bucket: aws.String(os.Args[1]),
        Key:     aws.String(file.Name()),
        Body:    file,
    })
    if err != nil {
        log.Fatalln("Upload failed:", err)
    }

    // Get the URL of the queue that the message will be posted to
    svc := sqs.New(session.New())
    urlRes, err := svc.GetQueueUrl(&sqs.GetQueueUrlInput{
        QueueName: aws.String(os.Args[2]),
    })
    if err != nil {
        log.Fatalln("GetQueueURL failed:", err)
    }

    // Send the message to the queue
    _, err = svc.SendMessage(&sqs.SendMessageInput{
        MessageBody: &uploadRes.Location,
        QueueUrl:     urlRes.QueueUrl,
    })
    if err != nil {
        log.Fatalln("SendMessage failed:", err)
    }
}

```

Download Manager

The Amazon S3 download manager determines if a file can be split into smaller parts and downloaded in parallel. You can customize the number of parallel downloads and the size of the downloaded parts.

Example: Download a File

The following example uses the Amazon S3Downloader to download a file. Using Downloader is similar to the `s3.GetObject()` operation.

```

downloader := s3manager.NewDownloader(session.New())
numBytes, err := downloader.Download(downloadFile,
    &s3.GetObjectInput{
        Bucket: &downloadBucket,
        Key:     &downloadFileKey,
    })

```

The `downloadFile` parameter is an `io.WriterAt` type. The `WriterAt` interface enables the `Downloader` to write multiple parts of the file in parallel.

Configuration Options

When you instantiate a `Downloader` instance, you can specify several configuration options (`DownloadOptions`) to customize how objects are downloaded:

- `PartSize`— Specifies the buffer size, in bytes, of each part to download. The minimum size per part is 5 MB.
- `Concurrency`— Specifies the number of parts to download in parallel.

Tweak the `PartSize` and `Concurrency` configuration values to find the optimal configuration. For example, systems with high-bandwidth connections can receive bigger parts and more downloads in parallel.

For more information about `Downloader` and its configurations, see the [s3manager](#) package in the *AWS SDK for Go API Reference*.

Example: Download All Objects in a Bucket

The following example uses pagination to gather a list of objects from an Amazon S3 bucket. Then it downloads each object to a local file.

```
package main

import (
    "fmt"
    "os"
    "path/filepath"

    "github.com/aws/aws-sdk-go/aws/session"
    "github.com/aws/aws-sdk-go/service/s3"
    "github.com/aws/aws-sdk-go/service/s3/s3manager"
)

var (
    Bucket      = "MyBucket" // Download from this bucket
    Prefix      = "logs/"    // Using this key prefix
    LocalDirectory = "s3logs" // Into this directory
)

func main() {
    manager := s3manager.NewDownloader(session.New())
    d := downloader{bucket: Bucket, dir: LocalDirectory, Downloader: manager}

    client := s3.New(session.New())
    params := &s3.ListObjectsInput{Bucket: &Bucket, Prefix: &Prefix}
    client.ListObjectsPages(params, d.eachPage)
}

type downloader struct {
    *s3manager.Downloader
    bucket, dir string
}

func (d *downloader) eachPage(page *s3.ListObjectsOutput, more bool) bool {
    for _, obj := range page.Contents {
        d.downloadToFile(*obj.Key)
    }
}
```

```
    return true
}

func (d *downloader) downloadToFile(key string) {
    // Create the directories in the path
    file := filepath.Join(d.dir, key)
    if err := os.MkdirAll(filepath.Dir(file), 0775); err != nil {
        panic(err)
    }

    // Set up the local file
    fd, err := os.Create(file)
    if err != nil {
        panic(err)
    }
    defer fd.Close()

    // Download the file using the AWS SDK for Go
    fmt.Printf("Downloading s3://%s/%s to %s...\n", d.bucket, key, file)
    params := &s3.GetObjectInput{Bucket: &d.bucket, Key: &key}
    d.Download(fd, params)
}
```

Document History

This topic describes important changes to the *AWS SDK for Go Developer Guide* over the course of its history.

To view the list of changes to the AWS SDK for Go and its documentation, see the [CHANGELOG.md](#) file in the `aws/aws-sdk-go` repository in GitHub.

Last documentation update: Mar 24, 2018

Jul 10, 2017

Added the Document History topic.

June 14, 2017

Updated error handling example code.

April 14, 2017

Added S3 policy example.

April 11, 2017

Added S3 bucket ACL examples.

April 7, 2017

Added samples of getting/setting bucket and bucket object ACLs.

March 28, 2017

Updated basic Amazon S3 bucket examples.

February 27, 2017

Added an example of AssumeRoleTokenProvider.

November 2, 2016

Added SDK setters examples.

October 23, 2016

Changed deprecated `session.New` to `session.NewSession`.

September 29, 2016

Added a new topic for handling service errors from GitHub repository.

September 28, 2016

Added an example of Go extending SDK for `context.Context`.

August 31, 2016

Initial release of the *AWS SDK for Go Developer Guide*.