



# The Software Development Process at Amazon

**Jonathan Weiss**

**Managing Director Amazon Web Services Germany GmbH**

**AWS OpsWorks, AWS Resource Groups, AWS Systems Manager**

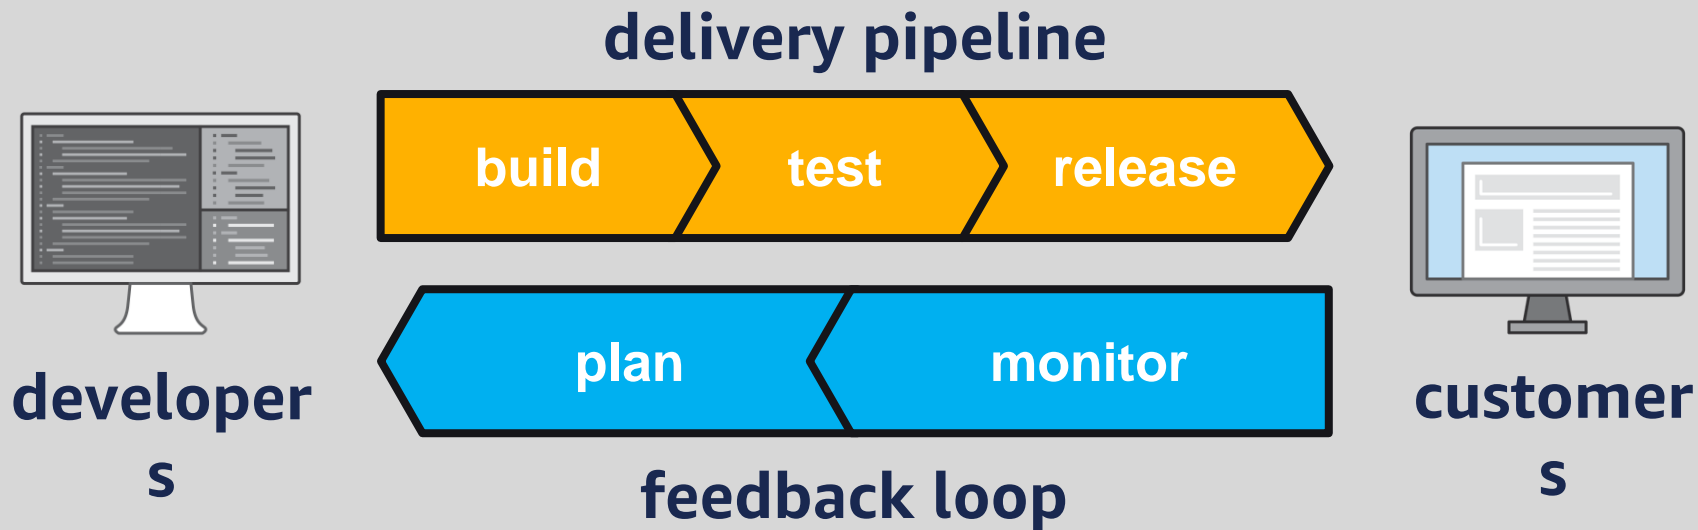


**Amazon is hundreds  
of different  
businesses**

**AWS did over 1,430  
major feature  
releases in 2017**

**No signs of  
slowing down**

# The Software Development Lifecycle



# Ten Lessons Learned

1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**

# Ten Lessons Learned

1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**

# Let's go back 15 years ago...

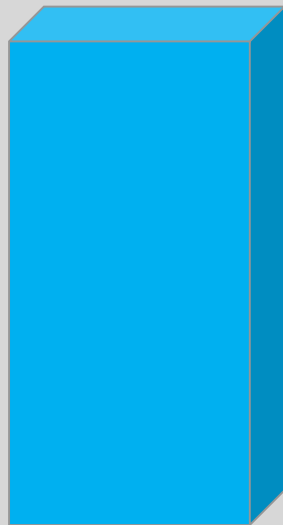


# Monolith development lifecycle



**developer**

**S**



**app**



**delivery pipeline**



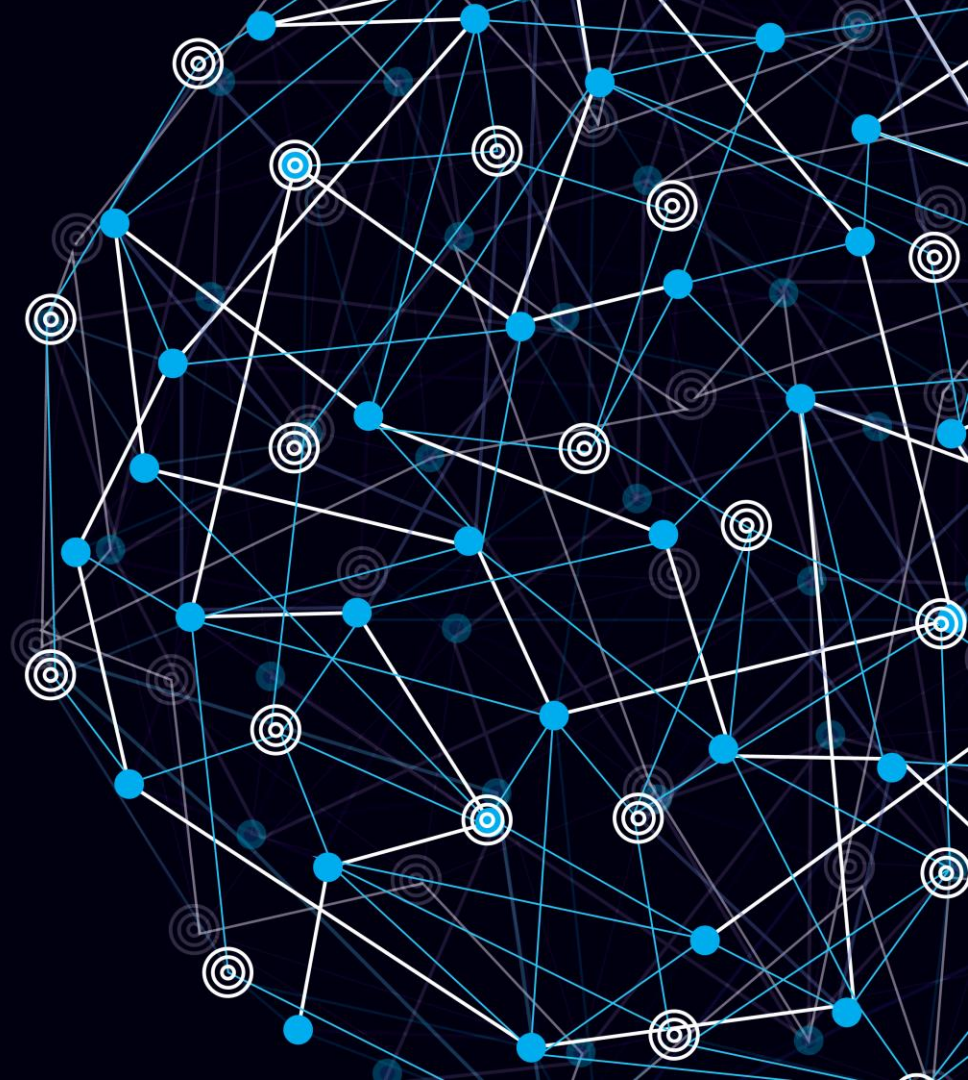
**Service-Oriented  
Architecture (SOA)**

**Single-purpose  
Primitive**

**Connected w/ APIs**

**Highly decoupled**

**“Microservices”**



# Ten Lessons Learned

1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**



**Two-pizza teams**

**Decentralized**

**Agility**

**Autonomy**

**Accountability**

**Ownership**

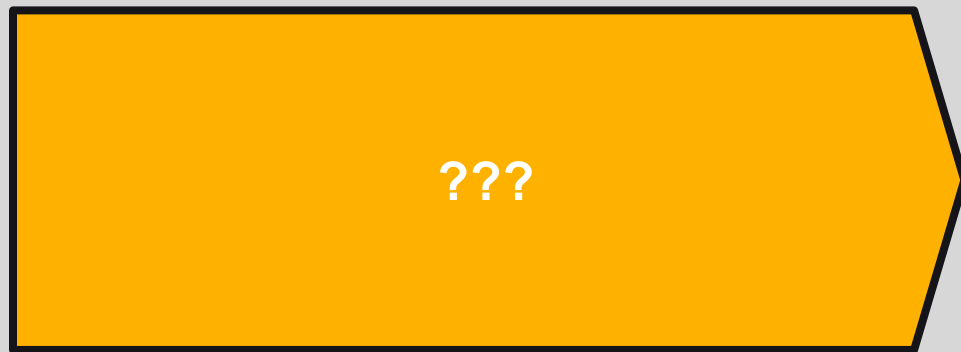
**“DevOps”**

# Missing Tools



developer

services



delivery pipeline

# Ten Lessons Learned

1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**



Self-service

Technology-agnostic

Encourage best practices

Single-purpose services







**Deployment service**

**No downtime  
deployments**

**Health tracking**

**Versioned artifacts  
and rollbacks**

# Ten Lessons Learned

1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**



# Pipelines

A large, insulated pipeline runs diagonally across the frame from the bottom left towards the top right. The pipeline is covered in a thick layer of snow. It is supported by several dark, vertical metal posts. The background is a dense forest of trees, all heavily laden with snow. The sky is a pale, hazy blue with a soft orange glow on the horizon, suggesting a sunset or sunrise. The overall scene is quiet and cold.

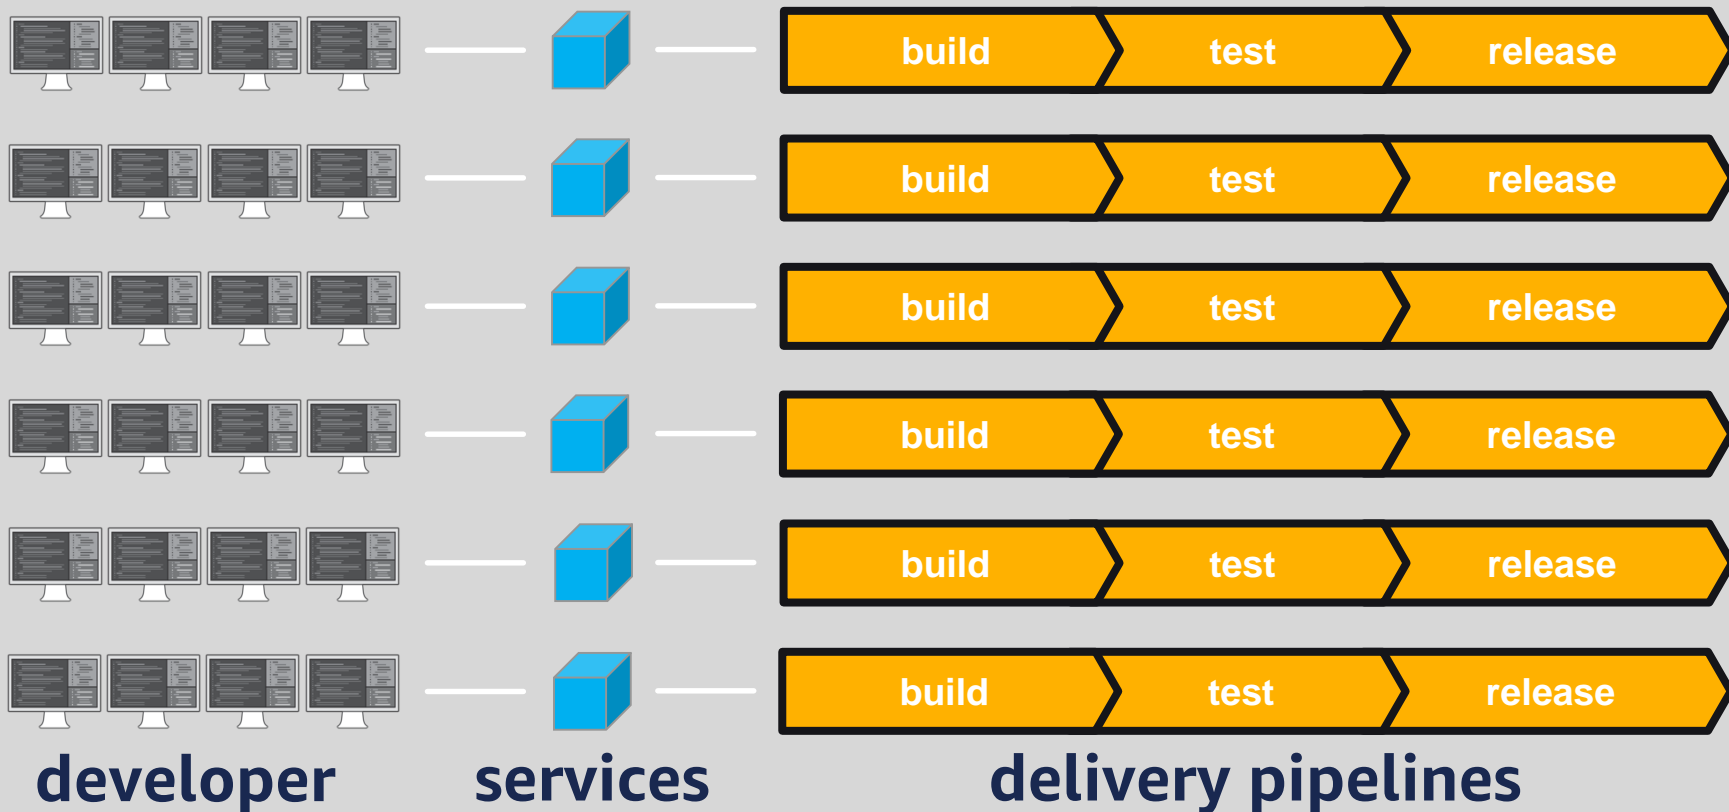
**Continuous delivery**

**Automated release process**

**Faster & more reliable releases**

**Adopted by virtually all teams**

# Microservice development lifecycle



**Thousands of teams  
x Microservice architecture  
x Continuous delivery  
x Multiple environments**

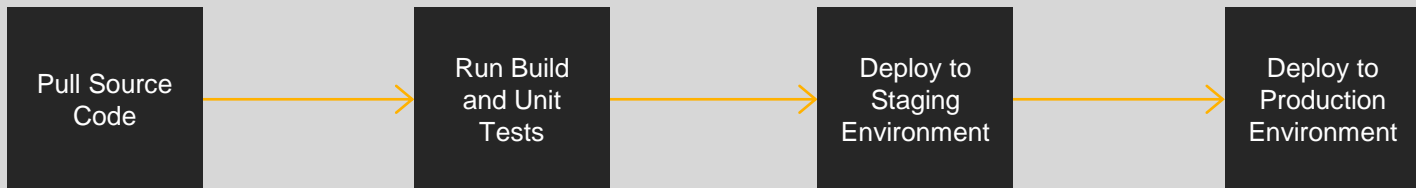
---

**= millions of deployments**

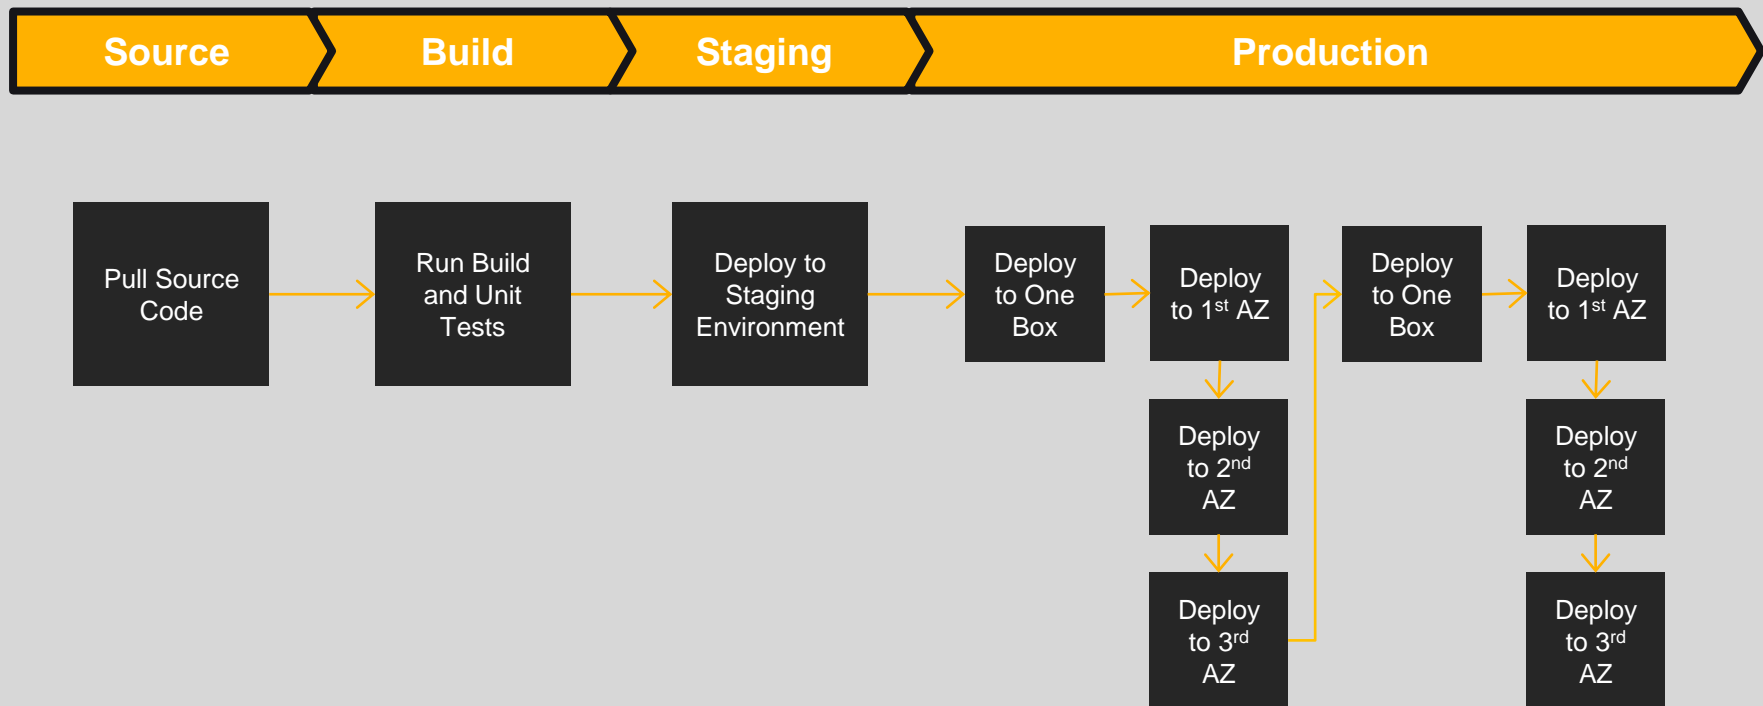
# Ten Lessons Learned

1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**

# Release Pipelines



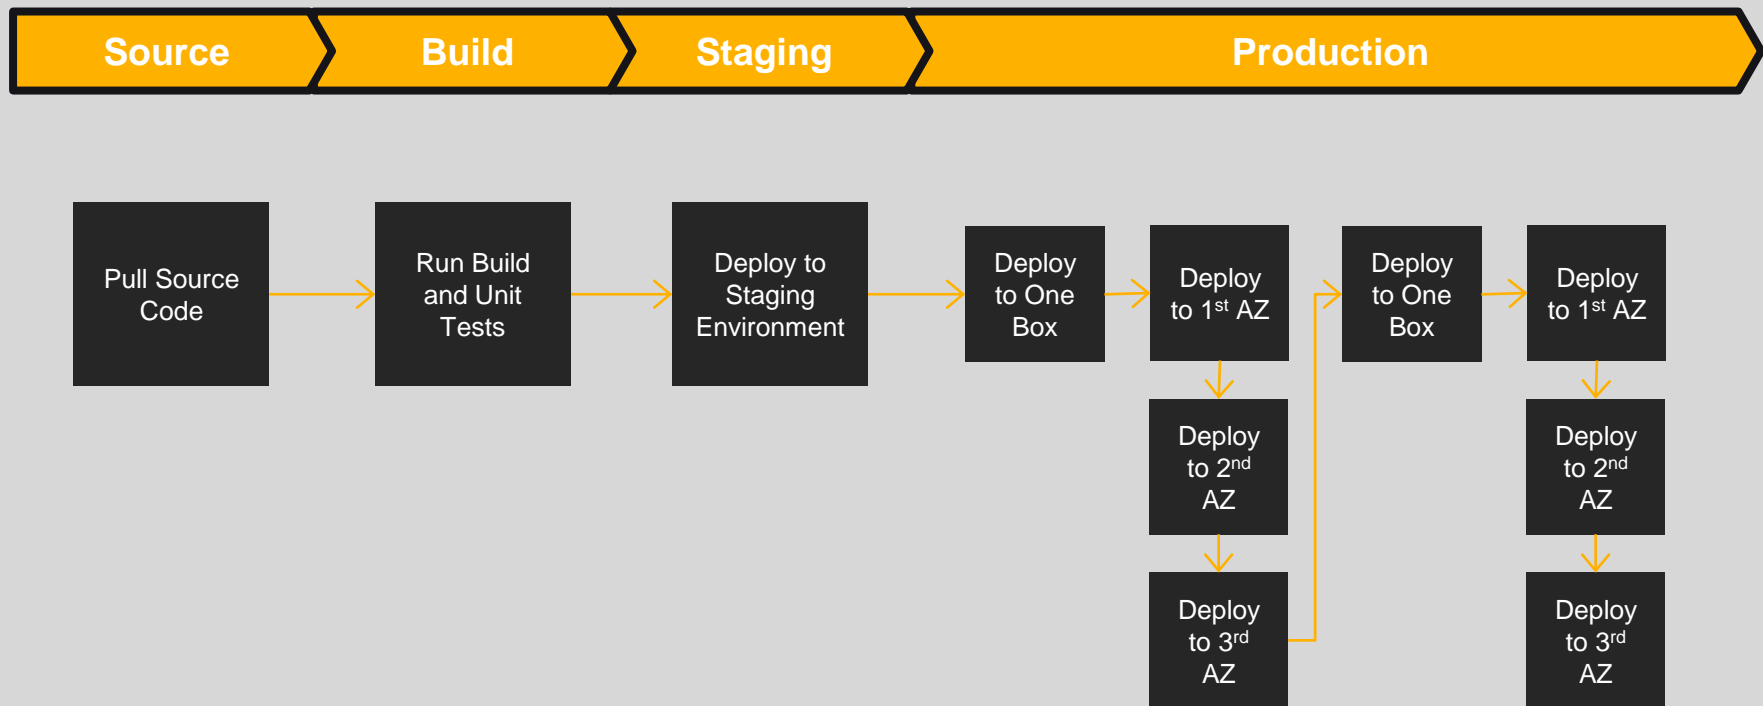
# Release Pipelines



# Ten Lessons Learned

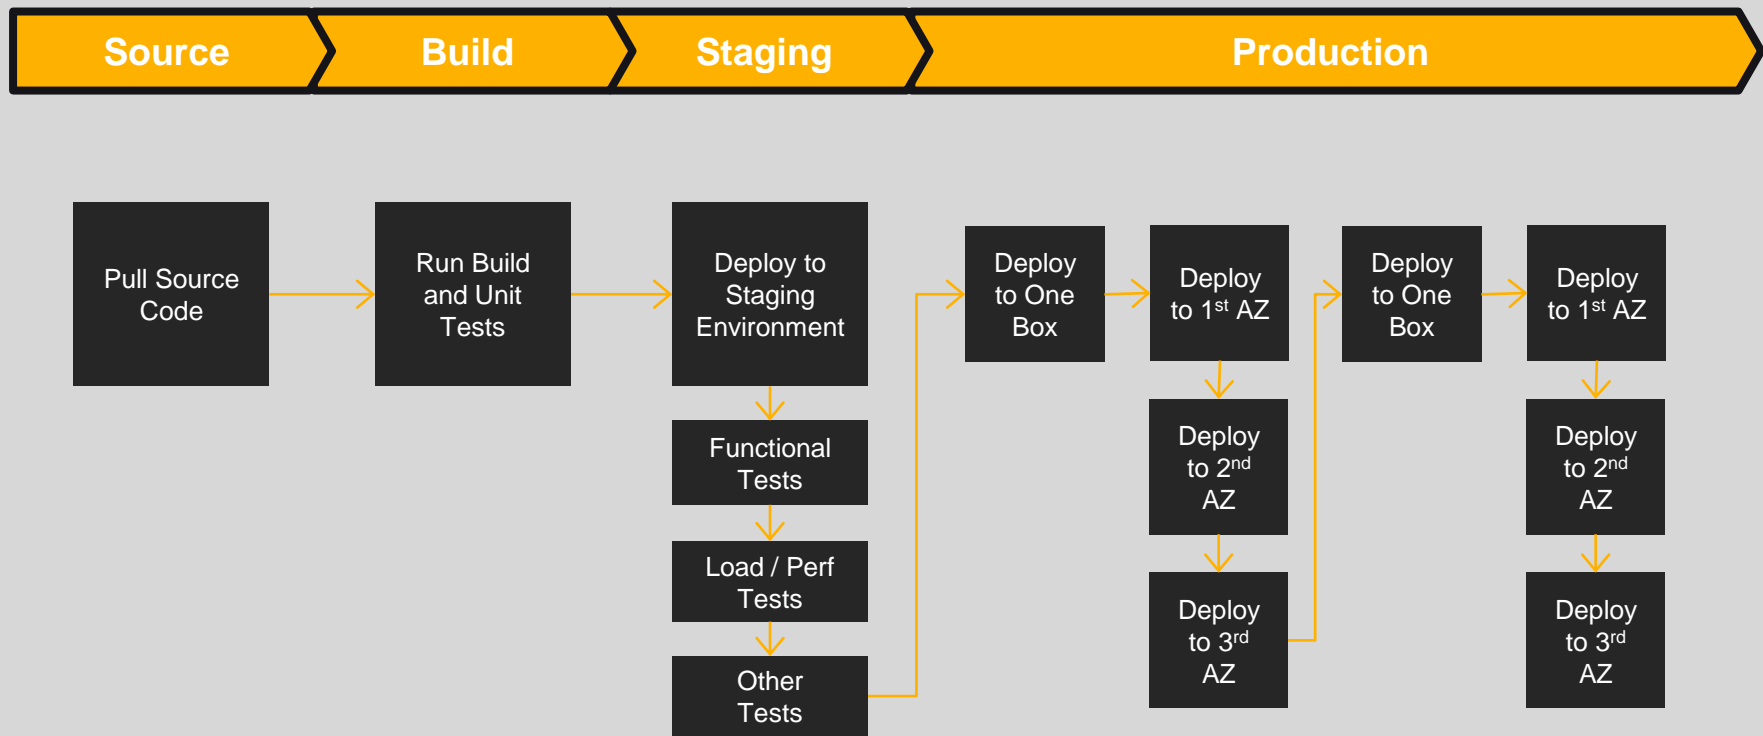
1. **Microservices**
2. **DevOps Teams**
3. **Self-Service Tools**
4. **Continuous Delivery**
5. **Pessimistic Deployments**
6. **Automated Testing**
7. **Optimize ECT Loop**
8. **Monitor Everything**
9. **Measure Everything**
10. **Listen to Customers**

# Release Pipelines

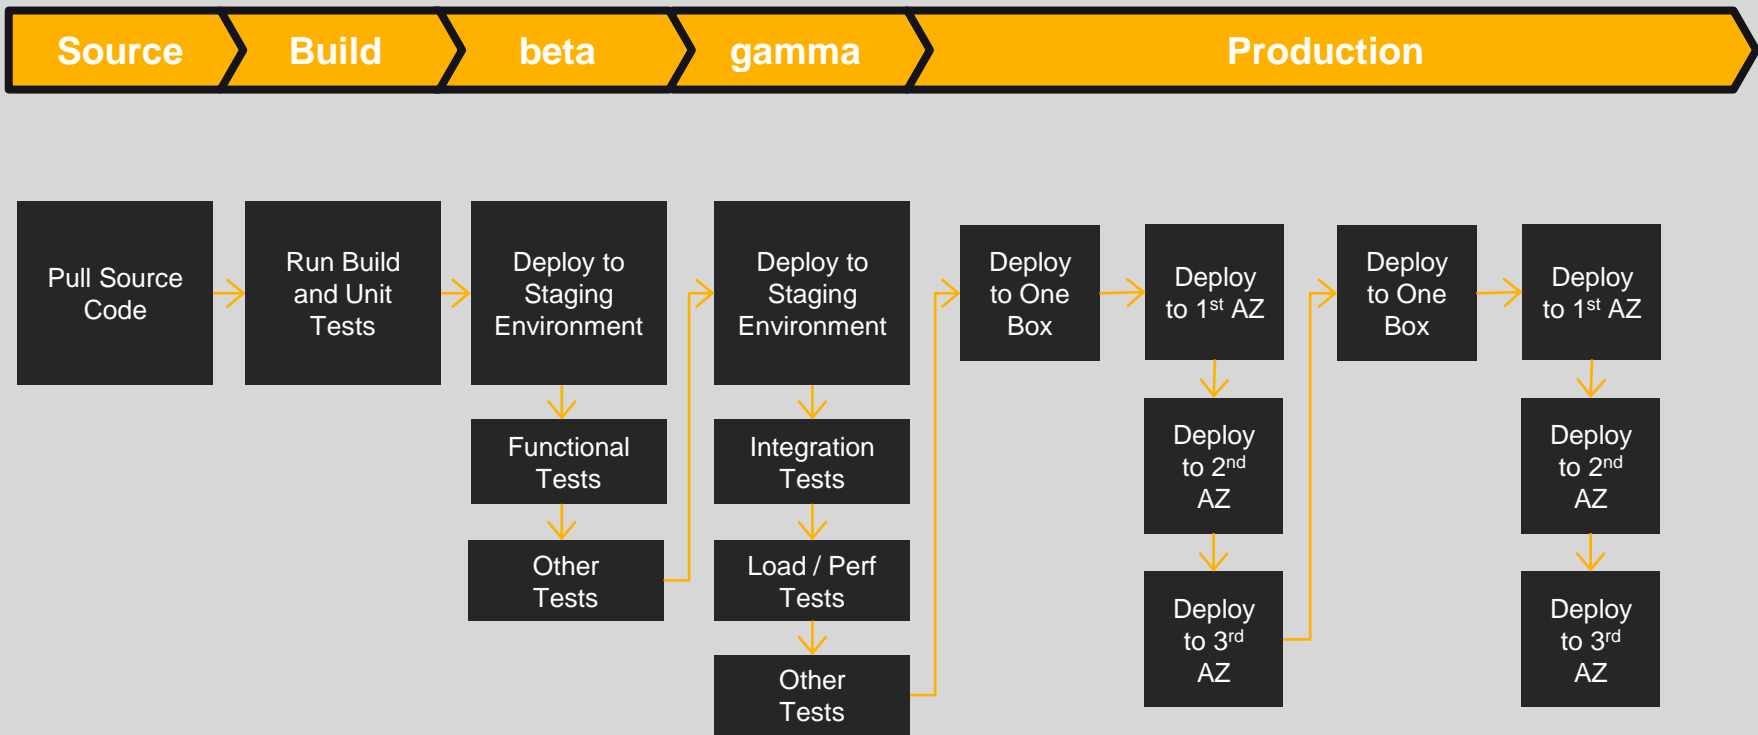




# Release Pipelines



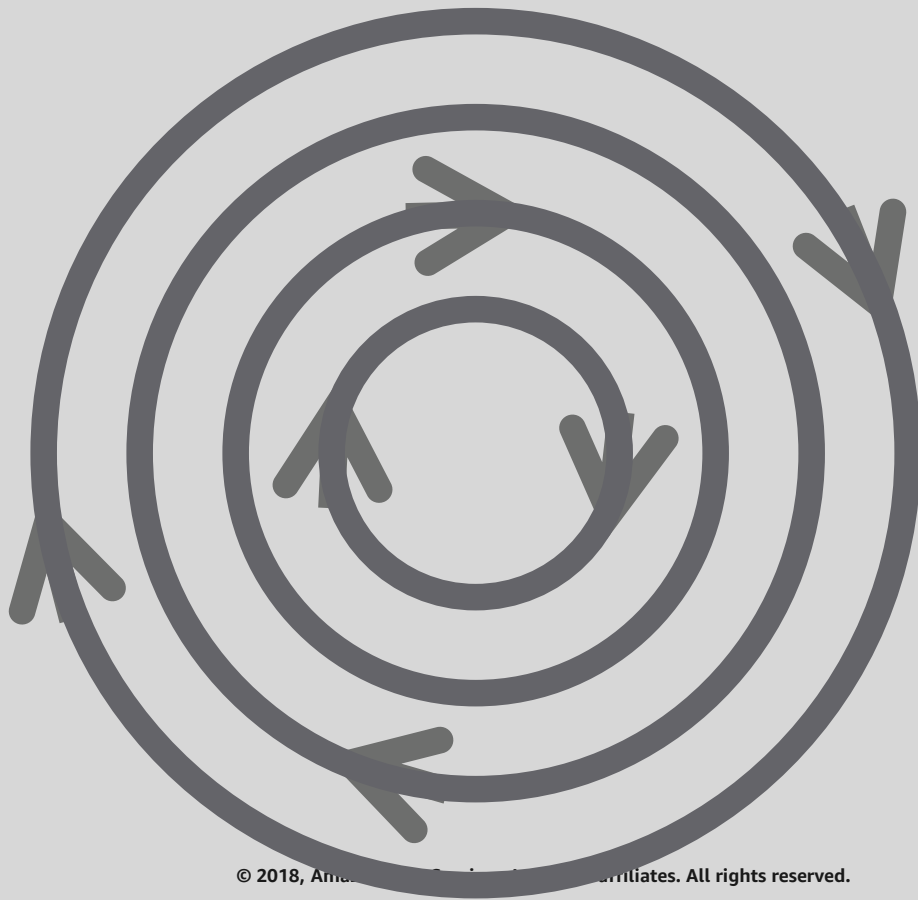
# Release Pipelines



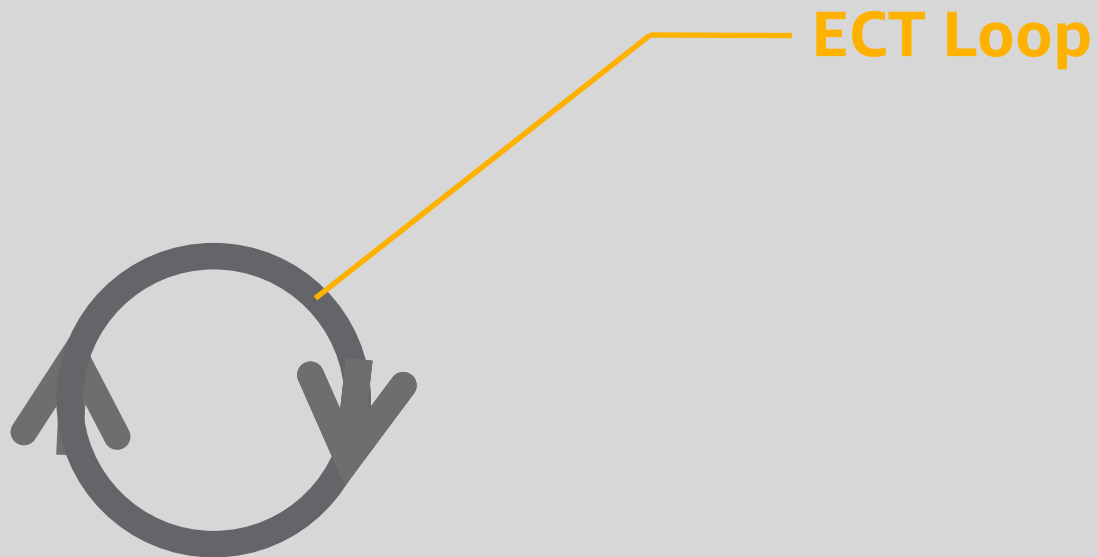
# Ten Lessons Learned

1. Microservices
2. DevOps Teams
3. Self-Service Tools
4. Continuous Delivery
5. Pessimistic Deployments
6. Automated Testing
7. Optimize ECT Loop
8. Monitor Everything
9. Measure Everything
10. Listen to Customers

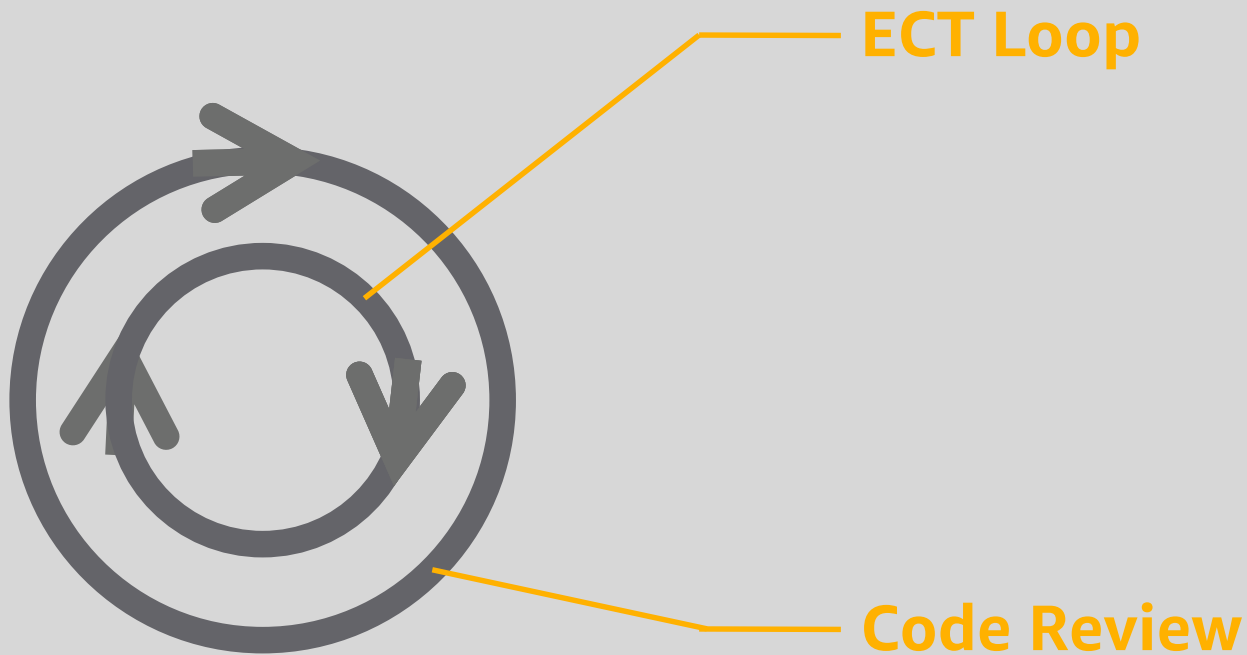
# The Edit-Compile-Test Loop



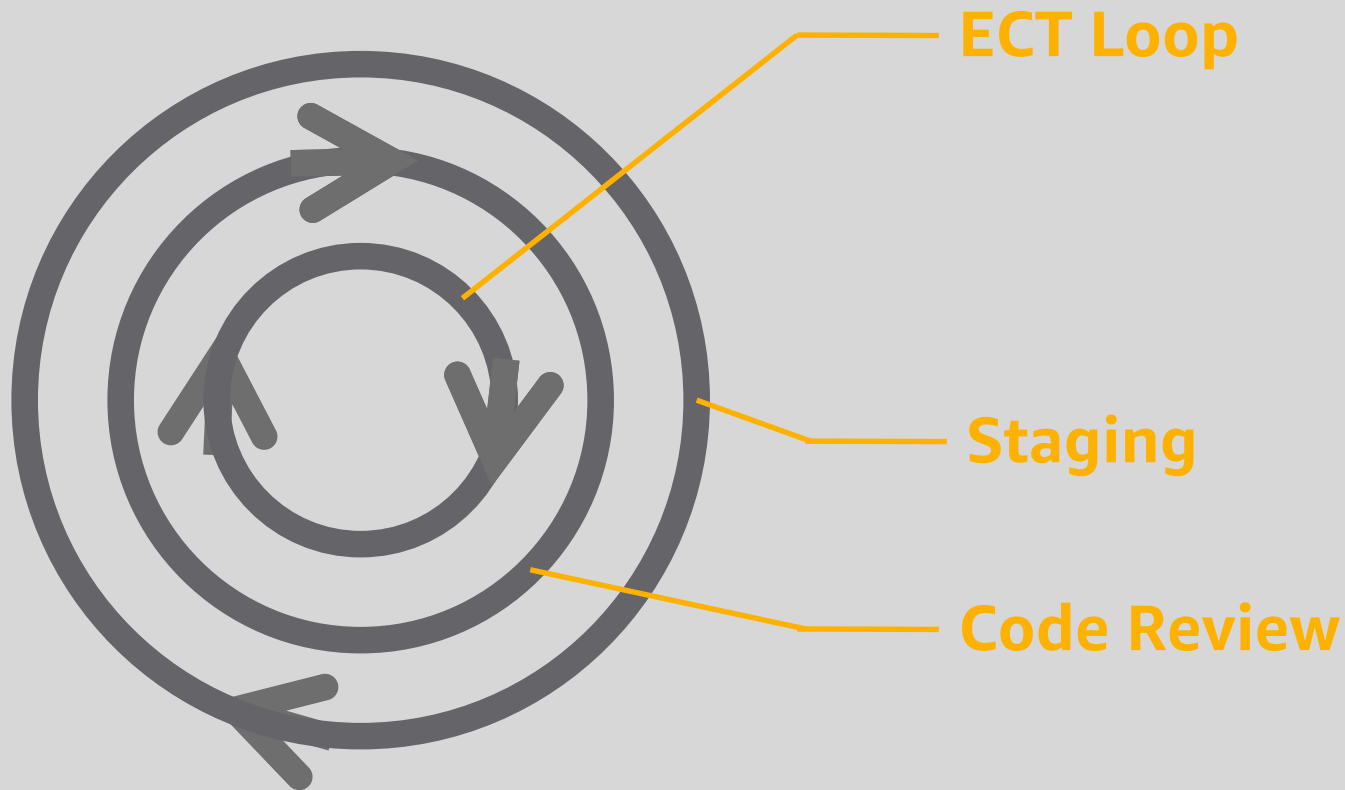
# The Edit-Compile-Test Loop



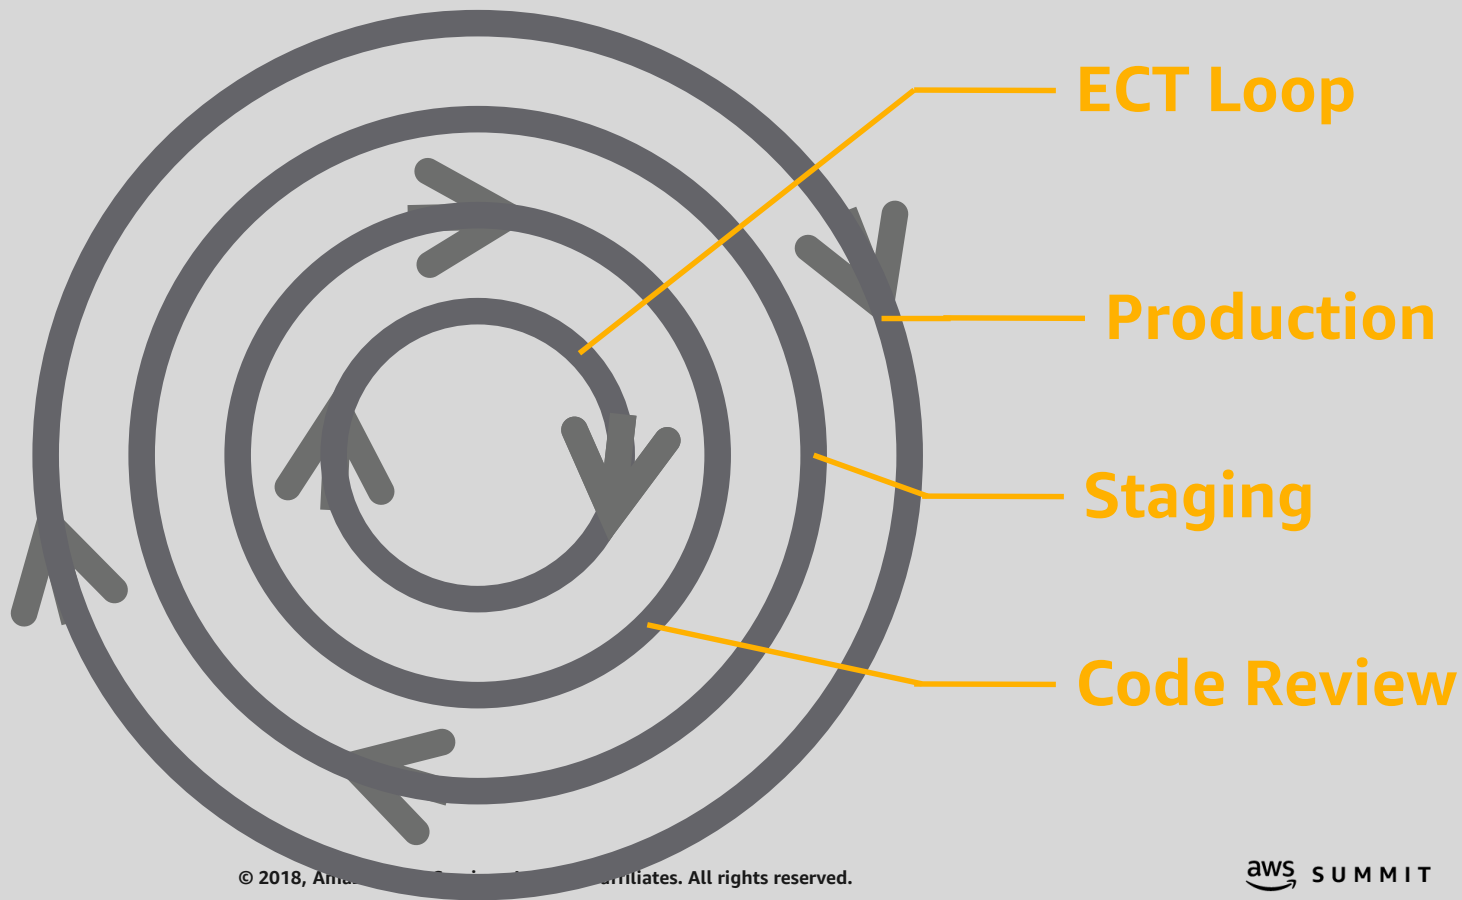
# The Edit-Compile-Test Loop



# The Edit-Compile-Test Loop



# The Edit-Compile-Test Loop





# Catching Problems in the Development Lifecycle



**Ideal Place**



**Great Place**



**Good Place**



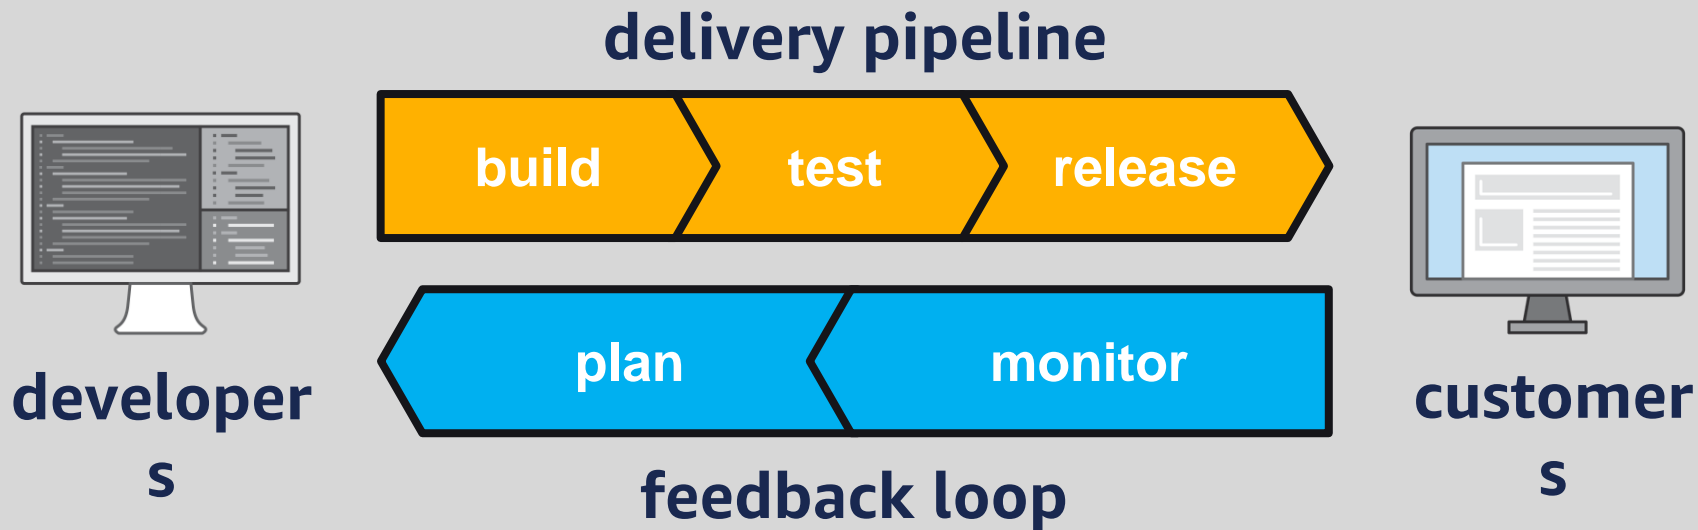
**Last Resort**



# Ten Lessons Learned

1. Microservices
2. DevOps Teams
3. Self-Service Tools
4. Continuous Delivery
5. Pessimistic Deployments
6. Automated Testing
7. Optimize ECT Loop
8. Monitor Everything
9. Measure Everything
10. Listen to Customers

# The Software Development Lifecycle



The screenshot displays the 'Monitoring' page with a grid of 20 charts. The top row shows a summary for 'CPU total', 'Memory total', 'Load (1min)', and 'Processes'. Below this, there are four rows of charts, each representing a different server or application: 'Rails App Server', 'Ganglia', 'Rails App Server', and 'Ganglia'. Each row contains four charts: 'CPU total', 'Memory total', 'Load (1min)', and 'Processes'. The charts show data over a 24-hour period, with a legend indicating the time range (24 hours, 1 hour, 8 hours, 24 hours, 1 week, 2 weeks). The charts are color-coded: CPU total (blue), Memory total (green), Load (1min) (orange), and Processes (purple). The 'CPU total' chart for the first 'Rails App Server' shows a peak in CPU usage around 12h. The 'Memory total' chart for the same server shows a peak in memory usage around 12h. The 'Load (1min)' chart for the same server shows a peak in load around 12h. The 'Processes' chart for the same server shows a peak in processes around 12h. The 'Ganglia' charts show similar trends but with lower overall values.

# Ops Meetings

# Ops Dashboards

- **resource monitoring**
- **application monitoring**
- **canary monitoring**
- **client-side monitoring**
- **“order drop rate”**

# Ten Lessons Learned

1. Microservices
2. DevOps Teams
3. Self-Service Tools
4. Continuous Delivery
5. Pessimistic Deployments
6. Automated Testing
7. Optimize ECT Loop
8. Monitor Everything
9. Measure Everything
10. Listen to Customers



# Data Culture

## WBR Meetings

## Metrics Decks

- feature adoption
- segment growth
- funnel analysis
- cost structure

# Ten Lessons Learned

1. Microservices
2. DevOps Teams
3. Self-Service Tools
4. Continuous Delivery
5. Pessimistic Deployments
6. Automated Testing
7. Optimize ECT Loop
8. Monitor Everything
9. Measure Everything
10. Listen to Customers

# Listen

**It's easy to be right if you just listen to customers.**

- 1. Ask them what they want**
- 2. Get feedback on what you build**
- 3. Course correct if need to**





# Feature vs. Need



# All Things Distributed

Werner Vogels' weblog on building scalable and robust distributed systems.

## Working Backwards

By Werner Vogels on 01 November 2006 05:42 AM | [Permalink](#) | [Comments \(0\)](#)

In the fine grained services approach that we use at Amazon, services do not only represent a software structure but also the organizational structure. The services have a strong ownership model, which combined with the small team size is intended to make it very easy to innovate. In some sense you can see these services as small startups within the walls of a bigger company. Each of these services require a strong focus on who their customers are, regardless whether they are externally or internally. To ensure that a service meets the needs of the customer (and not more than that) we use a process called "*Working Backwards*" in which you start with your customer and work your way backwards until you get to the minimum set of technology requirements to satisfy what you try to achieve. The goal is to drive simplicity through a continuous, explicit customer focus.

The product definition process works backwards in the following way: we start by writing the documents we'll need at launch (the press release and the faq) and then work towards documents that are closer to the implementation.

The Working Backwards product definition process is all about is fleshing out the concept and achieving clarity of thought about what we will ultimately go off and build. It typically has four steps:

1. **Start by writing the Press Release.** Nail it. The press release describes in a simple way what the product does and why it exists - what are the features and benefits. It needs to be very clear and to the point. Writing a press release up front clarifies how the world will see the product - not just how we think about it internally.
2. **Write a Frequently Asked Questions document.** Here's where we add meat to the skeleton provided by the press release. It includes questions that came up when we wrote the press release. You would include questions that other folks asked when you shared the press release and you include questions that define what the product is good for. You put yourself in the shoes of someone using the product and consider all the questions you would have.
3. **Define the customer experience.** Describe in precise detail the customer experience for the different things a customer might do with the product. For products with a user interface, we would build mock ups of each screen that the customer uses. For web services, we write use cases, including code snippets, which describe ways you can imagine people using the product. The goal here is to tell stories of how a customer is solving their problems using the product.
4. **Write the User Manual.** The user manual is what a customer will use to really find out about what the product is and how they will use it. The user manual typically has three sections, concepts, how-to, and reference, which between them tell the customer everything they need to know to use the product. For products with more than one kind of user, we write more than one user manual.

Once we have gone through the process of creating the press release, faq, mockups, and user manuals, it is amazing how much clearer it is what you are planning to build. We'll have a suite of documents that we can use to explain the new product to other teams within Amazon. We know at that point that the whole team has a shared vision on what product we are going to build.



### Contact Info

Werner Vogels  
CTO - Amazon.com

[werner@allthingsdistributed.com](mailto:werner@allthingsdistributed.com)

### Other places

Follow werner on [twitter](#) if you want to know what he is current reading or thinking about.

At [werner.ly](#) he posts material that doesn't belong on this blog or on twitter.

### Syndication

[Subscribe to this weblog's atom feed](#) or [rss feed](#)

### Archives

[All postings](#)

### Recent Entries

[Driving Storage Costs Down for AWS Customers](#)

[Expanding the Cloud - The AWS Storage Gateway](#)

[Amazon DynamoDB - a Fast and Scalable NoSQL Database Service Designed for Internet Scale Applications](#)

[Countdown to What is Next in AWS](#)

[Expanding the Cloud - Introducing the](#)

## “Working Backwards” description by Werner Vogels, Amazon CTO

[www.allthingsdistributed.com/2006/11/working\\_backwards.html](http://www.allthingsdistributed.com/2006/11/working_backwards.html)

# Ten Lessons Learned

1. Microservices
2. DevOps Teams
3. Self-Service Tools
4. Continuous Delivery
5. Pessimistic Deployments
6. Automated Testing
7. Optimize ECT Loop
8. Monitor Everything
9. Measure Everything
10. Listen to Customers

# Thank you

## Two reminders:

**Please complete the session  
survey in the summit mobile app.**

**AWS DevDay with Chef in Berlin  
on  
June 12<sup>th</sup> - <https://bit.ly/2sp8kxo>**