



# Container-based Architectures on AWS

**Sascha Möllering**

**Solutions Architect, @sascha242, Amazon Web Services Germany GmbH**

**Steffen Grunwald**

**Solutions Architect, @steffeng, Amazon Web Services Germany GmbH**

## Amazon EKS – Now Generally Available

by [Jeff Barr](#) | on 05 JUN 2018 | in [Amazon Elastic Container Service For Kubernetes](#), [Launch](#), [News](#) | [Permalink](#) | [Comments](#) | [Share](#)

We announced [Amazon Elastic Container Service for Kubernetes](#) and invited customers to take a look at a preview during re:Invent 2017. Today I am pleased to be able to let you know that Amazon EKS is available for use in production form. It has been certified as Kubernetes conformant, and is ready to run your existing Kubernetes workloads.

Based on the [most recent data](#) from the [Cloud Native Computing Foundation](#), we know that AWS is the leading environment for Kubernetes, with 57% of all companies who run Kubernetes choosing to do so on AWS. Customers tell us that Kubernetes is core to their IT strategy, and are already running hundreds of millions of containers on AWS every week. Amazon EKS simplifies the process of building, securing, operating, and maintaining Kubernetes clusters, and brings the benefits of container-based computing to organizations that want to focus on building applications instead of setting up a Kubernetes cluster from scratch.

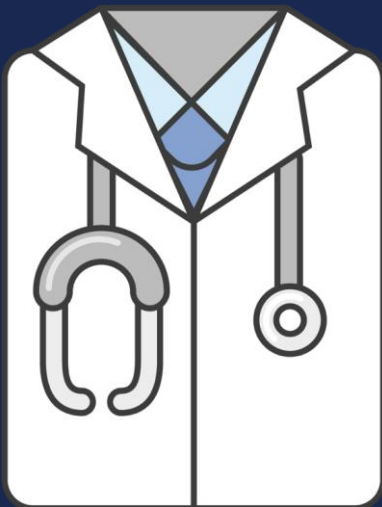
# We started in 2014



# Customers running Docker on EC2 from the very beginning...



# But there were pain points.



Things like scheduling, placing, managing and deploying containers were difficult.

They wanted something to make those pain points better.

# So we built ECS



Amazon Elastic  
Container  
Service

Highly scalable,  
high performance  
container  
management  
system



Cluster  
managemen  
t

A managed  
platform



Container  
orchestratio  
n



Deep AWS  
integration

# So we built ECS



Amazon Elastic  
Container  
Service

Highly scalable,  
high performance  
container  
management  
system



**AWS VPC  
networking mode**



**Advanced task  
placement**



**Deep integration  
with AWS platform**



**ECS CLI**



**Global footprint**



**Powerful  
scheduling engines**



**Auto scaling**

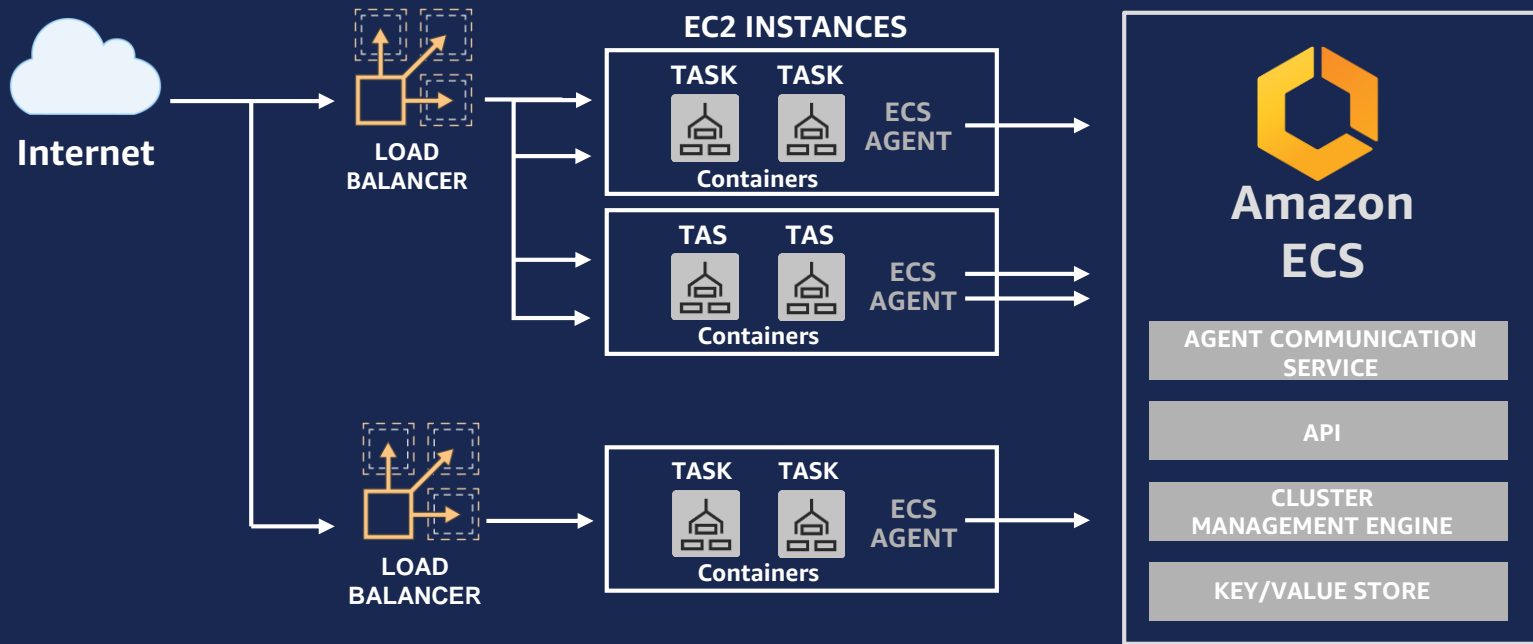


**CloudWatch metrics**



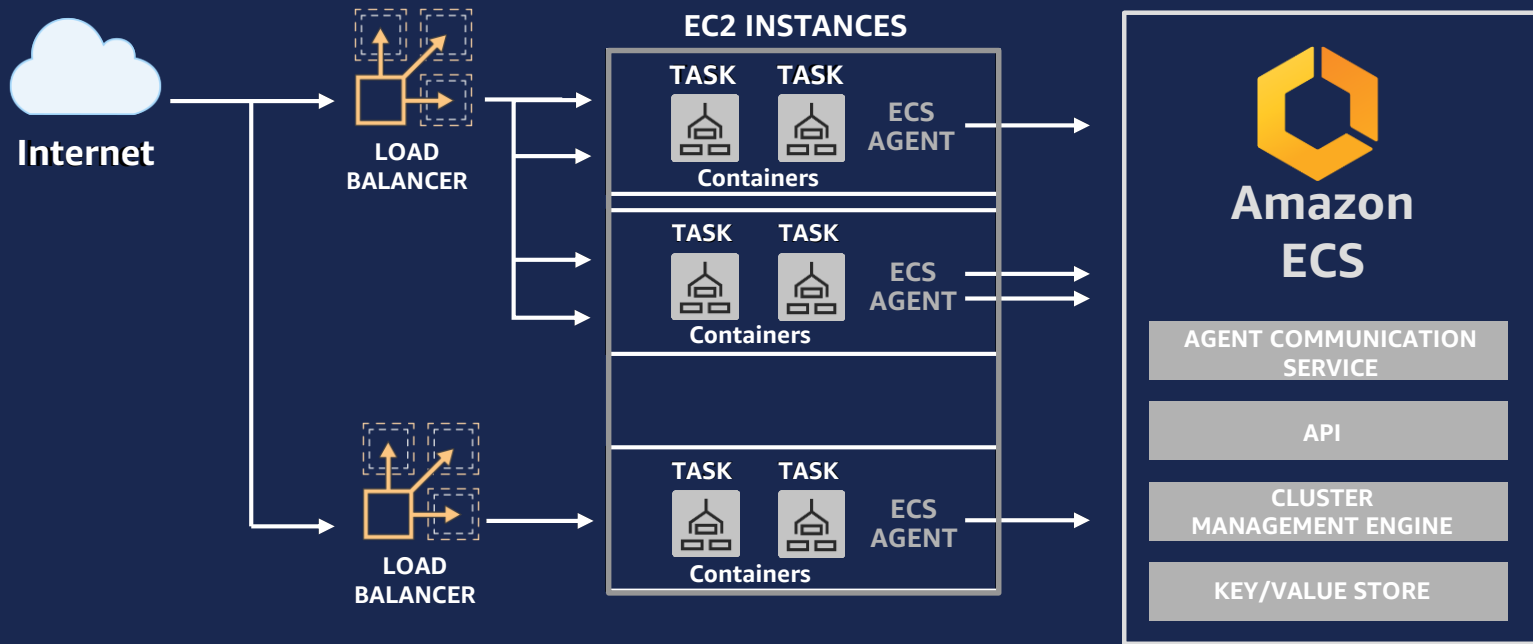
**Load balancers**

# Amazon ECS

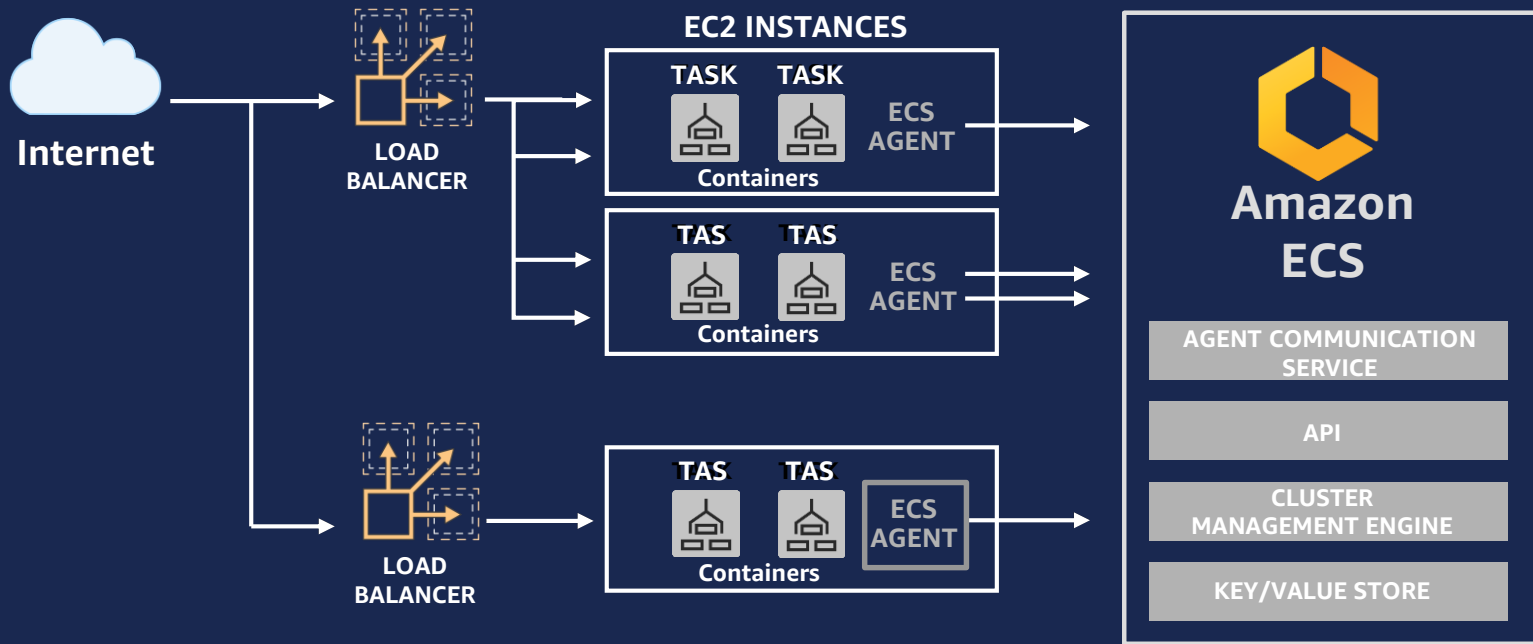




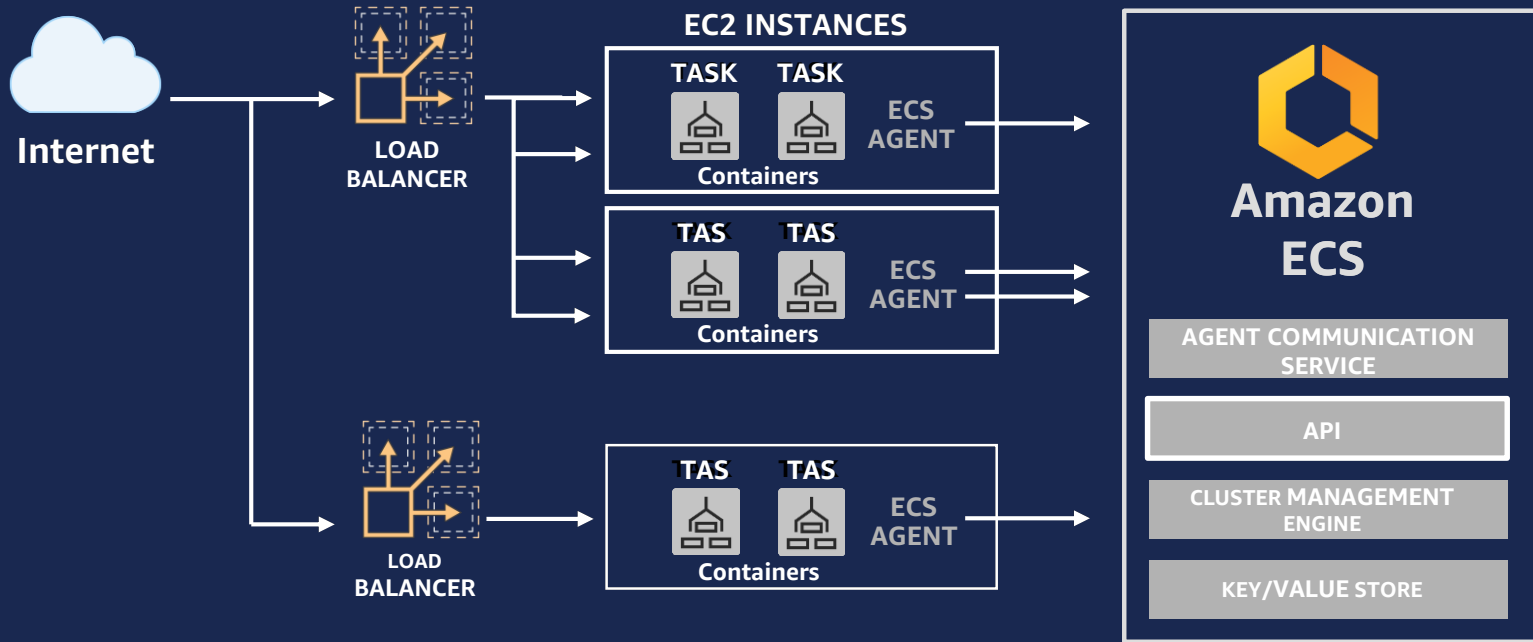
# Cluster of hosts



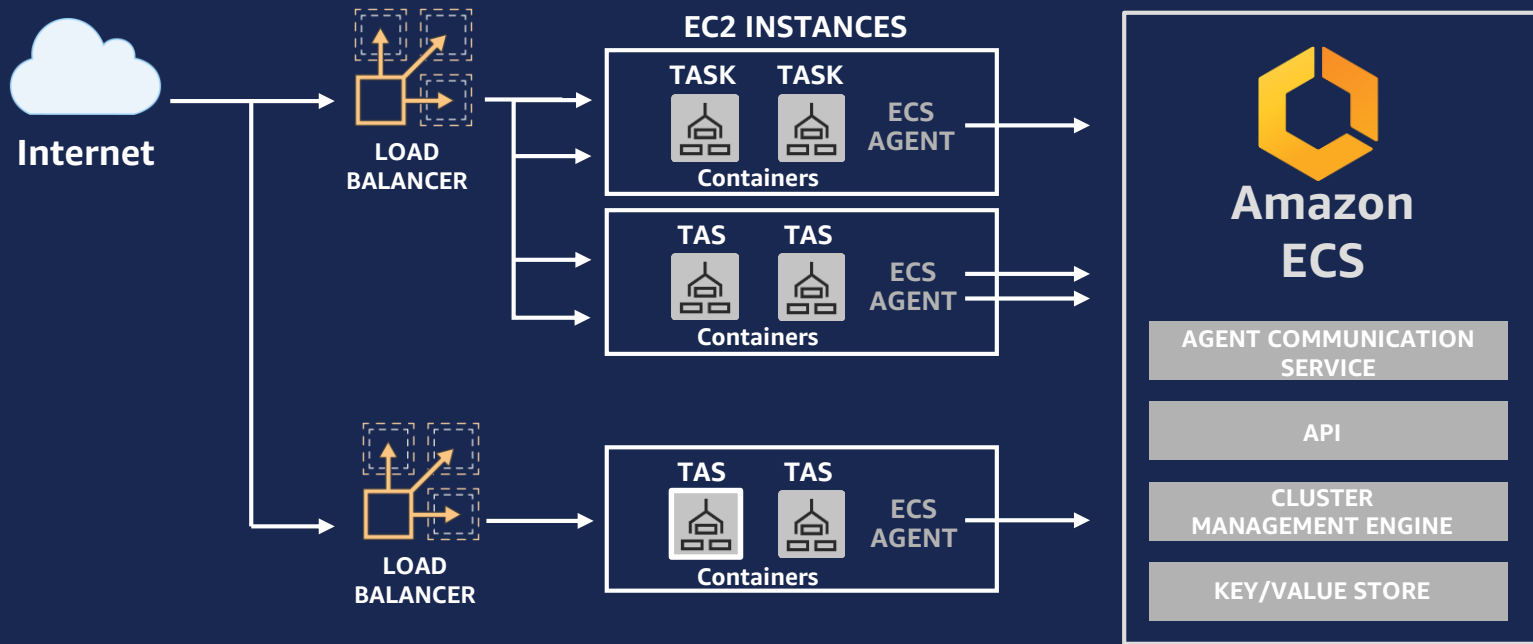
# Lightweight agent on each host



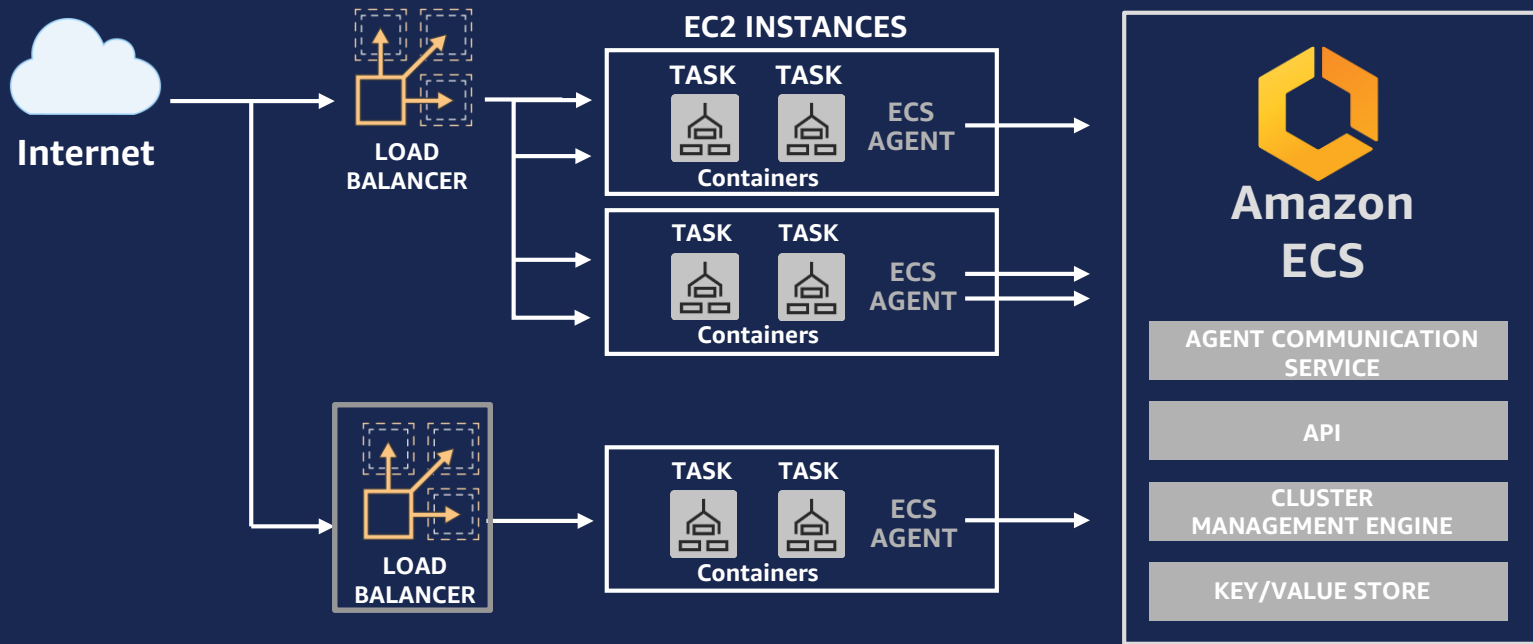
# API for launching containers on the cluster



# Container task is placed on a host



# Traffic is sent to your host



# ECS Optimized Amazon Machine Images (AMIs)

- Optimized AMIs available for Linux & Windows
- Bring your own images based on it
- Expects ECS cluster name in user-data
- Update images on SNS update notifications

 Menu



Contact Sales

Products ▾

Solutions

Pricing

More ▾

English ▾

My Account ▾

Sign In to the Console

ABOUT AWS

About AWS >

Global Infrastructure >

What's New >

AWS in the News >

Events & Webinars >

RELATED LINKS

Amazon ECS Support for Windows Server Containers is Generally Available

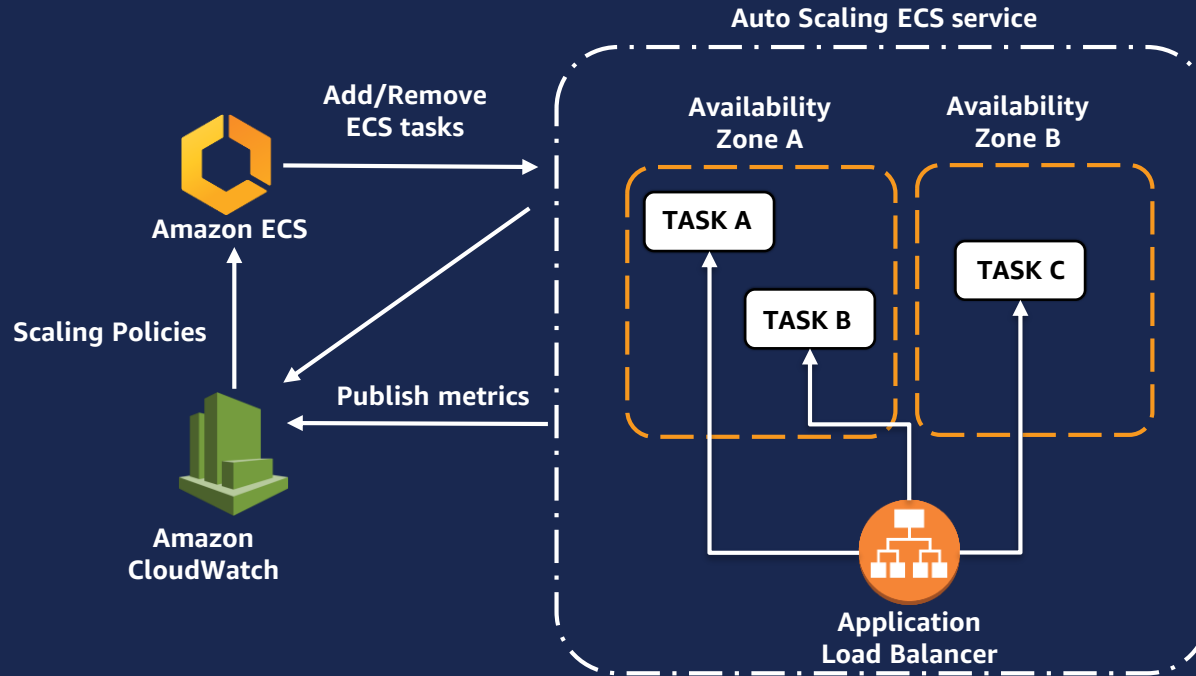
Posted On: Dec 5, 2017

Amazon Elastic Container Service (Amazon ECS) now supports running Windows Server containers for production workloads.

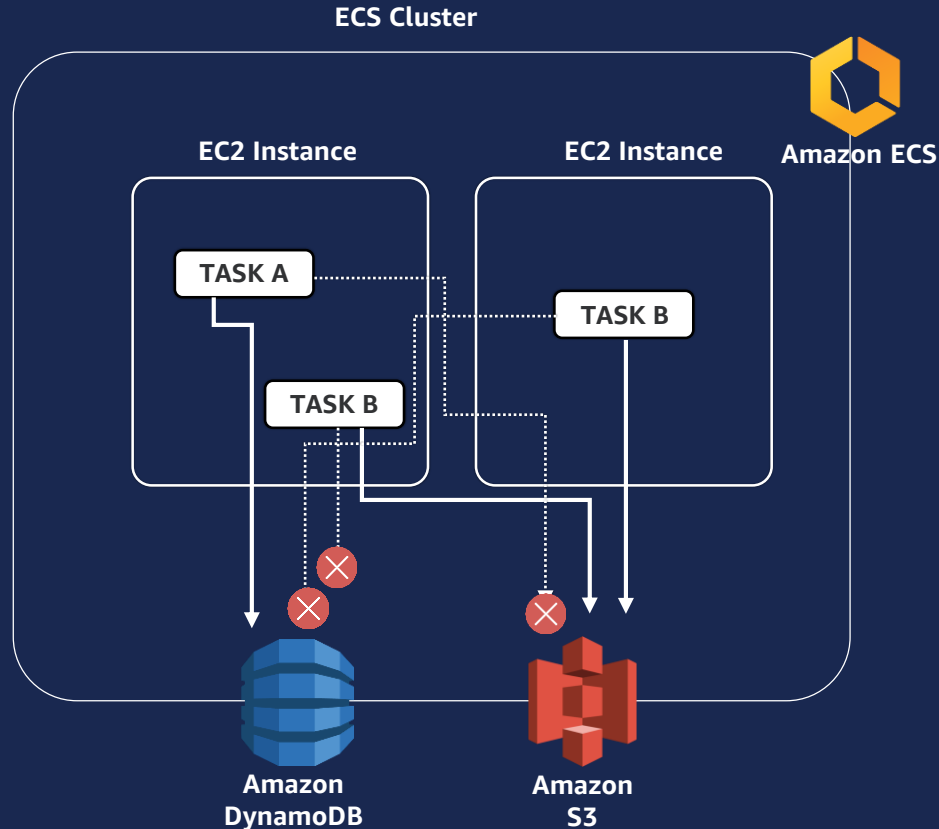
Previously, Windows containers were supported by Amazon ECS as a [public beta](#).

Now, Amazon ECS provides an ECS-Optimized Windows Server Amazon Machine Image (AMI), which is based on the EC2 Windows Server 2016 AMI and includes Docker 17.06 Enterprise Edition and the ECS Agent 1.16. This AMI provides improved instance and container launch time performance. The ECS Agent runs as a Windows Service, which provides an improved ECS Agent lifecycle

# Automatic Service Scaling

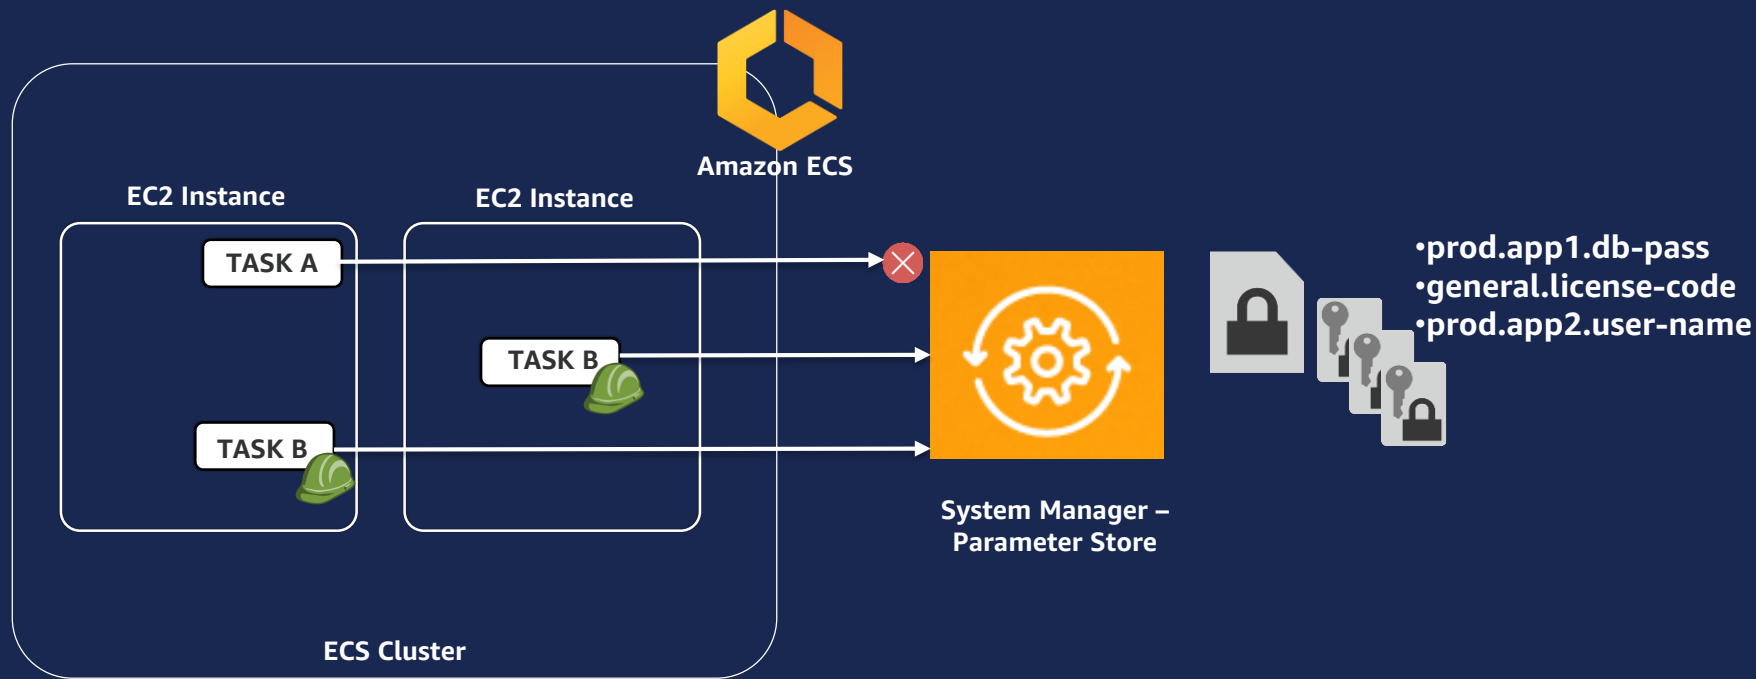


# IAM Roles For Tasks





# Secrets Management



# Amazon Elastic Container Registry (Amazon ECR)

- Cloud-based Docker image registry
- Fully managed
- Secure – images encrypted at rest, integrated with IAM
- Scalable and Highly Available
- Integrated with Amazon ECS and the Docker CLI



Amazon ECR

# AWS Fargate



**Underlying technology for  
container management**



**No cluster or  
infrastructure to  
manage or  
scale**



**Everything is  
handled at the  
container level**



**Scale  
seamlessly on  
demand**

# What does Fargate mean?

**Not worrying about scaling, underlying infrastructure, cluster resources, capacity, setup.**

**Just give it a task definition or pod (in 2018), set some resource limits, and away you go.**



# Task Definitions Repository on GitHub

The screenshot shows the GitHub repository page for `aws-samples / aws-containers-task-definitions`. The repository has 40 stars, 23 unstars, and 8 forks. It is currently on the `master` branch. The repository description is "Task Definitions for running common applications Amazon ECS" with a link to <https://aws.amazon.com/ecs>. The repository includes tags for `amazon`, `ecs`, `task`, `containers`, `container`, `docker`, `docker-container`, `elasticcontainerservice`, and `ec2-container-service`. It has 21 commits, 1 branch, 0 releases, 4 contributors, and is licensed under Apache-2.0. The repository was updated by `smoell` on 15 Dec 2017. The file list includes `.github`, `consul`, `gunicorn`, `jetty`, `kibana`, `nginx`, `tomcat`, `wildfly`, `LICENSE`, `NOTICE`, and `README.md`.

aws-samples / aws-containers-task-definitions

Unwatch 40 Unstar 23 Fork 8

<> Code Issues 0 Pull requests 0 Projects 0 Insights

Task Definitions for running common applications Amazon ECS <https://aws.amazon.com/ecs>

amazon ecs task containers container docker docker-container elasticcontainerservice ec2-container-service

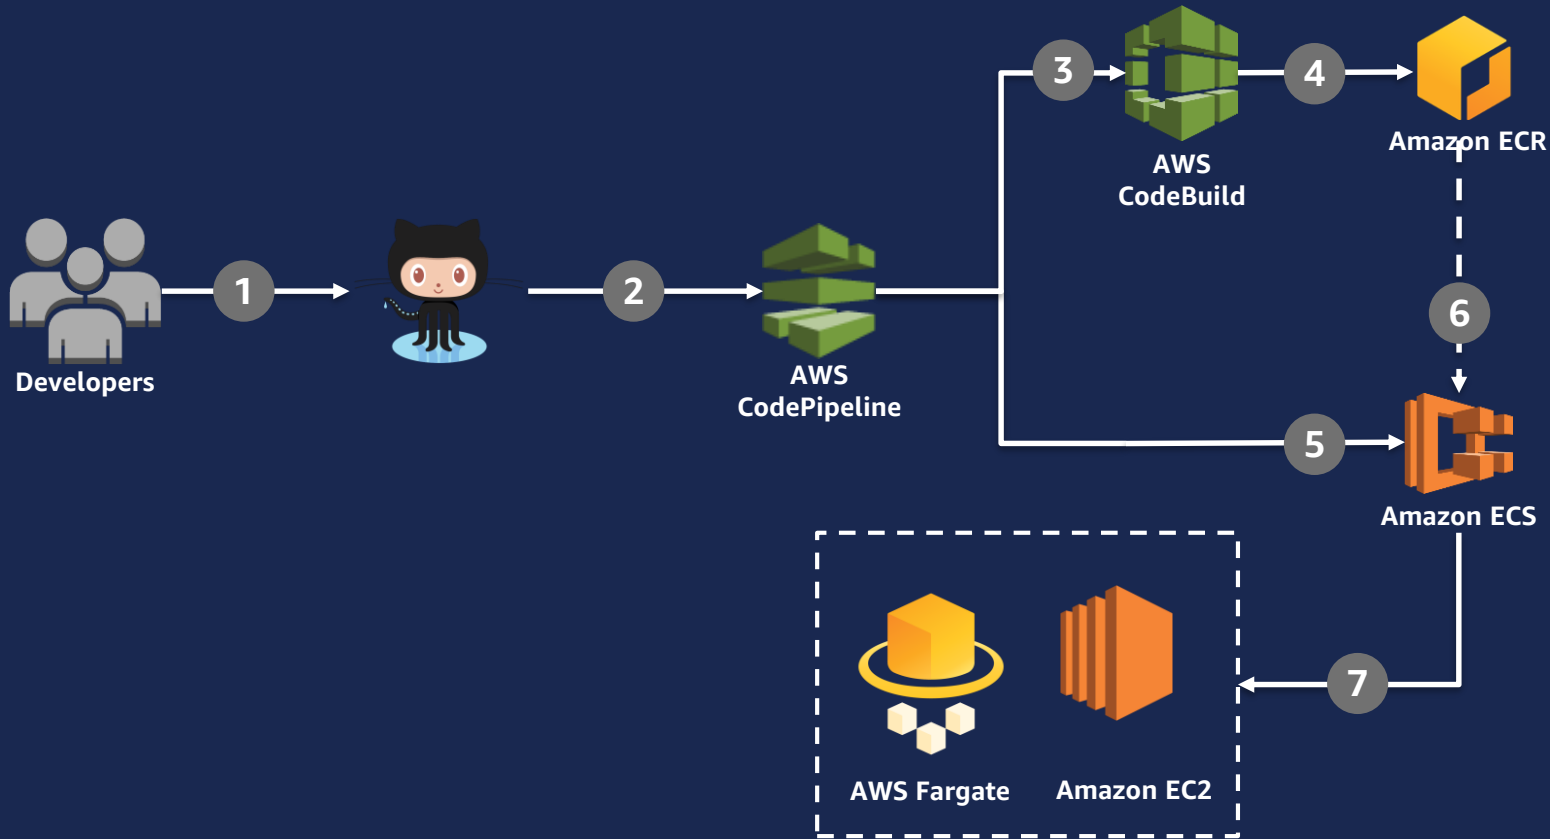
21 commits 1 branch 0 releases 4 contributors Apache-2.0

Branch: master New pull request Create new file Upload files Find file Clone or download

smoell Update README.md Latest commit 8b7ed1d on 15 Dec 2017

.github	Creating initial file from template	4 months ago
consul	Add files via upload	3 months ago
gunicorn	Cleanup, fixed log groups	3 months ago
jetty	Cleanup, fixed log groups	3 months ago
kibana	Cleanup, fixed log groups	3 months ago
nginx	first commit	4 months ago
tomcat	Tomcat fix for awsvpc, update documentation	3 months ago
wildfly	Cleanup, fixed log groups	3 months ago
LICENSE	Creating initial file from template	4 months ago
NOTICE	Creating initial file from template	4 months ago
README.md	Update README.md	3 months ago

# Continuous Deployment in Amazon ECS





**Europe's leading app for ordering taxis**  
**Majority of services on ECS**



**10+ million users with 45,000+ taxis across 40+ cities.**

**With the microservice architecture (140+ services) built on AWS, mytaxi can provide new features to users faster than ever before.**

**Running entirely on Spot.**

© 2018, Amazon Web Services, Inc. or its a







„In November 2015 we moved our Docker **container architecture to Amazon ECS**, and for the first time ever in December we were able to celebrate a new year in which our system could handle the **huge number of requests** without any crashes or interruptions.“

–Sebastian Herzberg, System Engineer



kubernetes



“Run Kubernetes for me.”

# Amazon Elastic Container Service for Kubernetes: EKS



**Managed Kubernetes on  
AWS**



**Managed  
Kubernetes  
control  
plane**



**Highly  
available**



**Automated  
version  
upgrades**



**Integration  
with other  
AWS services**

CloudTrail,  
CloudWatch, ELB,  
IAM, VPC,  
PrivateLink

# Elastic Container Service for Kubernetes

- Platform for enterprises to run production grade workloads
- Provides a native and upstream Kubernetes experience
- Not forced to use additional AWS services, but offer seamless integration
- EKS team actively contributes to the Kubernetes project

# Elastic Container Service for Kubernetes

EKS > Clusters > jeff1

jeff1 Delete

**jeff1 settings**

API server endpoint https://[redacted].us-west-2.eks.amazonaws.com <span>copy</span> <span>copied to clipboard</span>	Cluster ARN arn:aws:eks:us-west-2:[redacted]:cluster/demo <span>copy</span>
Kubernetes version 1.10	Status <span>ACTIVE</span>
Role ARN arn:aws:iam:[redacted]:role/eks-service-role <span>copy</span>	Subnets subnet-3db5cb67, subnet-ee1695a5, subnet-fd7d1c84
Security groups sg-4b43d83a	Certificate authority [redacted] <span>copy</span>

# EKS Architecture



# Jenkins – CI/CD with Kubernetes

```
1 node {
2
3   stage 'Checkout'
4   git 'https://github.com/omarlari/aws-container-sample-app.git'
5
6   stage 'Build Dockerfile'
7   docker.build('hello')
8
9   stage 'Push to ECR'
10  sh ("eval \$(docker run awscli aws ecr get-login --region ${REGION} --no-include-email | sed '
11  docker.withRegistry('https://${ECR_REPO}') {
12    docker.image('hello').push('${BUILD_NUMBER}')
13  }
14
15  stage 'update application'
16
17  kubernetes: { node {
18  docker.image('kubectl').inside("--volume=/home/ec2-user/.kube:/config/.kube"){
19    sh 'kubectl describe deployment ${APP}'
20    sh 'kubectl set image deployment/${APP} hello=${ECR_REPO}/hello:${BUILD_NUMBER}'
21    sh 'kubectl describe deployment ${APP}'
22  }
23  }}
24 }
```

## Pipeline hello-jenkins

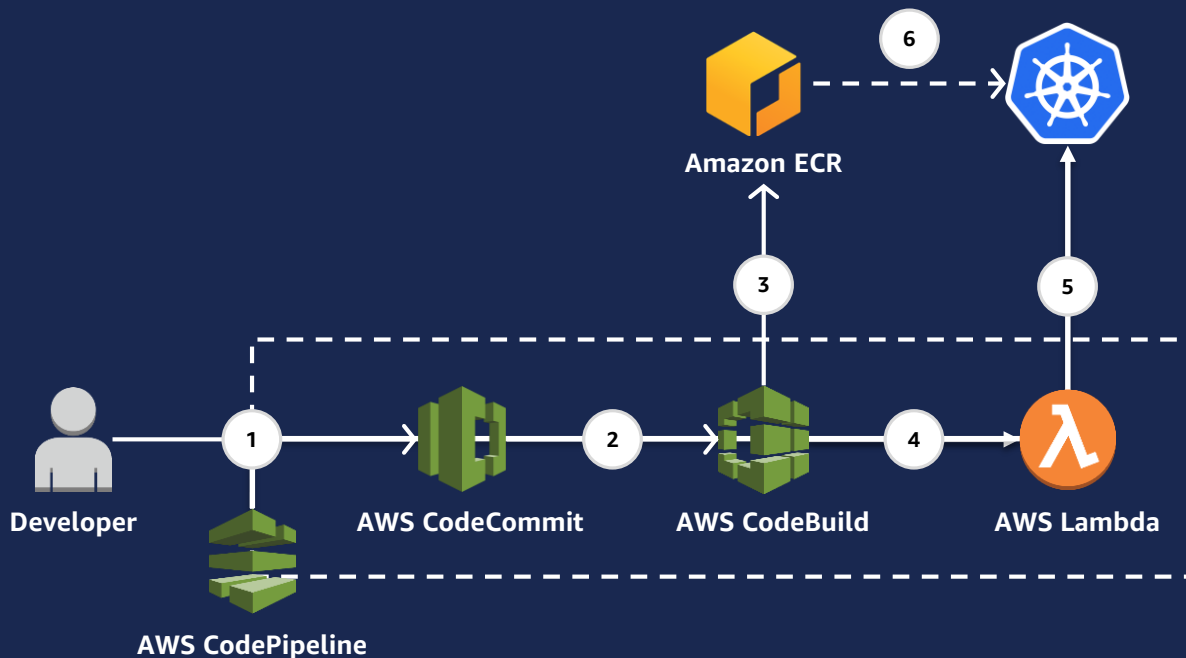


### Stage View

	Checkout	Build Dockerfile	Push to ECR	update application
Average stage times: (Average full run time: ~7s)	305ms	9s	13s	4s
#6 Nov 10 18:46 No Changes	242ms	552ms	2s	3s
#5 Nov 10 18:37 No Changes	328ms	808ms	3s	3s failed
#4 Nov 10 18:28 No Changes	220ms	589ms	2s	2s failed
#3 Nov 10 18:22 No Changes	238ms	631ms	1min 9s	8s failed
#2 Nov 10 18:18 No Changes	258ms	966ms	3s	failed



# AWS CodePipeline – CI/CD with Kubernetes



# Batch Processing with Containers

# What is batch computing?

**Run jobs  
asynchronously and  
automatically across  
one or more  
computers.**



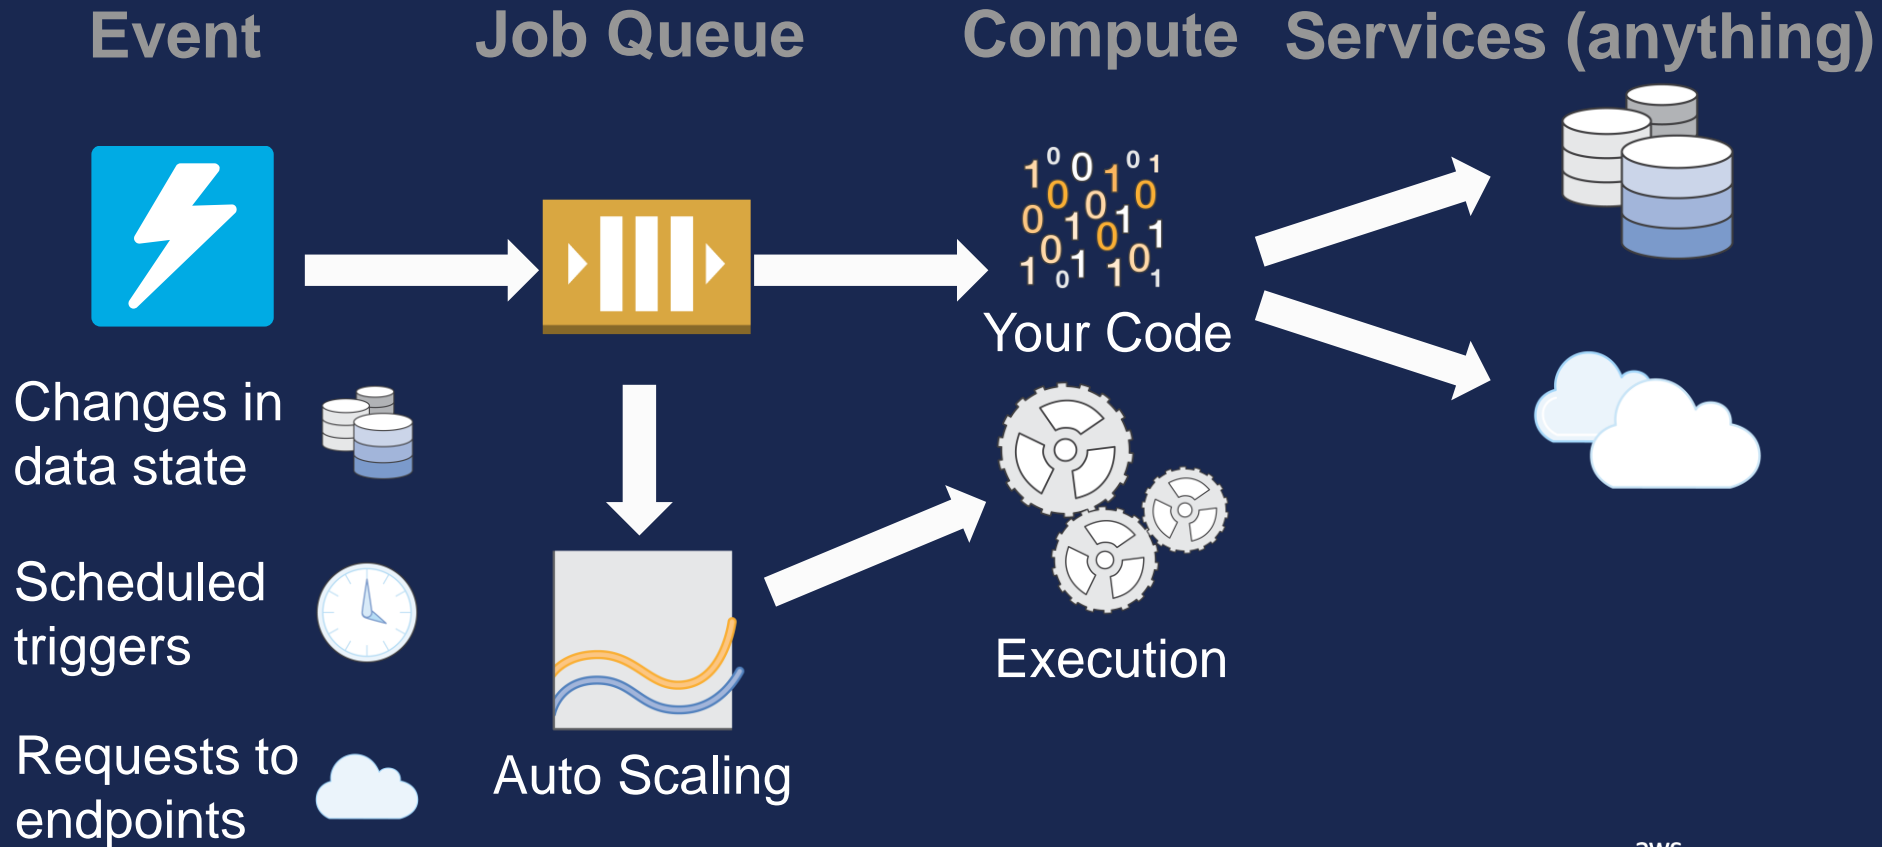
**Jobs may have  
dependencies,  
making the  
sequencing and  
scheduling of  
multiple jobs complex  
and challenging.**



# Cloud makes Sense for Batch

- Scalable
- Reliable
- Choice:
  - Compute resources (GPUs, RAM- or CPU-bound)
  - Storage resources (filesystems, performance characteristics)
  - Downstream services (e.g. databases, streaming services)
  - Pricing models
- Pay as you go (per second)

# Anatomy of a Batch



# Options for Batch Workloads on AWS



**Amazon ECS**



**AWS Batch**

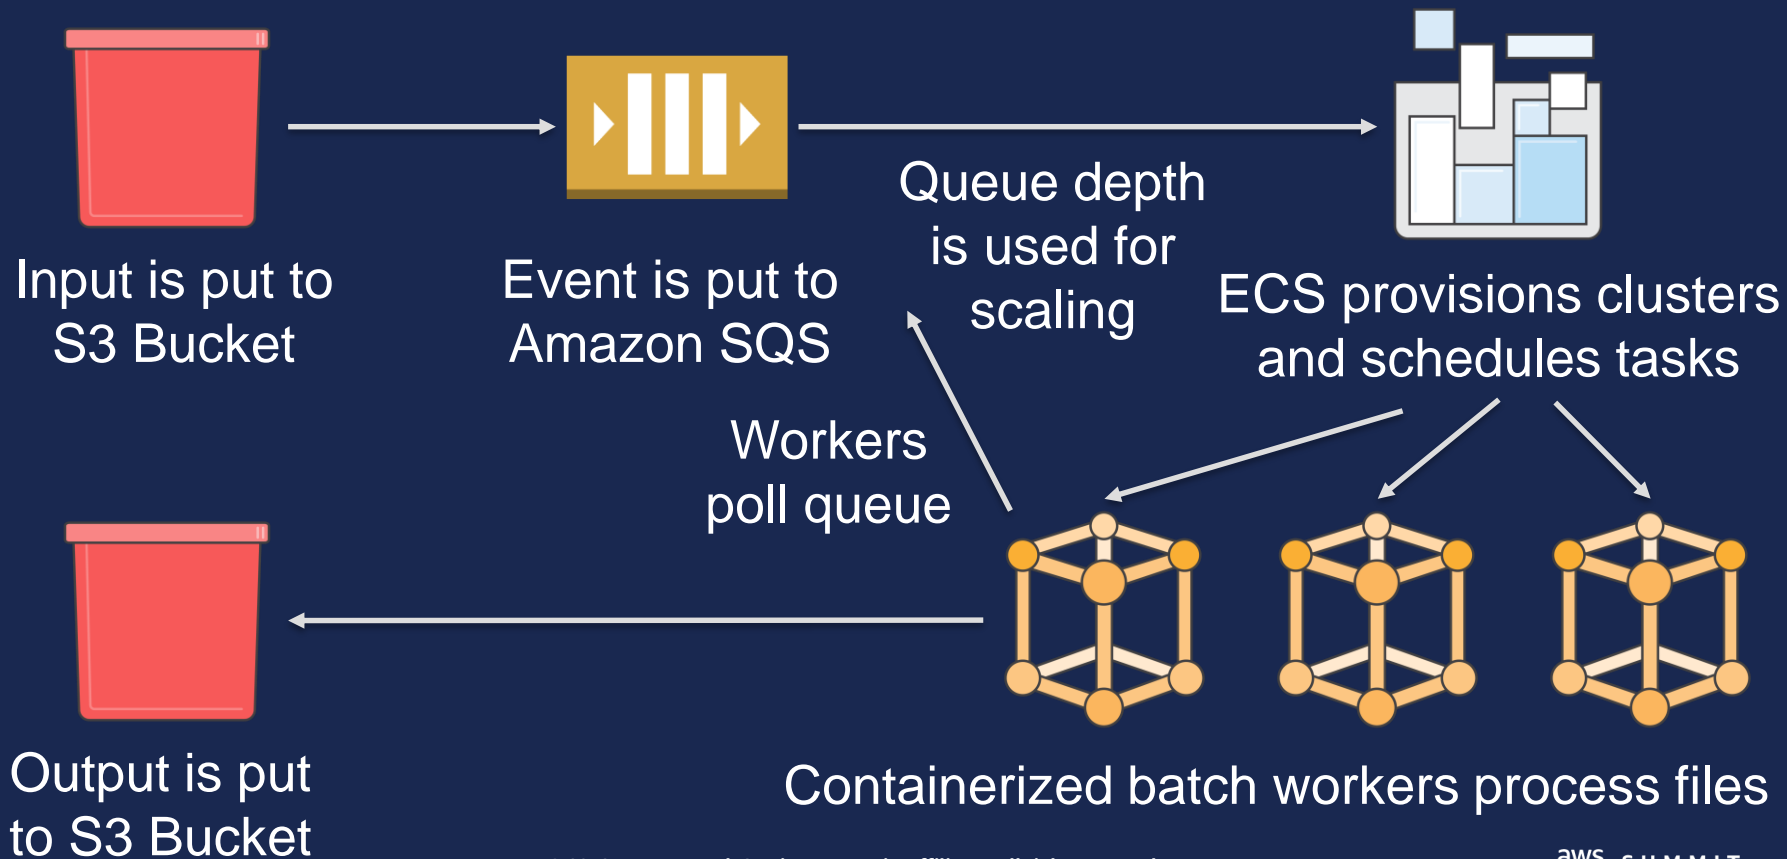


**AWS Lambda**  
(Amazon SQS as built-in  
event source coming soon)

# Containers make Sense for Batch

- **Benefits from container development model**
- **Polyglot**
- **Do one thing well**
- **Black Box – and easy to model**

# Basic Batch Workflow with ECS



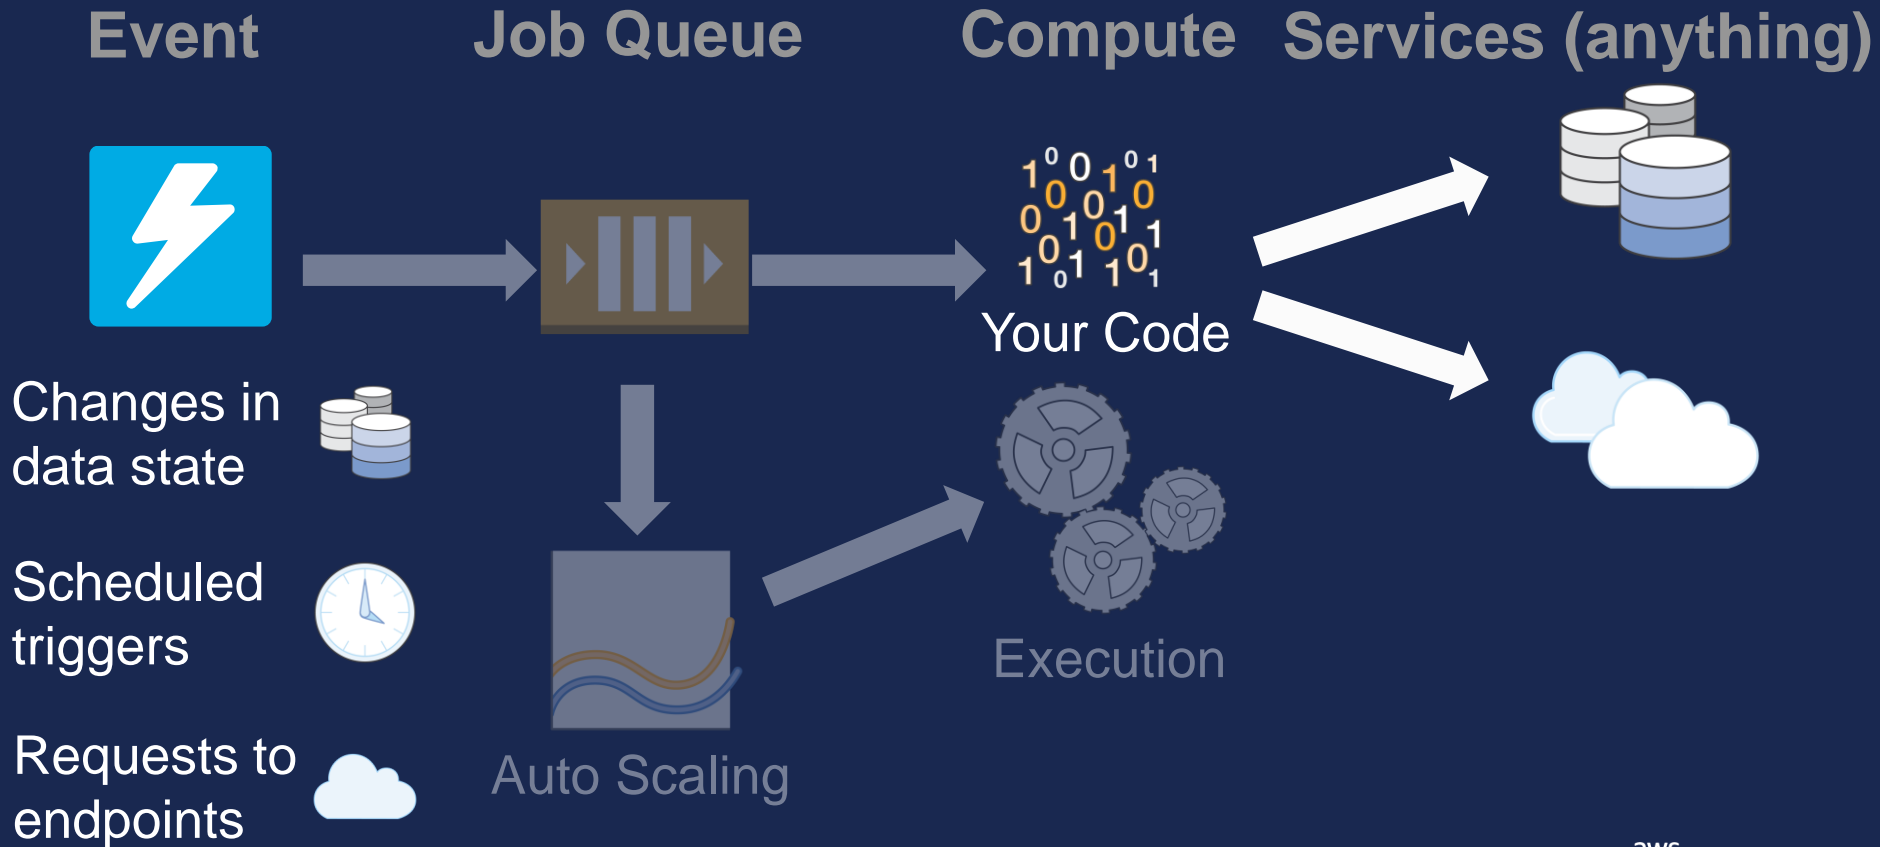


# Introducing AWS Batch

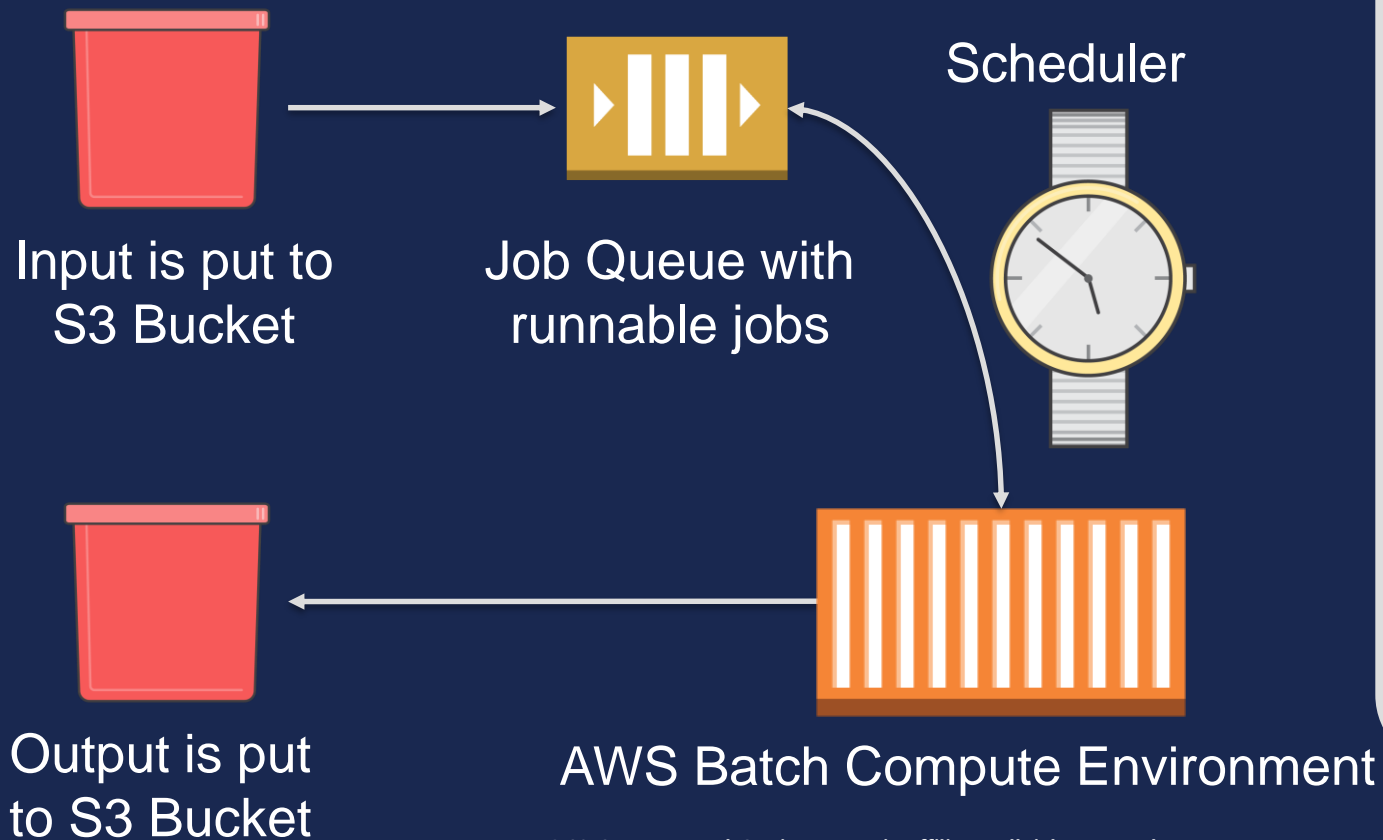
- Fully managed batch primitives
- Focus on your applications (shell scripts, Linux executables, Docker images) and their resource requirements
- We take care of the rest!



# What you need to take care of ...



# Typical AWS Batch Job Architecture



Job definition



Application Image + config



IAM role

# „Flip Image“ Job's Dockerfile

```
FROM amazonlinux:latest
```

```
RUN yum update -y
```

```
RUN yum install ImageMagick aws-cli -y
```

```
ADD flip.sh /usr/local/bin/flip.sh
```

```
WORKDIR /tmp
```

```
USER nobody
```

```
ENTRYPOINT ["/usr/local/bin/flip.sh"]
```

# flip.sh

```
#!/bin/bash
```

```
SRCBKT=$1 # SOURCE BUCKET
```

```
OBJ=$2 # OBJECT KEY
```

```
TRGBKT=$3 # TARGET BUCKET
```

```
error_exit() { echo "${1}" >&2; exit 1; }
```

```
tmpfile=$(mktemp /tmp/image.XXXXXXX)
```

```
aws s3 cp "s3://$SRCBKT/$OBJ" "$tmpfile" \ # DOWNLOAD
```

```
|| error_exit "Download failed $SRCBKT/$OBJ"
```

```
convert -flip "$tmpfile" "$tmpfile" \ # FLIP w/ IMAGEMAGICK
```

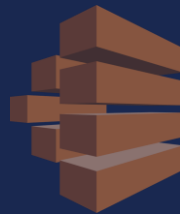
```
|| error_exit "Failed to flip file"
```

```
aws s3 cp "$tmpfile" "s3://$TRGBKT/$OBJ" \ # UPLOAD
```

```
|| error_exit "Upload failed $TRGBKT/$OBJ"
```

# AWS Batch Concepts

- Job definitions
- Jobs
- Job queue
- Compute environments



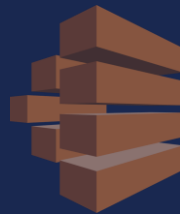
# Job Definitions

AWS Batch **job definitions** specify how jobs are to be run.

Some attributes in a job definition:

- **Container Image**
- IAM role associated with the job
- vCPU and memory requirements
- Mount points
- Environment variables
- Retry strategy

# Jobs



**Jobs** are the unit of work executed by AWS Batch.

Set/ overwrite **Job Definition** attributes, e.g.:

- Command
- Parameters
- Dependencies

## E.g. job's Command:

Space delimited

JSON

```
["batch-demo-input","steffeng.jpg","batch-demo-output"]
```

This JSON array is passed to the job as your command to execute.

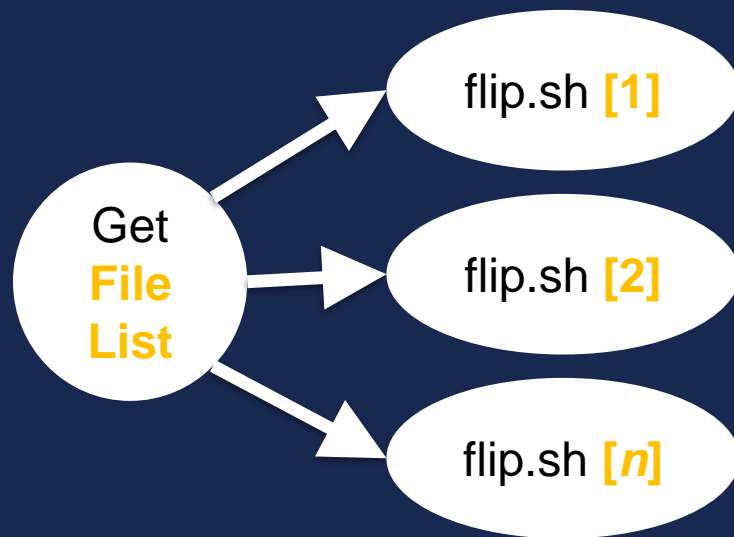


# Easily run massively parallel jobs

Start up to 10,000 copies of an application with a single call using **Array Jobs**.

Efficient way to run:

- Parametric sweeps
- Monte Carlo simulations
- Processing a large collection of objects



# Job Retries

AWS Batch supports **up to 10 attempts** per job:

- errors in the AWS Batch job
- termination of the Spot Instance

**The `AWS_BATCH_JOB_ATTEMPT` environment variable is set to the container's corresponding job attempt number.**

# Compute Environments

## Managed

**AWS scales and configures your instances for you.**

### Optional choice:

- **On demand/ Spot**
- **Instance Types/ Mix**
- **Amazon Machine Image**

## Unmanaged

**You control and manage the instance configuration, provisioning, and scaling.**

**Full control over scaling and instance provisioning for the ECS cluster used by AWS Batch.**

# Bring your own AMIs

**Customer Provided AMIs** let you set the AMI that is launched as part of a **managed compute environment**.

Makes it possible to configure Docker settings, mount **EBS/ EFS** volumes, and configure drivers for **GPU** jobs.

AMIs must be Linux-based, HVM and have a working ECS agent installation.

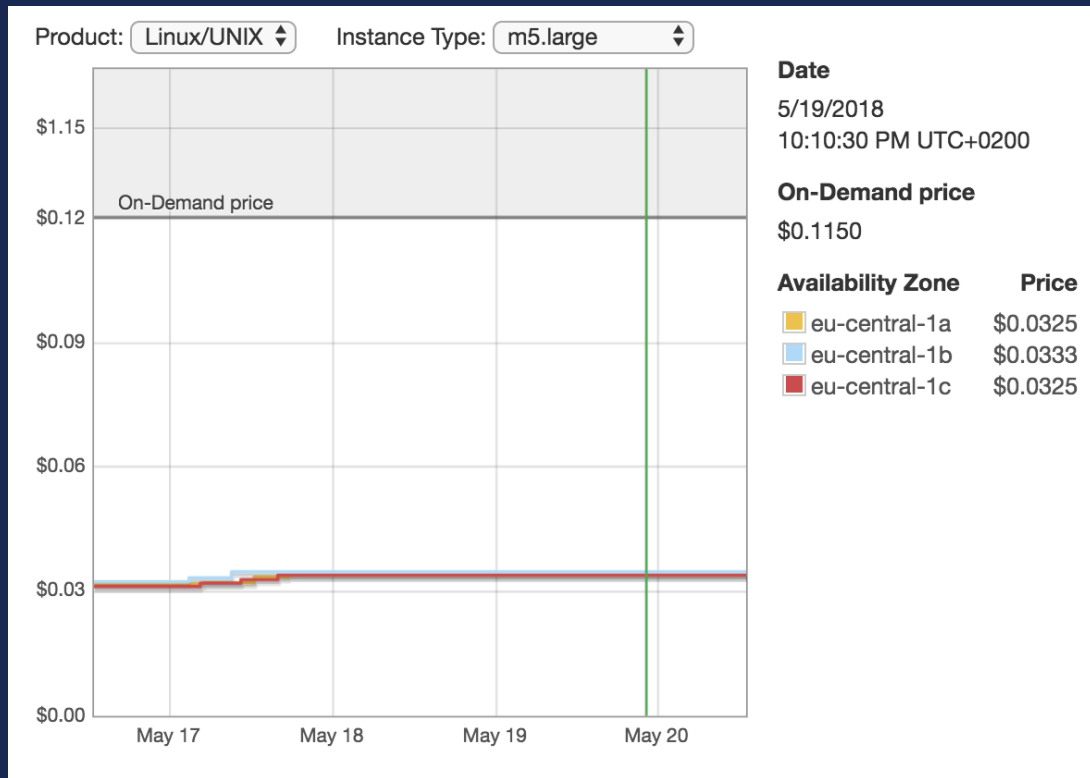
# Job Queues

Jobs are submitted to **Job Queues**.

Job queues are assigned to one or more compute environments.

Each job queue has a priority assigned. Jobs in queues with higher priority take precedence.

# Use Spot Instances



Get your jobs done  
**faster** or **cheaper**.

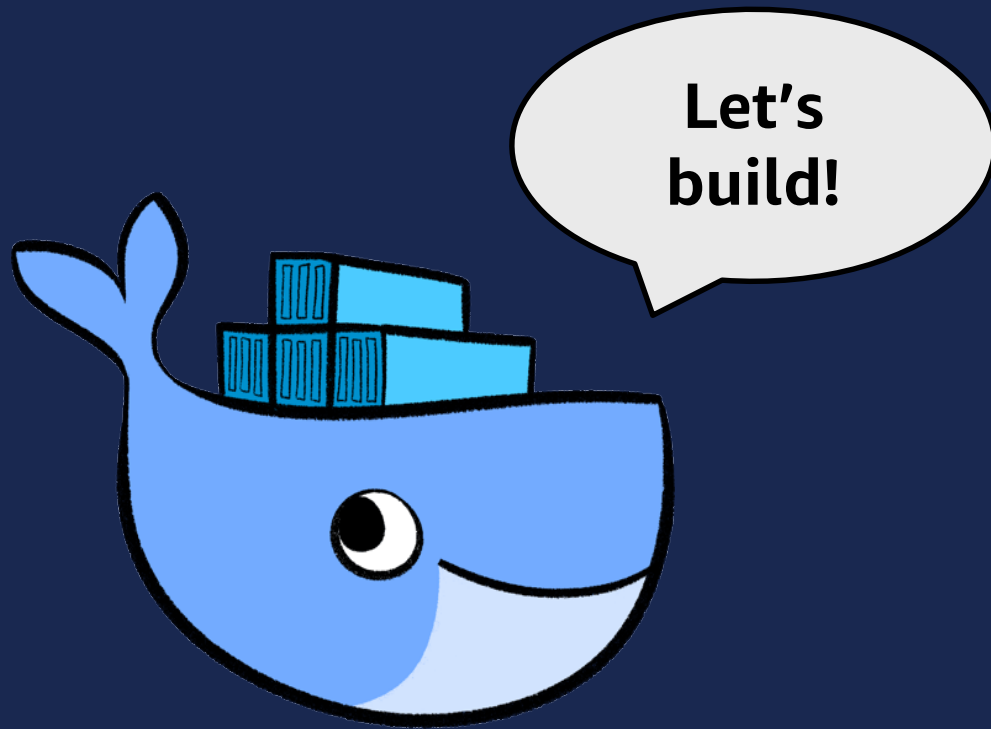
AWS Batch retries  
jobs on instance  
terminations and  
selects from  
multiple instance  
types.

# Pricing

**No additional charge for AWS Batch or Amazon ECS.**

**You only pay for the underlying resources that you consume!**







Questions? Ask these  
guys at the **Ask an  
Architect** booth:



Please **complete the session survey**  
in the summit mobile app.

Next session in this room:  
**14:00 – Kubernetes Running on AWS**

# Thank you!